Confidential

Development of a Versioned SQL Database System for Parameter Management in the Common Powertrain Controller

Entwicklung eines versionierten SQL-Datenbanksystems für das Parametermanagement im Common Powertrain Controller

Master's thesis

In the study program
Mechanical Engineering

by
Mahesh Kollati
Matr.-Nr. 3635029

Examiner: Prof. Dr.-Ing. Hans-Christian Reuss

Institutional Supervisor: Dan Greiner

External Supervisor: Nikolas Schönfelder,
                            Daimler Truck AG

Submitted to the University of Stuttgart
Institute of Automotive Engineering Stuttgart
Chair in Automotive Mechatronics
2025

Universität Stuttgart
Institut für Fahrzeugtechnik Stuttgart

Automotive Mechatronics
Prof. Dr.-Ing. H.-C. Reuss

+49 711 685 65601
+49 711 685 65710
info@ifs.uni-stuttgart.de

Master's thesis
for
Mr. Mahesh Kollati

Matr.-Nr. 3635029

Title: Development of a Versioned SQL Database System for Parameter Management in the Common Powertrain Controller

Entwicklung eines versionierten SQL-Datenbanksystems für das Parametermanagement im Common Powertrain Controller

In current vehicles, the coordination of variant diversity plays an important role. For cost reasons, attempts are made to use common parts across model series and for different equipment variants within a model series. For control units, this approach also extends to the software, where possible. Different characteristics are then implemented using parameters. Parameter management is typically done using databases.

This master's thesis aims to develop a concept for an SQL database for managing parameters for a central control unit in vehicle applications. Various aspects shall be considered.
First, various versioning approaches shall be tested and evaluated. The goal is to document changes in a traceable manner.
Based on the requirements of software developers, various use cases shall be defined. Based on this, a concept for the tables within the database shall be developed.
Furthermore, it shall be described how the database can later be used by both software developers and development engineers. This also includes the development of a user interface.
Finally, the limitations and restrictions of the developed approach should be described.

The results must be documented and presented.

Confidentiality note: This thesis must be treated confidentially until .

Institutional Supervisor: Dan Greiner

External Supervisor: Nikolas Schönfelder,

Daimler Truck AG

Examiner: Prof. Dr.-Ing. Hans-Christian Reuss

Start Date:

Submission Date:

_____

Prof. Dr.-Ing. H.-C. Reuss

Declaration

I, Mahesh Kollati, hereby confirm that I have written the present work, or the parts marked with my name, independently and that I have complied with the relevant provisions in the preparation of the work, in particular on the copyright protection of third-party contributions as well as contributions of "artificial intelligence" (AI), in particular from generative models, and have only used the sources and aids indicated.

As far as my work contains third-party contributions (e.g. images, drawings, text passages, contributions related to "artificial intelligence"), I declare that I have marked these contributions as such (e.g. quotation, reference) and that I have obtained any necessary consents from the authors to use these contributions in my work.

In the event that my work violates the rights of third parties, I declare that I will compensate the University of Stuttgart for any resulting damage and to indemnify the University of Stuttgart against any claims by third parties respectively.

The thesis has not been the subject of any other examination procedure, either in its entirety or in essential parts. Furthermore, it has not yet been published in full or in parts. The electronic copy matches the other specimens.

Stuttgart,

Mahesh Kollati

Contents

Acronyms

1NF	First Normal Form
2NF	Second Normal Form
3NF	Third Normal Form
AUTOSAR	AUTomotive Open System ARchitecture
CPC	Common Powertrain Controller
DBMS	Database Management Systems
ECU	Electronic Control Unit
ER	Entity-Relationship
NoSQL	Not Only SQL
PDD	Parameter Definition Database
PID	Parameter ID
RBAC	Role-Based Access Control
RDBMS	Relational Database Management Systems
SQL	Structured Query Language
UI	User Interface
VCD	Vehicle Configuration Database
VMAP	Variant Management and Parametrization

Symbols

List of Figures

List of Tables

Kurzfassung

Moderne Nutzfahrzeuge basieren auf hochentwickelten elektronischen Steuergeräten (ECUs), die tausende konfigurierbare Parameter zur Steuerung von Antriebsstrangfunktionen benötigen. Der derzeitige Excel-basierte Ansatz für das Parametermanagement in der Automobilsoftwareentwicklung birgt erhebliche Risiken, darunter Dateninkonsistenzen, Versionskonflikte und eingeschränkte Nachvollziehbarkeit [41]. Diese Arbeit adressiert diese Herausforderungen durch die Entwicklung einer umfassenden Datenbankarchitektur für Variantenmanagement und Parametrisierung (VMAP) in PostgreSQL.

Die Forschung implementiert einen phasenbasierten Versionierungsansatz, der an den Entwicklungszyklen der Automobilindustrie ausgerichtet ist und gegenüber traditionellen änderungsbasierten Versionierungsmodellen erhebliche Leistungsvorteile bietet, während die Datenintegrität gewährleistet bleibt. Empirische Tests zeigen, dass dieser Ansatz Abfrageantwortzeiten unter 100ms selbst bei Datensätzen mit über 100.000 Parametern erzielt, bei akzeptabler Speichereffizienz [4]. Das hybride Rollen-Berechtigungszugriffsmodell des Systems kombiniert rollenbasierte Sicherheit mit modulspezifischen Zugriffskontrollen, was durch umfangreiche Benutzervalidierung mit realen Engineering-Workflows verifiziert wurde.

Die Integration mit Unternehmenssystemen ermöglicht synchronisierte Parameterdefinitionen und Fahrzeugkonfigurationsdaten, entsprechend den von Hohpe und Woolf [18] beschriebenen Enterprise-Integration-Patterns. Die Leistungsanalyse ergab, dass die Datenbank interaktive Antwortzeiten für die meisten Operationen auch bei Produktionsdatensätzen mit 830 Varianten und 167.990 Segmenten beibehält, wobei bestimmte komplexe Operationen wie Phasenvergleiche eine suboptimale Skalierung aufweisen, die von weiterer Optimierung profitieren würden.

Das resultierende System bietet eine 6,5- bis 21,8-fache Leistungsverbesserung gegenüber nicht-indexierten Implementierungen und löst kritische Einschränkungen des Excel-basierten Ansatzes, einschließlich gleichzeitigem Mehrbenutzer-Zugriff, automatisierter Validierung und umfassender Änderungsverfolgung. Die Leistungsanalyse zeigt, dass, obwohl Änderungshistoriendatensätze den Speicherbedarf dominieren (60,8

Abstract

Modern commercial vehicles rely on sophisticated Electronic Control Units (ECUs) that require thousands of configurable parameters to manage powertrain functions. The current Excel-based approach to parameter management in automotive software development introduces significant risks including data inconsistency, version conflicts, and limited traceability [41]. This thesis addresses these challenges through the development of a comprehensive database architecture for variant management and parametrization (VMAP) in PostgreSQL.

The research implements a phase-based versioning approach aligned with automotive development cycles, providing significant performance advantages over traditional change-based versioning models while maintaining data integrity. Empirical testing demonstrates that this approach yields query response times below 100ms even with datasets exceeding 100,000 parameters, while maintaining acceptable storage efficiency [4]. The system's hybrid role-permission access control model combines role-based security with module-specific access controls, verified through extensive user management validation with real-world engineering workflows.

Integration with enterprise systems enables synchronized parameter definitions and vehicle configuration data, implementing the enterprise integration patterns described by Hohpe and Woolf [18]. Performance analysis revealed that the database maintains interactive response times for most operations even with production-scale datasets containing 830 variants and 167,990 segments, though certain complex operations like phase comparison exhibit suboptimal scaling that would benefit from further optimization.

The resulting system provides a 6.5x-21.8x performance improvement over non-indexed implementations and resolves critical limitations of the Excel-based approach, including multi-user concurrent access, automated validation, and comprehensive change tracking. Performance analysis demonstrates that while change history records dominate storage requirements (60.8% of database size), the storage-performance tradeoff effectively supports the audit requirements common in regulated automotive development environments [37].

1 Introduction

Modern commercial vehicles represent a quintessential example of cyber-physical systems, where sophisticated software enables precise control over complex mechanical components. The software controlling these vehicles has grown exponentially in complexity over recent decades, evolving from simple engine management to comprehensive control of virtually all vehicle functions. At the core of this evolution is the ECU—a specialized computer that executes software to manage specific vehicle functions [8]. Contemporary commercial vehicles contain dozens of interconnected ECUs working in concert to ensure optimal performance, efficiency, and safety across diverse operating conditions.

1.1 Background and Context

The automotive industry has undergone a profound transformation over the past decades, evolving from predominantly mechanical systems to highly sophisticated mechatronic platforms [29]. This evolution has been particularly pronounced in the commercial vehicle sector, where modern trucks rely on complex networks of ECUs to manage everything from engine performance to safety systems [8]. These systems must adapt to a wide range of operational conditions, regulatory requirements, and market-specific configurations, creating a significant challenge in managing software variability.

At the heart of this variability management lies the Common Powertrain Controller (CPC)—a central ECU managing critical functions related to engine and transmission control. The CPC's operation is governed by thousands of configurable parameters that determine how the powertrain behaves under specific conditions [37]. These parameters influence everything from basic engine timing to sophisticated emission control strategies, making their precise configuration essential for vehicle performance, efficiency, and regulatory compliance.

The parameter management challenge is further complicated by the global nature of modern vehicle development. Commercial vehicles must conform to different emissions regulations, operate in diverse environmental conditions, and meet varying customer expectations across global markets. Consequently, a single vehicle model may require numerous parameter configurations, each tailored to specific combinations of market requirements, hardware configurations, and customer specifications [41].

1.2 Problem Statement

The current approach to parameter management in commercial vehicle development relies predominantly on distributed Excel spreadsheets, a methodology that emerged during a period when parameter counts were manageable and development teams were smaller [41]. However, as software complexity has increased exponentially, this fragmented approach has introduced significant limitations and risks to the development process.

Development teams distributed across different locations must coordinate changes to thousands of parameters, track their versions, and ensure consistency across multiple vehicle platforms. The absence of a centralized version control system makes it exceptionally difficult to track changes effectively and manage releases. This situation becomes particularly critical when dealing with safety-critical parameters that directly influence vehicle performance and regulatory compliance.

The manual nature of current processes, combined with the lack of automated validation mechanisms, introduces substantial risks of data inconsistency, version conflicts, and delayed implementation of critical parameter updates. Parameter changes are not consistently verified against established rules and constraints, potentially leading to incompatible configurations or non-compliant behavior [37].

Integration with critical enterprise systems presents another significant challenge. The current process of synchronizing data with internal database systems involves several manual steps, consuming valuable development resources and introducing potential points of failure in the configuration management workflow. The absence of automated data validation and synchronization mechanisms creates additional risks for data integrity and consistency across these interconnected systems.

Furthermore, the increasing emphasis on rapid development cycles and continuous integration in the automotive industry demands a more sophisticated approach to parameter management [8]. The existing system's limitations become particularly apparent when considering the need for simultaneous development of multiple vehicle variants, each requiring specific parameter configurations for different markets and regulatory environments.

These challenges collectively underscore the urgent need for a modern, database-driven solution that can address the complexities of contemporary automotive software development while providing a scalable foundation for future growth and adaptation.

1.3 Research Objectives

This thesis aims to address the fundamental challenges in automotive parameter management through the development of database architecture for Variant Management and Parametrization (VMAP), a web-based application for powertrain parameter configuration. The research objectives encompass both theoretical foundations and practical implementation considerations, focusing on creating a robust solution that meets the complex demands of modern vehicle development processes.

The primary research objective centers on developing a centralized database architecture that can effectively manage the complexity of powertrain parameters while maintaining data integrity and traceability [42]. This architecture must support sophisticated version control mechanisms that can handle parameter variations across different development stages and vehicle variants. The system should provide comprehensive audit trails and change history, enabling development teams to track modifications and understand the evolution of parameter configurations over time.

A second crucial objective focuses on the implementation of a sophisticated version control system that addresses the unique requirements of parameter management in automotive software development. This system must go beyond traditional source code version control approaches to handle the complex relationships between parameters, their variants, and their applications across different vehicle platforms [37]. The version control mechanism should support parallel development streams while maintaining consistency and preventing conflicts in parameter configurations.

The research also aims to establish a comprehensive role-based access control system that supports the diverse needs of different user groups within the development process. This includes creating specialized interfaces and permissions for Module Developers, Documentation Team members, Administrators, and Read-only Users, each with specific capabilities and restrictions aligned with their responsibilities [32]. The access control system must balance security requirements with the need for efficient collaboration among development teams.

Integration with existing enterprise systems represents another critical objective of this research. TheVMAP system must establish seamless data exchange mechanisms with internal database systems, ensuring consistent information flow while minimizing manual intervention [8]. This integration should support automated validation of parameter changes and provide mechanisms for maintaining data consistency across different systems.

A final key objective involves the development of database interfaces and query optimization strategies that will support the web-based interface implementation. While

the actual User Interface (UI) development falls outside the scope of the thesis, the research will focus on designing efficient database structures, stored procedures, and APIs that enable seamless integration with the planned web interface [29]. This includes developing optimized query patterns for complex operations such as parameter comparison, variant management, and release workflows, while ensuring robust data validation and business rule enforcement.

1.4 Significance of the Study

The significance of this research extends beyond addressing immediate technical challenges in parameter management. By developing a comprehensive database solution for variant management and parametrization, this work contributes to the broader field of automotive software engineering in several important ways.

First, the research advances the understanding of version control in parameter-centric systems, extending traditional concepts of software versioning to accommodate the unique characteristics of automotive parameter configurations. While considerable research has been conducted on code versioning, the versioning of parameter data presents distinct challenges that require specialized approaches [4]. This thesis contributes to closing this gap by developing and evaluating new methods for parameter versioning in complex automotive systems.

Second, the work addresses critical industry needs for improved quality and efficiency in vehicle development. Commercial vehicle manufacturers face increasing pressure to reduce development time while managing growing software complexity and ensuring regulatory compliance across global markets [8]. By providing a more robust and efficient parameter management solution, this research directly contributes to these industry priorities, potentially reducing development costs and improving vehicle quality through more consistent parameter configurations.

Third, the research advances the integration of database technology with domain-specific engineering processes. By developing specialized database structures and functions tailored to the unique requirements of automotive parameter management, this work demonstrates how database technology can be adapted to support complex engineering workflows [14]. This integration perspective is valuable not only for automotive applications but also for other engineering domains facing similar challenges in managing complex, highly variable system configurations.

Finally, the research contributes to the growing field of model-based systems engineering by providing a structured approach to managing the parametric aspects of system

models. As the automotive industry continues to adopt model-based approaches for system development, the management of parameter configurations becomes increasingly critical for maintaining model integrity and traceability [37]. This thesis provides insights and solutions that support this evolution toward more systematic model-based development practices.

1.5 Thesis Structure

The thesis is organized into six chapters that systematically address the research objectives and present a comprehensive solution for automotive parameter management. The structure follows a logical progression from theoretical foundations through practical implementation, ensuring thorough coverage of both academic and industry perspectives.

Following this introduction, Chapter 2 presents a comprehensive review of the theoretical background relevant to automotive electronic control systems and database management approaches. This chapter examines the hierarchical organization of automotive electronic systems, explores database management systems with a focus on relational databases, investigates database design methodologies including Entity-Relationship (ER) modeling and normalization, and evaluates access control models and version control concepts applicable to parameter management.

Chapter 3 presents a comprehensive review of the state of the art in database version control systems and automotive parameter management. This chapter examines existing approaches to software configuration management in the automotive industry, analyzes current database versioning techniques, and evaluates their applicability to parameter management systems. The review encompasses both academic research and industry practices, providing a solid foundation for the proposed solution.

Chapter 4 details the methodology and concept development, beginning with a thorough requirements analysis based on industry needs and academic best practices. This chapter explores the system architecture design, consisting of the database schema, version control mechanisms, and user management frameworks. Particular attention is given to the integration requirements with existing systems and the development of robust validation mechanisms for parameter management.

The implementation strategy and technical design are presented in Chapter 5, which outlines the practical realization of the VMAP system. This chapter describes the development of the database structure, the implementation of version control mechanisms, and the creation of the database interfaces. The chapter also details the integration

approaches with internal database systems, highlighting the technical challenges and solutions developed during the implementation phase.

Chapter 6 focuses on system evaluation and validation, presenting a comprehensive assessment of the VMAP system against the defined research objectives. This chapter includes detailed performance analyses, user acceptance testing results, and comparative evaluations against existing parameter management solutions. The evaluation framework incorporates both quantitative metrics and qualitative assessments to provide a thorough understanding of the system's effectiveness.

The thesis concludes with Chapter 7, which summarizes the research findings and presents recommendations for future development. This chapter reflects on the contributions of the research to both academic knowledge and industry practice, discussing the implications for automotive software development and configuration management. Additionally, it outlines potential areas for future research and system enhancement based on the insights gained during the project.

Throughout these chapters, the research methodology combines theoretical analysis with practical implementation, ensuring that the resulting system meets both academic standards and industry requirements. Special attention is given to database versioning approaches, user role management, and integration strategies with existing systems, addressing the unique challenges of automotive software configuration management.

1.6 Project Plan

The research project follows a structured approach spanning six months from November 2024 to April 2025, organized into three distinct phases: Exposé, Implementation, and Finalization. The comprehensive timeline ensures systematic progression through all research objectives while maintaining academic rigor and quality standards.

1.6.1 Exposé Phase (November - January)

The initial phase focuses on establishing strong theoretical foundations and gathering comprehensive requirements. Literature review constitutes a significant portion of this phase, extending over six weeks to ensure thorough coverage of current database versioning approaches, parameter management systems, and industry practices in automotive applications. This review encompasses analysis of existing version control

systems, examination of industry standards for software configuration management, and evaluation of current parameter management solutions.

Requirements analysis follows the literature review, spanning three weeks to capture detailed system specifications. This phase involves extensive stakeholder consultation to document system requirements, analyze existing Excel-based workflows, define integration requirements with internal database systems, and establish user roles and access control specifications. The Exposé phase concludes with the submission of a comprehensive research proposal at the end of Week 3 in January.

1.6.2 Implementation Phase (December - March)

The implementation phase encompasses four major components, each allocated four weeks for development and refinement. Database design initiates this phase, focusing on developing the schema for parameter management, designing version control mechanisms, creating data models for user management, and planning integration interfaces with existing systems.

System architecture development follows, concentrating on overall system design, version control workflows, user management frameworks, and validation mechanisms. This stage establishes the foundational structure for the entire system while ensuring alignment with identified requirements and industry standards.

Prototype development constitutes the third component, involving implementation of core database functionality, development of version control features, creation of user management interfaces, and construction of system integration components. This stage transforms theoretical designs into practical implementations while maintaining focus on system usability and performance.

The final component of this phase involves comprehensive testing and validation, including database performance testing, validation of version control mechanisms, testing of user management functions, and verification of system integration capabilities. This stage ensures all implemented features meet specified requirements and performance standards.

1.6.3 Finalization Phase (April)

The concluding phase focuses on documentation and thesis preparation over four weeks. The first two weeks are dedicated to comprehensive documentation, including compilation of implementation details and system architecture documentation.

The subsequent two weeks concentrate on thesis writing, involving comprehensive documentation of research findings, inclusion of test results and analysis, preparation of conclusions and recommendations, and thorough content review and refinement. The phase concludes with thesis submission in Week 16 and final project presentation in Week 17.

University of Stuttgart
Institute of Automotive Engineering

Figure 1.1: Gantt Chart of the planned work schedule.

2 Theoretical Background

This chapter establishes the theoretical foundation necessary for understanding the design and implementation of the VMAP database system. It begins with an overview of automotive electronic control systems and parameter management, explaining the fundamental concepts that drive the requirements for the VMAP system. The chapter then explores database management systems, database design methodologies, and access control models relevant to the implementation. Finally, it discusses version control concepts and temporal database management approaches, which are critical for the parameter versioning requirements in automotive software development.

2.1 Automotive Electronic Control Systems

Modern commercial vehicles contain dozens of ECUs that manage various vehicle subsystems. Each ECU is a specialized computing device that controls specific functions through software parameters [37]. Understanding the structure and organization of these systems is essential for designing an effective parameter management solution.

2.1.1 ECU Hierarchy and Parameter Organization

Automotive electronic systems follow a hierarchical organization that structures parameters into logical groupings. Figure 2.1 illustrates this hierarchical structure, showing how parameters are organized in automotive electronic systems. At the top level, ECUs represent distinct hardware components controlling specific vehicle functions such as engine management, transmission control, or brake systems [22]. Within each ECU, modules represent functional software units that implement specific capabilities such as cruise control, adaptive power steering, diagnosis. Each module contains Parameter IDs (PIDs) that group related parameters, and finally, individual parameters define specific configuration values that affect system behavior [38].

As shown in Figure 2.1, parameters can either use their default values or be customized through variants and segments. When a parameter is used without customization, its default value is applied. However, when specific vehicle configurations require different parameter settings, variants are created, and segments define the modified parameter values within these variants. This hierarchical structure is not merely organizational but reflects the actual architecture of automotive electronic systems, where software

Figure 2.1: Hierarchical Organization of Automotive Electronic Systems

components are modularized for maintainability, reusability, and functional separation. The CPC, a central ECU in modern trucks, manages critical powertrain functions through thousands of configurable parameters organized into this hierarchical structure [37].

Parameters themselves have complex characteristics beyond simple values. They can be scalar values, one-dimensional arrays (curves), or multi-dimensional arrays (maps or tables). Each parameter has specific attributes defining its data type, valid range, engineering units, scale factors, and default values [29]. For instance, an engine timing map might be represented as a two-dimensional array where engine speed and load are the independent variables, and ignition timing angle is the dependent variable. This complexity in parameter structure creates specific requirements for the database system designed to manage them.

2.1.2 Parameter Variants and Customization

A fundamental challenge in automotive parameter management is supporting multiple parameter configurations for different vehicle variants, regional requirements, and operating conditions. Rather than maintaining separate complete parameter sets for each configuration, which would lead to significant redundancy, automotive systems

implement a variant mechanism that allows selective overriding of parameter values based on specific conditions [37].

In this approach, each parameter has a default value defined in the baseline configuration. Variants are created to represent specific vehicle configurations or conditions, and segments define modified parameter values within these variants. If no segment exists for a particular parameter in an applicable variant, the default value is used. This approach minimizes redundancy by storing only the modified values rather than complete parameter sets for each configuration [8].

The lower part of Figure 2.1 illustrates this concept, showing how parameters can follow two paths: either using their default values (left branch) or being customized through variants and segments (right branch). Variants are associated with code rules—boolean expressions that determine when a variant applies based on vehicle configuration codes. For example, a variant might apply only to vehicles with a specific engine type and transmission combination, or to vehicles destined for a particular market with unique regulatory requirements. The code rule evaluation process selects the appropriate variants for a specific vehicle configuration during parameter file generation [38].

This variant approach creates specific requirements for the database system, which must efficiently store and retrieve variant definitions and segment values while maintaining the relationships between parameters, variants, and segments. The system must also implement a parameter resolution process that correctly applies variants based on vehicle configuration codes, ensuring that the right parameter values are used for each specific vehicle.

2.1.3 Release and Phase Management

Automotive software development follows a structured release process with well-defined phases representing increasing levels of maturity and stability [8]. For parameter management, this translates into a phase-based development process where parameter configurations evolve through sequential stages before being released for production.

Figure 2.2: Automotive Parameter Release Cycle

Figure 2.2 illustrates the typical release cycle in automotive parameter development, consisting of bi-annual releases (e.g., "24.1" and "24.3" for first and third quarters of 2024), with each release progressing through four sequential phases: Phase1 (Initial definition & Configuration), Phase2 (Initial Testing), Phase3 (Refined Testing), and Phase4 (Ready for Production). The diagram shows the linear progression through these phases, with each phase building upon the work completed in the previous phase. Different ECUs may progress through these phases at different rates, requiring the parameter management system to support concurrent work on multiple phases.

Each phase represents a milestone in the development process with specific activities and quality gates. The Phase1 involves the creation of new parameters and initial configuration. Phase2 and Phase3 involve refinement based on testing feedback, with increasing levels of validation. The Phase4 represents the completed configuration ready for production release [37].

When a phase transitions to the next stage, parameter configurations are copied forward, establishing a new baseline for continued development. Changes made in earlier phases should propagate to later phases unless explicitly overridden, creating a complex versioning requirement for the parameter management system [29]. Additionally, at specific development milestones, phases may be "frozen" to create stable reference points for documentation and testing, requiring the parameter management system to enforce read-only access to frozen phases while still allowing continued development in active phases.

This phase-based release process establishes specific requirements for the database system's versioning model, which must maintain distinct parameter configurations for each phase while supporting phase transitions, change propagation, and selective freezing. The versioning approach must align with this development process rather than implementing a generic temporal model, ensuring that the system supports the actual workflows used in automotive parameter development.

2.2 Database Management Systems

Database Management Systems (DBMS) serve as the foundation for structured information storage and retrieval. They provide mechanisms for storing, organizing, and accessing data while ensuring integrity, security, and concurrent access [14]. For the VMAP system, selecting an appropriate database approach is critical for meeting the complex requirements of automotive parameter management.

2.2.1 Relational Database Management Systems

Relational Database Management Systems (RDBMS) organize data into structured tables composed of rows and columns, based on the relational model proposed by E.F. Codd in 1970 [10]. The relational model establishes a mathematical foundation for representing data as relations (tables) with well-defined operations for data manipulation. This approach has dominated database technology for decades due to its solid theoretical foundation and practical advantages for structured data management.

Primary Key
Column

Product Table

ModelNumber	ProductName	ProductPrice	UnitCost
1111	Chair	$129.99	$65.00
1112	Desk	$599.99	$325.00
1113	Table	$299.99	$100.00
1114	File Cabinet	$159.99	$80.00

one-to-many
relationship

Sales Table

OrderNumber	ModelNumber	OrderQuantity	OrderDate
10-1256	1111	1	01/01/2021
10-1328	1111	3	02/15/2021
10-1524	1114	5	02/20/2021
10-1657	1113	1	03/02/2021

Foreign Key
Column

Figure 2.3: Example of a Relational Schema [39]

Figure 2.3 illustrates a practical example of a relational schema with two related tables. The Product Table contains product information with ModelNumber as the primary key, while the Sales Table captures order information with OrderNumber as its primary key. The relationship between these tables is established through the ModelNumber field in the Sales Table, which serves as a foreign key referencing the Product Table. This

one-to-many relationship indicates that a single product can appear in multiple sales orders. The diagram clearly shows how primary and foreign keys establish relationships between tables, implementing the referential integrity that ensures consistency across related data.

In relational databases, tables adhere to predefined schemas that specify the structure, data types, and constraints applicable to the data. Each table typically includes a primary key that uniquely identifies each row, while foreign keys establish relationships between tables, implementing the referential integrity that ensures consistency across related data [14].

A key strength of relational databases is their adherence to ACID properties (Atomicity, Consistency, Isolation, Durability), which ensure reliable transaction processing. Atomicity guarantees that transactions are treated as indivisible units that either complete entirely or have no effect. Consistency ensures that transactions maintain database integrity by transforming the database from one valid state to another. Isolation prevents interference between concurrent transactions, making them appear as if executed sequentially. Durability ensures that committed transactions persist even after system failures [14].

ACID compliance makes relational databases particularly suitable for automotive parameter management, where data integrity and consistency are paramount. Incorrect parameter values could potentially affect vehicle safety and performance, making the strong consistency guarantees of relational databases essential for maintaining data integrity [37]. Additionally, the hierarchical structure of automotive parameter systems—with well-defined relationships between ECUs, modules, PIDs, and parameters—aligns naturally with the relational model's representation of structured data and relationships.

2.2.2 Non-Relational Database Systems

Non-relational databases, often referred to as Not Only SQL (NoSQL) databases, emerged as alternatives to the relational model, particularly for use cases involving large-scale distributed systems, unstructured data, or schema flexibility requirements. Unlike relational databases, NoSQL systems typically sacrifice some aspects of ACID compliance in favor of scalability, flexibility, and performance characteristics suited to specific application domains [4].

Figure 2.4 illustrates the four major types of NoSQL databases, each with distinct data models and storage approaches. Key-value databases (top-left) store data as simple key-value pairs, offering fast lookups but limited query functionality. Graph

Key-value database

k	v
k	v
k	v
k	v
k	v
k	v
k	v

Graph database

Column family database

Document database
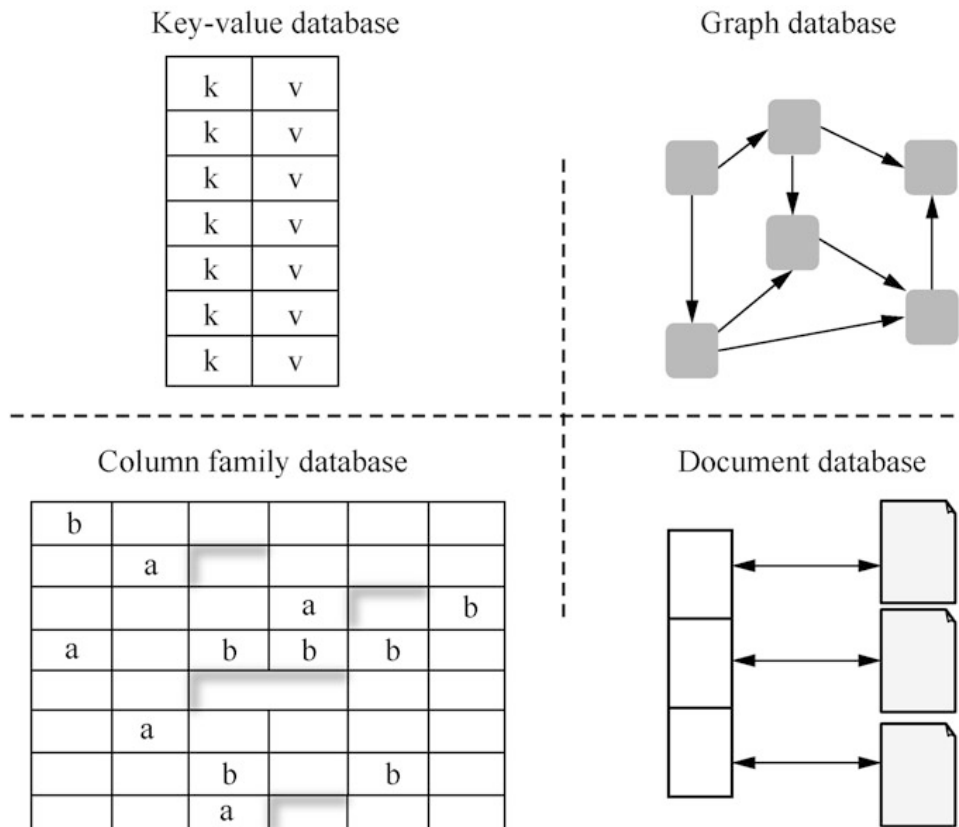
Figure 2.4: Major Types of NoSQL Databases [17]

databases (top-right) organize data as networks of interconnected nodes and relationships, excelling at representing complex relationships. Column-family databases (bottom-left) store data in column-oriented structures that can efficiently handle sparse data. Document databases (bottom-right) store semi-structured data as documents, typically in formats like JSON or XML, providing schema flexibility while maintaining query capabilities.

NoSQL databases can be categorized into these four types based on their data models: document databases (MongoDB, CouchDB), key-value stores (Redis, DynamoDB), column-family stores (Cassandra, HBase), and graph databases (Neo4j, Amazon Neptune). Many NoSQL systems follow the BASE principle (Basically Available, Soft state, Eventually consistent) rather than ACID, prioritizing availability and partition tolerance over immediate consistency [7].

While NoSQL databases excel in specific domains such as high-volume web applications, real-time analytics, and social networks, they present challenges for applications requiring complex transactions, strict data integrity, or sophisticated query capabilities across related entities [23]. For automotive parameter management, these limitations

make NoSQL systems generally less suitable than relational databases.

The potential for eventual consistency rather than immediate consistency in many NoSQL systems could lead to incorrect parameter configurations being used during development or testing, creating significant risks for vehicle performance and safety. Additionally, the hierarchical nature of automotive electronic systems, with well-defined relationships between entities, aligns naturally with the relational model's approach to representing structured data and relationships. The ability to enforce these relationships through foreign key constraints provides important safeguards against data inconsistency that would be more difficult to implement in many NoSQL systems.

2.3 Database Design Methodologies

Database design methodologies provide structured approaches to creating efficient, reliable database systems. These methodologies help translate real-world information needs into technical implementations that can store and manage data effectively. This section explores fundamental approaches that form the theoretical foundation for database design, presented in a sequence that follows the natural progression from user requirements to technical implementation.

2.3.1 Use Case Modeling

Before designing a database structure, it is essential to understand how users will interact with the system. Use case modeling provides a technique for capturing user requirements by identifying who will use the system (actors) and what they need to accomplish (use cases). Developed by Ivar Jacobson, use case modeling has become a cornerstone of requirements analysis in system development [20].

Figure 2.5 illustrates a simple use case diagram for a telephone system. The diagram shows two actors, the "Caller" and "Callee," and two use cases, "Make Call" and "Transfer Call." The lines connecting the Caller to both use cases indicate that the Caller can initiate either a direct call or a call transfer, while the Callee is connected only to the "Make Call" use case, indicating their role as the recipient of calls. The rectangle surrounding the use cases represents the system boundary, clearly delineating the functional scope of the telephone system. This visual representation effectively communicates the basic functionality that the system must support, providing a foundation for further design and implementation decisions.

Figure 2.5: Simple Use Case Diagram for a Telephone System

A use case represents a specific goal that an actor wishes to achieve using the system. Actors can be human users with different roles (such as administrators or regular users) or external systems that interact with the database. The collection of all use cases defines the system's functional boundaries—what it must do to satisfy user needs [20].

Use case diagrams provide a visual representation of these relationships, showing actors as stick figures and use cases as ovals, with lines connecting actors to their associated use cases. This visual format makes the system's purpose accessible to non-technical stakeholders, facilitating communication between developers and users. As noted by Jacobson, "Use cases bridge the gap between the users' and the developers' views of the system" [20].

For automotive parameter management systems, use case modeling helps identify the different ways in which engineers, documentation specialists, administrators, and other stakeholders need to interact with parameter data. These use cases then inform the database design, ensuring that the resulting structure effectively supports all required operations.

2.3.2 Entity-Relationship Modeling

After understanding user requirements through use cases, the next step is to model the data itself. ER modeling provides a conceptual framework for representing the data structure needed to support the identified use cases. Introduced by Peter Chen in 1976, ER modeling has become the most widely used approach for conceptual database design [9].

Figure 2.6 presents a detailed ER diagram for a fitness studio database. The diagram shows six main entities: Members, Member Signups, Class Schedule, Classes, Types,
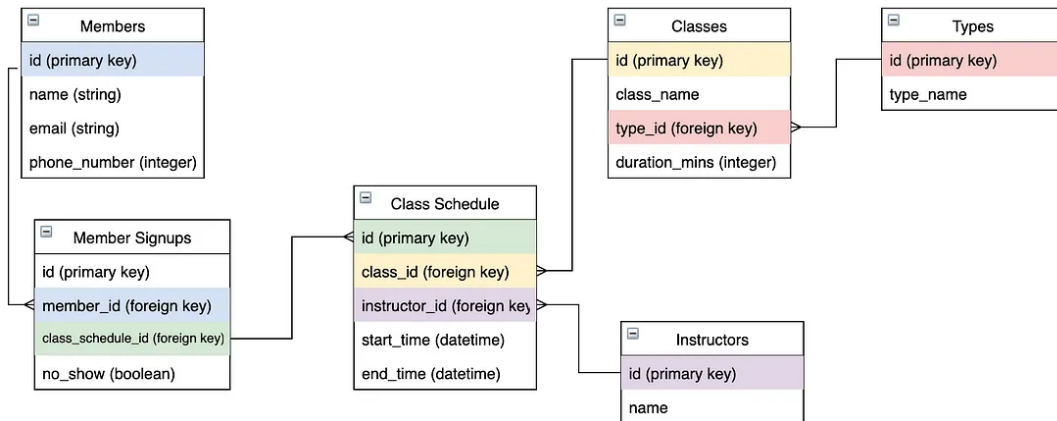
Figure 2.6: Entity-Relationship Diagram for a Fitness Studio Database

and Instructors. Each entity is represented as a rectangle containing its attributes, with primary keys clearly labeled at the top of each entity. The diagram uses color coding to visually distinguish different entity types and their related attributes. Foreign key relationships are explicitly shown with connecting lines and arrows pointing to the referenced primary keys, illustrating the database's referential integrity constraints.

The diagram demonstrates several important relational concepts. The Members entity contains basic member information including a unique identifier, name, email, and phone number. The Class Schedule entity serves as a central hub, connecting classes with their instructors and time slots. The Member Signups entity implements a many-to-many relationship between members and scheduled classes, with an additional attribute to track attendance (no_show). The Classes entity includes a foreign key reference to the Types entity, establishing a categorization system for different class offerings. This comprehensive diagram provides a clear blueprint for database implementation, showing both the structural components and their interrelationships in the fitness studio domain.

ER modeling identifies three main components: Entities represent the objects or concepts about which information needs to be stored. In an automotive context, these might include vehicles, ECUs, parameters, and users. Entities are represented as rectangles in ER diagrams. Attributes describe the specific properties or characteristics of each entity. For example, a parameter entity might have attributes like name, value, unit, and description. Attributes are shown as ovals connected to their entity. Relationships describe the associations between entities. For instance, "ECUs contain parameters" expresses a relationship between ECU and parameter entities. Relationships are shown as diamonds connecting the related entities, with cardinality notations indicating how many instances of each entity can participate in the relationship [14].

ER modeling is particularly valuable for complex domains like automotive systems because it provides a visual representation that stakeholders can understand while being precise enough to guide database implementation. Chen explains that "the ER model adopts the more natural view that the real world consists of entities and relationships" [9], making it an intuitive approach for modeling real-world systems.

2.3.3 Database Normalization

Once the conceptual model is established through ER modeling, database normalization helps refine this model into an efficient, consistent structure. Normalization is a systematic process developed by E.F. Codd that organizes data to minimize redundancy and avoid update anomalies [10].

Normalization proceeds through several "normal forms," each addressing specific types of data inconsistencies:

First Normal Form (1NF) requires that each cell in a table contains only a single value, not a list of values. For example, storing multiple phone numbers in a single field would violate 1NF. This ensures that data is atomic (indivisible) and can be manipulated consistently.

Second Normal Form (2NF) builds on 1NF by requiring that all non-key attributes depend on the entire primary key, not just part of it. This prevents situations where changing one piece of data requires multiple updates in different places.

Third Normal Form (3NF) further refines the structure by requiring that non-key attributes depend only on the primary key, not on other non-key attributes. This eliminates transitive dependencies that can lead to update anomalies [14].

For most practical applications, achieving 3NF provides a good balance between data integrity and system performance. As explained by Date, "Third normal form is considered adequate for most practical purposes; further normalization is usually performed only when necessary" [12].

In automotive parameter management, normalization helps organize complex data about ECUs, modules, and parameters into a structure that maintains consistency while supporting efficient access. For example, normalizing parameter data ensures that when a parameter value changes, that change only needs to be recorded in one place, eliminating the risk of inconsistent values across the database.

2.3.4 Role-Based Access Control Models

Database systems often contain sensitive information that should not be accessible to all users. Role-Based Access Control (RBAC) provides a structured approach to managing permissions within a database system. Introduced by David Ferraiolo and Richard Kuhn in the 1990s, RBAC has become the predominant model for access control in enterprise systems due to its balance of security and administrative simplicity [32].

The core concept of RBAC is that permissions are associated with roles, and users are assigned to appropriate roles rather than being granted permissions directly. A role represents a specific function within an organization, such as "administrator," "engineer," or "analyst." Each role is granted a set of permissions that allow users assigned to that role to perform specific operations on database objects like reading, creating, updating, or deleting records [15].

This structure provides several important advantages over direct permission assignment. First, it simplifies administration by allowing permissions to be managed at the role level rather than the individual user level. When a new user joins the organization, they can simply be assigned to the appropriate roles rather than requiring configuration of individual permissions. Second, it improves security by implementing the principle of least privilege, ensuring that users have only the permissions necessary for their specific responsibilities which reduces the risk of unauthorized access or accidental data modifications [32].

The theoretical foundation of RBAC includes several key components: users (individuals who need access to the system), roles (collections of permissions that correspond to job functions), permissions (defined operations on specific resources), and sessions (temporary bindings between users and their assigned roles). These components provide a flexible framework for implementing access control policies tailored to specific organizational needs while maintaining a clear separation between users and permissions through the role abstraction [31].

2.3.5 Version Control for Databases

Version control for databases addresses the challenge of tracking changes to data structures and content over time. Unlike traditional file-based version control systems designed for source code, database versioning must maintain complex relationships between entities while preserving historical states and supporting evolution through distinct development stages [4].

Several approaches have emerged for implementing version control in database systems. The snapshot approach captures complete database states at specific points in time, providing simple retrieval of historical states but potentially consuming significant storage resources. The change-based approach records only modifications to database content, reducing storage requirements but requiring reconstruction of historical states through the application of change records [4].

The temporal approach extends traditional database structures with time dimensions, enabling direct querying of historical states through time-based predicates. This approach typically introduces valid time (when facts are true in the real world) and transaction time (when facts are recorded in the database) dimensions, allowing sophisticated historical analysis but adding complexity to schema design and query formulation [35].

Phase-based versioning represents a domain-specific approach that aligns database versioning with development phases rather than continuous time. This approach explicitly models development stages as first-class entities in the database schema, associating data with specific phases rather than temporal timestamps. According to Bhattacherjee et al., "Domain-specific versioning approaches often provide better performance and usability than generic temporal database techniques when tailored to specific application requirements" [4].

2.3.6 Temporal Database Concepts

Many applications, including automotive parameter management, need to track how data changes over time. Temporal database concepts address these requirements by providing mechanisms for managing time-varying data. Unlike traditional databases that store only the current state, temporal databases maintain historical states and support queries based on time dimensions [35].

ENo	EStart	EEnd	EDept
22217	2010-01-01	2011-02-03	3
22217	2011-02-03	2011-09-10	4
22217	2011-09-10	2011-11-12	3

Figure 2.7: Temporal Table Tracking Employee Department History

Figure 2.7 illustrates a temporal database table that tracks an employee's department history over time. The table contains four columns: ENo (Employee Number), EStart (employment start date in a department), EEnd (employment end date in a department), and EDept (department identifier). The example shows the complete employment history of employee number 22217, who moved between departments over time. The first row indicates that the employee started in department 3 on January 1, 2010, and remained there until February 3, 2011. The second row shows that on February 3, 2011, the employee transferred to department 4, where they stayed until September 10, 2011. The third row demonstrates that on September 10, 2011, the employee returned to department 3, where they remained until November 12, 2011.

This temporal table implementation enables tracking of the complete history of the employee's department assignments without overwriting previous records. Each row represents a specific period during which the employee belonged to a particular department, with non-overlapping time intervals. The start date of each subsequent record matches the end date of the previous record, ensuring continuity in the historical record. This approach to temporal data storage supports historical queries like "Which department did employee 22217 belong to on March 15, 2011?" by evaluating which row's time interval (EStart to EEnd) contains the query date. Such temporal structures are essential for maintaining comprehensive historical records in domains where tracking changes over time is critical, such as automotive parameter management.

Temporal databases typically support two key time dimensions. Valid time represents when facts are true in the modeled reality—for example, when a particular parameter configuration becomes active in a vehicle. Transaction time represents when facts are recorded in the database—for example, when a parameter value was updated in the system. Databases that support both dimensions are known as bi-temporal databases [24].

Temporal database implementations often use specialized table structures called temporal tables. These tables extend traditional table structures with additional timestamp columns that define the time periods during which each record is valid. For example, a temporal parameter table might include ValidFrom and ValidTo columns that define when each parameter value is applicable, allowing the database to maintain a complete history of parameter changes [30].

Kulkarni and Michels explain that "temporal tables provide a systematic way to track and query historical data without requiring application-level version management" [24]. This capability is particularly valuable in regulated industries like automotive development, where traceability and auditability of parameter changes are essential for compliance and quality assurance.

In automotive parameter management, temporal database concepts can support critical requirements such as tracking parameter evolution throughout the development lifecycle, maintaining historical records for diagnostic and compliance purposes, and enabling historical analysis to understand how parameter configurations have evolved over time. These capabilities form an important theoretical foundation for designing systems that manage time-sensitive data in complex domains.

2.3.7 Strategic Denormalization

While normalization provides a theoretical foundation for database integrity, practical database design often requires balancing normalization principles with performance considerations. Strategic denormalization involves deliberately introducing controlled redundancy to improve performance for specific operations [4].

Consider a fully normalized database where information about parameters, their modules, and their ECUs is stored in separate tables. To retrieve a parameter with its associated module and ECU information would require joining all three tables—an operation that becomes increasingly expensive as the database grows. In cases where this retrieval happens frequently, storing the module and ECU names directly in the parameter table (introducing controlled redundancy) could significantly improve performance [33].

Molinaro emphasizes that "denormalization is not about abandoning normalization principles, but about making strategic exceptions for performance reasons" [25]. These exceptions should be carefully documented and justified based on specific performance requirements.

For automotive systems, where both data integrity and query performance are critical, finding the right balance between normalization and strategic denormalization is essential. This balance ensures that parameter data maintains consistency while providing the performance needed for engineering workflows.

2.3.8 Conceptual, Logical, and Physical Design Levels

Database design typically proceeds through three levels of abstraction, allowing designers to manage complexity by focusing on different aspects at each stage:

Conceptual design focuses on what data needs to be stored, without concern for implementation details. The ER model created at this stage captures entities, attributes, and

relationships from a business perspective, providing a foundation that both technical and non-technical stakeholders can understand [14].

Logical design transforms the conceptual model into structures specific to the chosen database model (typically relational), defining tables, columns, keys, and relationships. This stage applies normalization principles to refine the structure, independent of any specific database system [14].

Physical design addresses how the logical design will be implemented in a specific database management system, considering factors like storage structures, indexing strategies, and access methods. This stage optimizes the design for performance based on anticipated usage patterns [28].

Moving through these levels allows database designers to progressively refine the database structure, addressing different concerns at each stage. As Elmasri and Navathe observe, "The separation of conceptual, logical, and physical design allows database designers to focus on the appropriate level of abstraction at each stage" [14].

For automotive parameter management, this layered approach helps manage the complexity of the domain, ensuring that the resulting database effectively supports both the business requirements (storing and managing parameter configurations) and the technical requirements (performance, scalability, and maintainability).

3 State of the Art

This chapter examines the current state of the art in database version control systems and automotive parameter management. It begins by analyzing existing approaches to software configuration management in the automotive industry, followed by an evaluation of database versioning techniques and their applicability to parameter management systems. The chapter also explores role-based access control models and integration strategies for enterprise systems, establishing the theoretical foundation for theVMAP system design.

3.1 Parameter Management in Automotive Software Development

The complexity of automotive software has grown exponentially in recent decades, with modern vehicles containing up to 100 million lines of code distributed across dozens of ECUs [29]. This growth has significantly increased the importance and complexity of parameter management in automotive development.

3.1.1 Evolution of Automotive Parameter Management

Parameter management in automotive systems has evolved from simple calibration tables to sophisticated configuration frameworks managing thousands of parameters across multiple vehicle variants. Broy [8] describes the fundamental challenges in automotive software engineering, highlighting that software complexity is driven by the need to address multiple variants, market requirements, and technical functions. The parameter configuration problem is specifically identified as one of the key challenges in this domain.

Early approaches to parameter management relied on specialized tools provided by ECU suppliers, which typically stored parameters in proprietary formats with limited version control capabilities. Pretschner et al. [29] note that these tools evolved from simple memory editors to more sophisticated calibration environments, but remained focused on individual ECUs rather than system-wide parameter management.

Staron [37] describes how AUTOSAR (Automotive Open System Architecture) has contributed to more structured parameter management by defining standard interfaces

and component models that separate parameters from implementation. However, the practical implementation of these standards varies across organizations and ECU suppliers, creating integration challenges for comprehensive parameter management.

3.1.2 Challenges in Automotive Parameter Management

The management of parameters in automotive software development presents specific challenges that distinguish it from general software configuration management. Pretschner et al. [29] identify several key challenges related to variability management in automotive software, including the need to maintain multiple parameter configurations for different vehicle variants, markets, and operating conditions.

Broy [8] emphasizes the challenge of managing interdependencies between parameters, noting that changes to one parameter often require coordinated changes to related parameters to maintain system consistency. This creates a need for sophisticated dependency tracking mechanisms that go beyond traditional version control systems.

Another significant challenge relates to validation requirements for parameter changes. Unlike source code, which can be validated through compilation and static analysis, parameters require functional testing to verify their correctness. Pretschner et al. [29] describe how this validation often involves specialized hardware-in-the-loop or vehicle-level testing, creating a significant gap between parameter modification and validation.

Kiencke and Nielsen [22] discuss the specific challenges related to powertrain control parameters, noting the complex interactions between engine control parameters and their effects on vehicle performance, emissions, and fuel economy. These interactions create a need for sophisticated parameter testing and validation processes beyond simple version control.

3.1.3 Current Approaches and Tools

Current parameter management solutions in the automotive industry span a spectrum from general-purpose tools to specialized automotive calibration systems.

AUTomotive Open System ARchitecture (AUTOSAR) is a global partnership of automotive manufacturers, suppliers, and technology companies that defines a standardized software architecture for automotive electronic control units (ECUs). Established in

2003, AUTOSAR aims to create an open industry standard for automotive E/E (electrical/electronic) architecture to manage the growing complexity of automotive software development. Staron [37] discusses how AUTOSAR tools provide standardized interfaces for parameter management in modern automotive systems, but notes that these tools focus primarily on the technical aspects of parameter definition rather than the organizational processes of parameter development and validation through multiple development phases.

Broy [8] identifies the challenges of integrating parameter management into broader software development processes, noting that many organizations maintain separate workflows for software development and parameter calibration. This separation creates coordination challenges, particularly when parameter changes affect multiple software components or require software modifications.

Pretschner et al. [29] discuss how model-based development approaches are increasingly used in automotive development, with parameters linked to model elements to provide traceability and support automated validation. However, they note that the integration between parameter management tools and modeling environments remains incomplete in many organizations.

For database-oriented approaches to parameter management, Bhattacherjee et al. [4] provide a theoretical foundation by examining the principles of dataset versioning. They describe the fundamental trade-offs between storage efficiency and reconstruction performance, which are particularly relevant for systems that must maintain multiple parameter configurations across different development phases.

3.2 Database Version Control Systems

Version control for database content presents distinct challenges compared to traditional source code version control. While source code version control focuses on tracking changes to text files, database version control must address structured data with complex relationships and constraints [4]. This section examines current approaches to database version control and their applicability to automotive parameter management.

3.2.1 Traditional Database Versioning Approaches

Traditional approaches to database versioning fall into several categories, each addressing different aspects of the versioning challenge. Schema evolution tools focus on tracking and managing changes to database structure through migration scripts or schema manipulation languages. Curino et al. [11] describe an approach for automating database schema evolution in information system upgrades, focusing on maintaining data integrity during schema transitions.

Bhattacherjee et al. [4] provide a comprehensive analysis of dataset versioning approaches, identifying a fundamental trade-off between storage and recreation costs. They categorize versioning strategies into several approaches:

1. Version-first approaches maintain complete snapshots of datasets at specific version points, providing simple retrieval of historical states but requiring substantial storage space.

2. Delta-based approaches store only the changes between versions, reducing storage requirements but increasing the computational cost of reconstructing historical states.

3. Hybrid approaches combine elements of both strategies, typically storing periodic full snapshots with incremental deltas between snapshots.

The authors note that the optimal strategy depends on specific usage patterns, particularly the ratio between storage costs and the frequency and complexity of historical data access operations.

Mueller and Müller [26] describe a practical implementation of database versioning between research institutes, highlighting the challenges of maintaining consistency across systems with different update cycles. Their approach uses a combination of schema versioning and data synchronization mechanisms to maintain consistency while supporting independent evolution.

3.2.2 Temporal Database Approaches

Temporal database approaches provide a theoretical foundation for managing time-varying data in database systems. Kulkarni and Michels [24] describe the temporal features introduced in SQL:2011, which formalized support for period data types and temporal tables in the Structured Query Language (SQL) standard. These features enable tracking of both valid time (business time) and transaction time (system time) dimensions, supporting bi-temporal data management.

The valid time dimension represents when facts are true in the modeled reality, independent of when they are recorded in the database. This dimension supports business-oriented temporal queries such as "What was the value of this parameter in a specific phase?" or "When did this parameter change from value A to value B?" [6]. The transaction time dimension represents when facts are recorded in the database, supporting auditability through questions like "Who changed this parameter, and when did they change it?" [24].

Bi-temporal databases combine both dimensions, providing a comprehensive framework for tracking both when changes occurred in the system and when they became effective in the real world [6]. This approach is particularly valuable for regulated industries like automotive development, where both historical accuracy and change auditability are essential for compliance and quality assurance.

Snodgrass [35] provides a comprehensive guide to developing time-oriented database applications in SQL, describing practical techniques for implementing temporal functionality in relational database systems. The author presents various approaches to tracking historical data, including transaction-time tables, valid-time tables, and bi-temporal tables, with practical implementation guidance for each approach.

Biriukov [5] examines practical implementation aspects of bi-temporal databases, highlighting the challenges of schema design, query formulation, and performance optimization. The author notes that domain-specific temporal approaches often provide more practical solutions than generic bi-temporal frameworks, particularly for applications with specialized temporal requirements.

3.2.3 Version Control for Parameter Management

Version control for automotive parameter management presents specific requirements that differ from general database versioning needs. Drawing from the literature, several key requirements can be identified:

Broy [8] discusses the need for version control approaches that align with automotive development processes, which typically follow a structured progression through predefined development phases. Unlike source code versioning, which often follows continuous development with arbitrary version points, parameter versioning must support specific phase-based workflows.

Bhattacherjee et al. [4] examine the trade-offs between different versioning strategies, which are particularly relevant for parameter management systems that must maintain multiple configurations across different development phases. The authors' analysis

of storage versus reconstruction costs provides a theoretical foundation for designing efficient parameter versioning systems.

Snodgrass [35] describes techniques for tracking valid-time information in database systems, which aligns with the need to maintain parameter configurations that are valid for specific development phases or vehicle configurations. However, the author's focus on general temporal database approaches does not address the specific requirements of phase-based development.

Bhattacherjee et al. [4] note that domain-specific versioning systems often provide more effective solutions than generic versioning frameworks, particularly for domains with structured development processes and complex entity relationships. This observation supports the development of specialized versioning approaches tailored to automotive parameter management rather than adopting generic temporal database techniques.

3.3 Role-Based Access Control in Enterprise Systems

RBAC has become a dominant paradigm for managing access rights in enterprise systems, providing a structured approach to security management that aligns with organizational responsibilities [32]. For automotive parameter management, where different user roles have distinct responsibilities and access requirements, RBAC provides a foundation for implementing appropriate security controls.

3.3.1 RBAC Model and Extensions

The core RBAC model, as defined by Sandhu et al. [32], consists of users, roles, permissions, and sessions. Users are assigned to roles that correspond to job functions, and roles are granted permissions that authorize specific operations on protected resources. This indirect association between users and permissions through roles simplifies security administration while maintaining the principle of least privilege.

Several extensions to the basic RBAC model have been developed to address more complex security requirements. Sandhu et al. [32] describe hierarchical RBAC, which introduces role hierarchies that enable permission inheritance between roles, supporting organizational structures with senior roles inheriting permissions from junior roles.

Sandhu and Bhamidipati [31] present administrative RBAC (ARBAC), which addresses the management of the RBAC system itself, defining who can assign users to roles and

modify role permissions. This extension is particularly relevant for enterprise systems where role and permission management is distributed across different administrative domains.

Ferraiolo et al. [15] describe policy-enhanced RBAC, which combines role-based permissions with attribute-based policies to provide context-sensitive access control. This hybrid approach is particularly valuable for systems where access decisions depend on both user roles and context-specific factors such as time, location, or resource attributes.

3.3.2 RBAC in Database Systems

Modern database management systems provide varying levels of support for RBAC principles. Elmasri and Navathe [14] describe the evolution of database security mechanisms from simple user-based privileges to more sophisticated role-based models. Most enterprise database systems now include native support for roles, user-role assignments, and permission management through SQL statements like GRANT and REVOKE.

Obe and Hsu [28] detail PostgreSQL's implementation of RBAC concepts, including role hierarchies through role inheritance, permission management through fine-grained privileges, and row-level security policies for content-based access control. These capabilities provide a foundation for implementing domain-specific access control models on top of the database system's native security features.

However, database-level RBAC implementations typically focus on controlling access to database objects like tables, views, and functions, rather than providing application-level access control that considers domain-specific entities and operations. For complex applications like automotive parameter management, database-level RBAC must be complemented with application-level access control logic that maps domain-specific concepts to database operations [15].

3.3.3 Access Control for Parameter Management

Access control for automotive parameter management presents specific requirements that extend beyond basic RBAC models. Drawing from the literature, several key access control requirements can be identified:

Sandhu et al. [32] provide the theoretical foundation for role-based access control, which aligns with the organizational structure of automotive development teams. Different roles such as parameter engineers, module developers, and system integrators require different access rights to parameter data.

Ferraiolo et al. [15] describe policy-enhanced RBAC, which combines role-based permissions with attribute-based policies. This hybrid approach is particularly relevant for parameter management, where access rights may depend on both user roles and attributes of the parameters being accessed, such as their development phase or module assignment.

Hu et al. [19] discuss practical aspects of implementing and managing policy rules in attribute-based access control, providing insights into the challenges of combining role-based and attribute-based approaches. Their work highlights the importance of balancing security requirements with usability considerations, which is particularly relevant for parameter management systems used by diverse stakeholder groups.

Sandhu and Bhamidipati [31] address the administrative aspects of RBAC, which are important for parameter management systems where access control administration may be distributed across different organizational units. Their ARBAC97 model provides a framework for delegating administrative responsibilities while maintaining central governance.

3.4 Database Integration with Enterprise Systems

Integration between database systems and enterprise applications presents significant challenges in automotive development environments, where parameter management must interact with numerous other systems across the development lifecycle. Effective integration strategies must address both technical interoperability and semantic consistency while maintaining performance and security [18].

3.4.1 Enterprise Integration Patterns

Enterprise integration patterns, as documented by Hohpe and Woolf [18], provide a catalog of solutions for common integration challenges. These patterns address various aspects of system integration, including messaging styles, messaging channels, message construction, and message transformation.

For database-centric applications like parameter management systems, several integration patterns are particularly relevant. Fowler [16] describes the Repository pattern, which provides a structured approach to data access, abstracting the database implementation details behind a domain-focused interface. This abstraction simplifies integration by providing a stable API for other systems to interact with the parameter repository.

Fowler [16] also documents the Data Transfer Object (DTO) pattern, which addresses the challenge of transferring data between systems with different data models. By defining specialized objects for inter-system communication, this pattern enables consistent data exchange while isolating each system's internal representation.

The Canonical Data Model pattern, as described by Hohpe and Woolf [18], establishes a common data representation across multiple systems, simplifying data transformation and ensuring consistent interpretation. This pattern is particularly valuable for parameter management, where the same parameter concepts may be represented differently in various systems across the development lifecycle.

3.4.2 Database Synchronization Approaches

Database synchronization presents specific challenges when integrating parameter management systems with other enterprise data sources. Mueller and Müller [26] describe approaches to database versioning and synchronization between research institutes, highlighting the challenges of maintaining consistency across systems with different update cycles.

Bhattacherjee et al. [4] discuss the principles of dataset versioning, which are relevant for synchronization between parameter management systems and other enterprise databases. Their analysis of the trade-offs between storage and recreation costs provides insights into designing efficient synchronization mechanisms that minimize both data transfer volumes and processing overhead.

Seenivasan and Vaithianathan [34] examine change data capture (CDC) techniques, which enable incremental synchronization by identifying and propagating only changed data between systems. These techniques reduce synchronization overhead compared to full dataset transfers but require reliable change detection mechanisms and careful handling of interdependent changes.

Kleppmann and Beresford [23] address the challenges of conflict resolution in distributed data systems, which are relevant for parameter management systems that must synchronize with multiple enterprise data sources. Their work on conflict-free

replicated data types provides theoretical foundations for designing synchronization mechanisms that maintain consistency across distributed systems.

3.5 Summary and Research Gaps

The review of existing literature reveals several research gaps in the domain of database systems for automotive parameter management:

Current database versioning approaches, as described by Bhattacherjee et al. [4] and Snodgrass [35], provide general frameworks for managing time-varying data but do not specifically address the phase-based development processes common in automotive parameter management. There is a need for specialized versioning approaches that align directly with automotive development workflows while providing the traceability and auditability required for regulatory compliance.

The RBAC models described by Sandhu et al. [32] and Ferraiolo et al. [15] provide a foundation for access control but require extensions to address the specific requirements of parameter management, where access rights depend on both organizational roles and parameter-specific attributes such as module assignment and development phase.

Integration approaches documented by Hohpe and Woolf [18] and Mueller and Müller [26] provide general patterns for system integration but do not specifically address the challenges of integrating parameter management systems with automotive-specific enterprise systems such as parameter definition databases and vehicle configuration databases.

These research gaps highlight the need for domain-specific solutions that combine insights from database version control, access control models, and enterprise integration patterns with specialized knowledge of automotive development processes. TheVMAP system addresses these gaps by developing a database architecture tailored to the specific requirements of automotive parameter management, as will be detailed in subsequent chapters.

4 Methodology and Concept Development

This chapter presents the systematic approach taken in designing the VMAP system. The methodology follows established software engineering principles to address the complex requirements of automotive parameter management. Beginning with a requirements analysis, the chapter proceeds to detail the conceptual architecture design, data model, validation mechanisms, and integration approaches developed to ensure system robustness and compatibility with existing enterprise infrastructure.

4.1 Requirements Analysis

The foundation of theVMAP system design was a comprehensive requirements analysis conducted through a series of structured interviews with stakeholders, detailed examination of the existing Excel-based process, and workshops with domain experts. This multi-faceted approach, following Sommerville's framework for requirements engineering, ensured that both functional and non-functional requirements would be thoroughly identified and prioritized [36].

4.1.1 Functional Requirements

The primary functional requirements were derived from direct observation of engineers' current Excel-based workflow combined with semi-structured interviews conducted with module developers and documentation specialists. Through this process, several critical requirements emerged for the VMAP system.

The system must support the hierarchical organization of parameters within ECUs, Modules, and PIDs, mirroring the domain-specific structure of automotive electronic systems as described by Staron [37]. This hierarchical organization is essential for maintaining the logical structure of vehicle parameters and aligning with established engineering practices.

Users must be able to create variants for parameters with specific code rules determining their applicability, and define segments representing modified parameter values. If no segment exists, the system must default to Parameter Definition Database values—an approach that allows efficient storage by tracking only modifications rather than duplicating unchanged parameters, aligning with Bhattacherjee's principles of dataset versioning [4].

The system must track parameter values across four distinct release phases: Phase1, Phase2, Phase3, and Phase4, with changes in earlier phases propagating to later phases unless explicitly overridden. This phase-based approach represents a domain-specific adaptation particularly suited to automotive software development cycles as identified in Broy's research on automotive software engineering challenges [8].

All modifications require comprehensive logging with user information, timestamp, and detailed change data, supporting regulatory compliance and enabling parameter evolution tracking. Through the stakeholder interviews, it was determined that the system must also provide functionality to create parameter configuration snapshots at specific points, particularly at phase transitions, for documentation purposes—a capability identified as essential for quality assurance and regulatory compliance in automotive software development by Staron [37].

4.1.2 Integration with External Systems

The stakeholder interviews and process analysis revealed that VMAP must integrate with two critical external enterprise systems: the Parameter Definition Database (PDD) and the Vehicle Configuration Database (VCD).

The PDD serves as the authoritative source for the hierarchical structure of automotive electronic systems, containing definitions of ECUs, Modules, PIDs, and baseline parameter configurations. As noted by Pretschner et al. [29], maintaining this hierarchical structure is essential for automotive software development. The workshops with domain experts confirmed that while VMAP will manage parameter variants and customizations, it must rely on PDD for the underlying parameter definitions and structural relationships, requiring a robust synchronization mechanism to maintain consistency between the systems.

The VCD contains comprehensive vehicle specifications and configuration codes that determine which parameter variants apply to specific vehicle configurations. Integration with this system is necessary for two critical functions: validating the boolean code rules associated with parameter variants to ensure they reference valid vehicle codes, and supporting parameter file generation for specific vehicle configurations by resolving the applicable parameter variants based on vehicle codes. This integration requirement aligns with Staron's analysis of automotive software architectures, which emphasizes the importance of configuration management in supporting variant-rich vehicle platforms [37].

These integration requirements necessitated careful consideration of data synchronization approaches, leading to an exploration of different strategies for maintaining

consistency between VMAP and these external systems while minimizing performance impact and complexity.

4.1.3 User Role Requirements

A systematic analysis of the current Excel-based workflow, coupled with contextual inquiries with engineering teams, identified four distinct user roles with specific access requirements. This analysis included shadowing users in their daily work, documenting their tasks and access patterns, and conducting structured interviews to validate the observed patterns.

Module developers require write access to parameters within their assigned modules, with the ability to create and modify variants and segments. Documentation specialists need access to frozen data for documentation, comparison capabilities between phases, and comprehensive change history access. System administrators require comprehensive control over user management, release phases, and special operations like variant deletion and phase freezing. Read-only users need view access to all parameter data with parameter file generation capabilities but no modification rights.

These roles were defined based on the principle of least privilege as described by Sandhu [32], ensuring users have access only to functionality required for their specific responsibilities. This enhances system security while simplifying the user experience by presenting only relevant options.

4.1.4 Data Management Requirements

The system must maintain distinct parameter versions across different release phases, allowing simultaneous work on multiple phases while enabling access to parameter values from any point in the development lifecycle. As highlighted by Elmasri and Navathe [14], data integrity requires maintaining referential integrity across all related entities, particularly ensuring variants and segments associate with valid parameters.

Multi-dimensional parameter support is essential for complex automotive parameters such as mapping tables. Operations modifying multiple related entities must function as atomic transactions to maintain data consistency—particularly important for phase transitions where numerous parameters, variants, and segments may change simultaneously, a requirement that aligns with Bhattacherjee's research on dataset versioning approaches [4].

Query performance analysis, based on projected usage patterns from the current Excel-based process, identified critical query paths including parameter retrieval by ECU, module, PID, release phase, and parameter name. These requirements influenced schema design decisions regarding normalization and indexing strategies to optimize common query patterns.

4.2 Use Case Modeling

Following the requirements gathering process, use case modeling was employed to formalize the system's functional requirements from a user perspective. This approach, as described by Jacobson [20], provides a structured way to represent the system's capabilities and the interactions between users and the system.

The use case diagram in Figure 4.1 illustrates the primary actors and their interactions with theVMAP system. Four primary actor types are identified, corresponding to the user roles established during requirements analysis: Module Developers, who create and modify parameter variants; Documentation Team members, who access parameter data for documentation purposes; Administrators, who manage system settings and user access; and Read-Only Users, who view parameter data without making modifications.

The diagram demonstrates how these actors interact with key system functionalities. Module Developers primarily interact with variant creation and modification use cases, while having limited access to parameter viewing and comparison features. Documentation Team members focus on viewing frozen configurations, comparing parameters across phases, and accessing change history. Administrators have access to all system functions, including user management, release configuration, and system settings. Read-Only Users are limited to viewing parameters and generating parameter files.

This use case model provides a clear visual representation of the system's scope and functionality, serving as a bridge between user requirements and technical implementation. By mapping user roles to specific system functions, the model ensures that the database design will support all required user interactions while maintaining appropriate access controls.
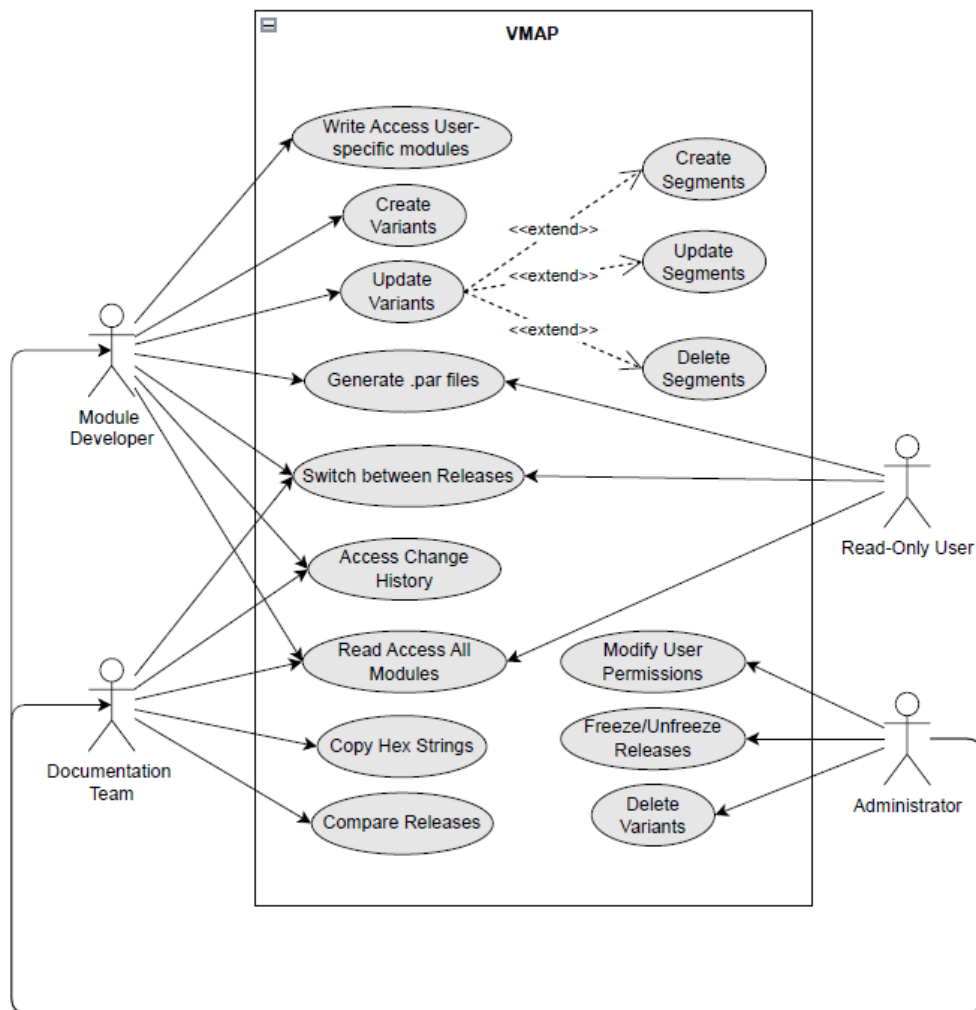
Figure 4.1: VMAP System Use Case Diagram

4.3 User Management Approaches

Based on the identified user role requirements, two distinct approaches to user management were considered for theVMAP system: a traditional role-based approach and a hybrid role-permission approach. Each approach offers different advantages in terms of flexibility, administrative complexity, and alignment with organizational needs.

Figure 4.2: Traditional Role-Based Access Control Approach

4.3.1 Traditional Role-Based Approach

The traditional role-based approach, illustrated in Figure 4.2, assigns users to prede-fined roles that contain fixed sets of permissions, following the classic RBAC model described by Sandhu [32]. In this approach, each user is assigned one or more roles (Administrator, Module Developer, Documentation Team, Read-Only User), and all permissions are granted through these role assignments without individual permission adjustments.

This approach offers administrative simplicity, as user management involves only assigning appropriate roles rather than configuring individual permissions. The role structure also provides clear organizational alignment, with roles directly corresponding to job functions within the development process. From an implementation perspective, this approach simplifies permission checking, typically requiring only verification of role membership rather than individual permission verification.

A key limitation of this approach is its reduced flexibility for accommodating exceptions or specialized access requirements. If a user requires a subset of permissions that doesn't align with existing roles, administrators must either create a new role specifically for that user or grant a role with more permissions than strictly necessary, potentially compromising the principle of least privilege [32].

4.3.2 Hybrid Role-Permission Approach

Figure 4.3: Hybrid Role-Permission Access Control Approach

The hybrid approach, illustrated in Figure 4.3, combines role-based permissions with direct permission assignments, similar to the model described by Ferraiolo et al. [15]. In this approach, users are assigned to primary roles defining their core permissions, but additional permissions can be granted on a per-user basis to address exceptional cases or specialized responsibilities.

This approach offers greater flexibility for accommodating exceptions without creating specialized roles, essential in environments where organizational structures evolve over time. It provides more granular permission control, allowing precise tailoring of access rights to individual responsibilities. However, this flexibility comes at the cost of increased administrative complexity, as both roles and individual permissions must be managed.

The hybrid approach is particularly valuable in the automotive parameter management context, where development responsibilities can vary between projects and temporary access adjustments may be needed for specific tasks or during transition periods. The approach balances structured role assignments with the flexibility to accommodate evolving access requirements, a common scenario in complex engineering environments like automotive development.

4.4 Parameter Synchronization Approaches

Integration with the PDD represents a critical aspect of theVMAP system, requiring careful consideration of synchronization approaches. Two different conceptual approaches were explored for maintaining parameter data across the release phases: the change-based approach and the phase-based approach.

4.4.1 Change-Based Synchronization Approach

Figure 4.4: Change-Based Parameter Synchronization Approach

The change-based approach, illustrated in Figure 4.4, maintains parameter values by recording changes between phases rather than storing complete parameter sets for each phase. In this model, parameters are initially created before the first phase (Phase1), and subsequent modifications are recorded as change entries associated with specific phases.

When a parameter changes in a later phase (like Phase2), only the specific change is recorded rather than creating a new complete copy of the parameter. As shown in the figure, Parameter1 is created before Phase1 with version V1, then modified in Phase2 (creating version V2), but remains unchanged in Phase3 and Phase4. Similarly, Parameter3 changes in Phase2 and Phase3, while Parameter5 changes in every phase except Phase3.

This approach is conceptually aligned with traditional version control systems as described by Bhattacherjee et al. [4], where efficiency is achieved by storing only the differences between versions rather than complete copies. The approach potentially offers storage efficiency advantages by minimizing data duplication across phases, which could be significant for parameter sets containing thousands of entries.

However, this approach introduces conceptual complexity for retrieving parameter values in a specific phase. To determine a parameter's value for a given phase, the system must identify the most recent version of that parameter up to and including the target phase. This reconstruction process introduces additional processing steps compared to direct parameter retrieval.

4.4.2 Phase-Based Synchronization Approach

Figure 4.5: Phase-Based Parameter Synchronization Approach

The phase-based approach, illustrated in Figure 4.5, maintains complete parameter sets for each phase independently. When parameters are initially created before Phase1, each parameter has its specific version (V1). When transitioning to a new phase, all parameters are copied forward, even if they haven't changed. If a parameter is subsequently modified in the new phase, it receives a new version specific to that phase.

As shown in the figure, each parameter exists in every phase with a phase-specific version, regardless of whether the parameter value actually changed between phases.

This creates a clear separation between phases, with each phase maintaining its complete parameter configuration independently.

This approach aligns more directly with the phase-oriented structure of automotive development described by Broy [8], where distinct development milestones form the primary organizational principle. The approach simplifies conceptual understanding and parameter retrieval, as parameters for a specific phase can be accessed directly without reconstructing their values from change history.

The phase-based approach also simplifies phase inheritance by copying parameter configurations forward during phase transitions, allowing subsequent modifications in each phase without affecting previous phases. This copying mechanism preserves the integrity of phase data while supporting the automotive development process, where configurations stabilize progressively through successive phases.

The primary consideration with this approach is the increased storage requirements, as parameters are duplicated across phases even when they haven't changed. However, this trade-off potentially provides benefits in terms of conceptual clarity, query simplicity, and alignment with the automotive development workflow.

4.5 Database System Considerations

Selecting an appropriate database management system for VMAP required consideration of different options against the specific requirements of automotive parameter management. Four major relational database systems were considered as potential platforms for theVMAP implementation: PostgreSQL, Oracle, Microsoft SQL Server, and MySQL.

4.5.1 Database System Requirements

The requirements analysis identified several critical database capabilities needed for effective parameter management:

1. Support for complex data types, including arrays for multi-dimensional parameters and structured types for variant definitions.

2. Robust transaction support for maintaining data consistency during operations affecting multiple related entities.

3. Advanced indexing capabilities to optimize the performance of common query patterns, particularly parameter retrieval across different dimensions.

4. Extensibility for implementing domain-specific operations and validation rules.

5. Comprehensive access control mechanisms supporting the role-based security model.

6. Efficient storage and retrieval of historical data for audit and traceability purposes.

These requirements guided the evaluation of different database systems, focusing on their respective strengths and limitations in addressing the specific needs of automotive parameter management.

4.5.2 Comparative Analysis of Database Systems

Table 4.1 presents a comparative analysis of the four database systems considered for theVMAP implementation, evaluating each against criteria relevant to automotive parameter management.

The comparative analysis revealed different strengths among the database systems. PostgreSQL offers excellent support for complex data types and extensibility, particularly valuable for representing multi-dimensional parameters and implementing domain-specific operations. Oracle provides robust enterprise features with sophisticated optimization capabilities but introduces licensing complexity. SQL Server offers strong integration with Microsoft technologies, while MySQL provides simplicity but has limitations for complex data management requirements.

This analysis provides a foundation for database system selection, considering both technical capabilities and practical factors such as licensing and total cost of ownership. The implementation chapter will detail the specific database system selected and how its capabilities are leveraged in the VMAP implementation.

4.6 Entity-Relationship Model

Based on the requirements analysis and architectural considerations, a comprehensive ER model was developed to capture the complex relationships between system entities. This model follows the approach described by Chen [9], providing a conceptual foundation for the database implementation.

Table 4.1: Comparison of Database Systems for Automotive Parameter Management

Feature	PostgreSQL	Oracle	SQL Server	MySQL
Complex Data Types	Excellent support for arrays, JSON, custom types [28]	Good support, additional licensing for advanced features [1]	Limited built-in support, extensions required [3]	Limited support, improved in recent versions [42]
Transaction Support	Comprehensive with serializable isolation [28]	Excellent with advanced options [1]	Robust support with multiple isolation levels [3]	Limited in some storage engines [33]
Indexing Capabilities	Diverse index types including GIN for text search [28]	Advanced indexing with optimizer hints [1]	Solid capabilities with columnstore indexes [3]	Basic indexing with some limitations [33]
Extensibility	Highly extensible with custom types and functions [28]	Extensible with proprietary mechanisms [1]	Extensible through .NET integration [3]	Limited extensibility [42]
Access Control	Fine-grained with role-based mechanisms [28]	Comprehensive with advanced security features [1]	Strong integration with Active Directory [3]	Basic capabilities with plugin architecture [42]
Licensing	Open source, PostgreSQL License [28]	Commercial, complex licensing model [1]	Commercial with edition-based pricing [3]	Dual licensing: GPL and commercial [42]

Figure 4.6: Comprehensive Entity-Relationship Diagram for theVMAP System

Figure 4.6 illustrates the complete entity-relationship model for theVMAP system, showing the logical organization of entities into functional groups and the relationships between them. The diagram captures the hierarchical structure of automotive parameter data, the phase-based versioning approach, the role-permission access control model, and the mechanisms for external system integration and change tracking.

4.6.1 Core Data Entities

The ER model includes several categories of entities representing different aspects of the system:

User management entities include Users, Roles, Permissions, and their relationships, implementing a role-permission model for access control. These entities support the authentication and authorization requirements identified during the requirements analysis.

Release management entities encompass Releases, Release Phases, and ECU Phase mappings, providing the foundation for the phase-based version control approach. These entities support the automotive development cycle with its distinct phases and milestones.

Parameter structure entities include ECUs, Modules, PIDs, Parameters, and Parameter Dimensions, representing the hierarchical organization of vehicle electronic systems as

described by Staron [37]. These entities capture both the structure and characteristics of parameters, including multi-dimensional parameters like maps and tables.

Variant management entities comprise Variants, Segments, and their relationships to parameters, implementing the core parameter customization functionality. These entities support the creation of parameter variations for different vehicle configurations and operating conditions.

Documentation entities include Documentation Snapshots, Snapshot Variants, and Snapshot Segments, supporting the preservation of historical parameter states for documentation and regulatory compliance. These entities enable the creation of complete parameter configuration snapshots at specific development milestones.

Integration entities consist of Synchronization Records, Vehicle Configurations, and Parameter File Records, supporting connectivity with the Parameter Definition Database and Vehicle Configuration Database. These entities track synchronization operations and store the necessary data for variant resolution and parameter file generation.

Audit entities encompass Change History, Transaction Records, and Phase Copy History, providing comprehensive traceability for all significant operations within the system. These entities support both regulatory compliance and diagnostic capabilities for investigating parameter evolution.

4.6.2 Relationship Structure

The relationships between entities in the ER model reflect the complex interactions between different aspects of automotive parameter management. Key relationships include:

Hierarchical relationships between ECUs, Modules, PIDs, and Parameters, representing the structural organization of automotive electronic systems. These relationships enforce the domain-specific hierarchy while supporting navigation from higher-level entities to their components.

Many-to-many relationships between parameters and phases, implemented through direct association rather than temporal versioning. This structure supports the phase-based versioning approach, allowing efficient retrieval of parameters for specific phases.

Complex relationships between variants, parameters, and segments, capturing the parameter customization process. These relationships ensure that segments are

associated with valid variants and parameters while supporting efficient resolution of effective parameter values based on vehicle configuration.

Temporal relationships for audit and history entities, capturing the evolution of parameter configurations over time. These relationships support both regulatory compliance and diagnostic capabilities, allowing reconstruction of parameter states at specific points in time.

4.6.3 Normalization and Optimization

The ER model was developed using data normalization principles to minimize redundancy while maintaining data integrity, following the approach described by Codd [10]. The model generally adheres to 3NF, ensuring that non-key attributes depend on the primary key rather than on other non-key attributes.

Strategic denormalization was considered in specific areas to optimize performance for common operations, following the principles described by Molinaro [25]. For example, certain frequently accessed attributes from parent entities might be duplicated in child entities to reduce join operations in common queries, providing performance benefits that outweigh the controlled redundancy.

4.7 Validation Mechanisms

To ensure data integrity and consistency, multiple validation mechanism layers were conceptualized for the VMAP system, from basic constraints to sophisticated business rule validation. These mechanisms work together to maintain parameter data quality and reliability throughout the system lifecycle.

4.7.1 Data Integrity Constraints

Database-level constraints were identified as the foundation for enforcing basic integrity rules, following the principles described by Elmasri and Navathe [14]. These constraints include primary key constraints ensuring unique entity identifiers, foreign key constraints maintaining referential integrity between related entities, not-null constraints ensuring required fields contain values, unique constraints preventing duplicate values in specified columns, and check constraints enforcing domain-specific rules such as valid date ranges and parameter value ranges.

These constraints are designed into the database schema as integral parts of entity definitions, ensuring consistent enforcement throughout the system regardless of access path. By implementing constraints at the database level rather than in application code, VMAP ensures that all data modifications adhere to fundamental integrity rules regardless of the source of those modifications.

4.7.2 Business Rule Validation

Domain-specific business rules were conceptualized to be implemented through database triggers and stored procedures, providing a second layer of validation beyond basic constraints. These rules include parameter range validation automatically checking modified values against defined minimum and maximum bounds, phase status validation preventing modifications to frozen phases, segment validation ensuring segments reference valid parameters and variants, and user access validation ensuring users can only modify parameters, variants, and segments for modules to which they have been granted access.

4.7.3 Audit and Traceability Mechanisms

Comprehensive audit and traceability mechanisms were identified as essential for regulatory compliance and quality assurance in automotive parameter management. The core of this capability is the change history tracking mechanism, which automatically captures both before and after states for entity modifications. For each change, the system records the entity being modified, the type of change, the user making the change, the timestamp, and detailed before/after values.

To optimize performance, selective filtering of change data was conceptualized, excluding non-essential fields such as timestamps and large binary data. Additionally, asynchronous audit recording for bulk operations was considered, to reduce the performance impact on high-volume operations while ensuring that all changes are eventually recorded.

Beyond change tracking, specialized audit mechanisms were conceptualized for specific scenarios: phase transition logging to record all phase propagation operations, freeze operation logging to record phase freeze and unfreeze operations, user access logging to capture authentication and authorization events, and integration operation logging to record all external system interactions.

5 Implementation

This chapter presents the technical implementation of the VMAP system conceptual architecture described in Chapter 4. The discussion focuses on key architectural components and technical decisions that enable efficient parameter versioning in automotive software development, with concrete implementation details demonstrating how theoretical concepts translate to practical solutions.

5.1 Database Structure Implementation

Following the comparative analysis of database systems presented in Chapter 4, PostgreSQL was selected as the implementation platform due to its superior support for complex data types, extensibility features, and advanced indexing capabilities [28]. The database implementation transforms abstract entities and relationships into concrete database structures, implementing the hierarchical organization of automotive electronic systems through physical tables and relationships.

5.1.1 Core Data Entities

The hierarchical structure of automotive electronic systems is implemented through four primary entity types: ECUs, Modules, PIDs, and Parameters. The ECU and Module entities form the top levels of the hierarchy, with a many-to-many relationship reflecting the reality that modules can exist across multiple ECUs. This structure aligns with the domain model described by Staron [37], where logical groupings of software functions must be maintained across different hardware configurations.

```
1  CREATE TABLE ecus (
2      ecu_id INTEGER PRIMARY KEY,
3      name VARCHAR(100) NOT NULL,
4      description TEXT,
5      byte_order VARCHAR(50),
6      created_at TIMESTAMP WITH TIME ZONE DEFAULT
         ↳ CURRENT_TIMESTAMP,
7      external_id INTEGER UNIQUE -- \ac{PDD}
         ↳reference ID
8  );
```

```
9
10  CREATE TABLE modules (
11      module_id INTEGER PRIMARY KEY,
12      shortcut VARCHAR(50) NOT NULL,
13      name VARCHAR(255) NOT NULL,
14      kind VARCHAR(255),
15      created_at TIMESTAMP WITH TIME ZONE DEFAULT
          ↳ CURRENT_TIMESTAMP,
16      external_id INTEGER UNIQUE -- \ac{PDD}
          ↳reference ID
17  );
18
19  CREATE TABLE ecu_modules (
20      ecu_id INTEGER REFERENCES ecus(ecu_id) ON
          ↳DELETE CASCADE,
21      module_id INTEGER REFERENCES modules(
          ↳module_id) ON DELETE CASCADE,
22      created_at TIMESTAMP WITH TIME ZONE DEFAULT
          ↳ CURRENT_TIMESTAMP,
23      PRIMARY KEY (ecu_id, module_id)
24  );
```

Listing 5.1: ECU and Module Table Implementation

The junction table `ecu_modules` implements the many-to-many relationship using foreign key constraints as described by Elmasri and Navathe [14]. Each entity includes an `external_id` attribute to maintain mapping with the Parameter Definition Database (PDD), facilitating integration through persistent identifier correlation as recommended by Hohpe and Woolf [18].

PIDs and parameters constitute the lower levels of the hierarchy, with PIDs grouping related parameters for specific modules. The parameter table implements additional attributes for handling the complexity of automotive parameter specifications:

```
1  CREATE TABLE parameters (
2      parameter_id BIGINT PRIMARY KEY,
3      pid_id BIGINT REFERENCES pids(pid_id) ON
          ↳DELETE CASCADE,
4      ecu_id INTEGER,
5      phase_id INTEGER,
```

```
6    name VARCHAR(255) NOT NULL,
7    parameter_name VARCHAR(255),
8    type_id INTEGER REFERENCES
      ↳parameter_data_types(data_type_id),
9    array_definition VARCHAR(50),
10   position INTEGER,
11   factor DECIMAL,
12   unit VARCHAR(50),
13   bias_offset DECIMAL,
14   is_active BOOLEAN DEFAULT true,
15   external_id INTEGER, -- \ac{PDD}reference
      ↳ID
16   FOREIGN KEY (ecu_id, phase_id) REFERENCES
      ↳ecu_phases(ecu_id, phase_id)
17 );
```

Listing 5.2: Parameter Table Implementation

The parameter table incorporates a form of strategic denormalization by including direct references to `ecu_id` and `phase_id` alongside the PID foreign key. While this introduces some redundancy, this approach aims to improve query execution for parameter queries filtered by release phase, a critical operation in the system. Bhattacherjee et al. [4] note that such denormalization can be justified when query performance benefits outweigh the overhead of maintaining consistency, especially for frequently accessed paths in the data model.

A particularly challenging aspect of the implementation was supporting multi-dimensional parameters, which are common in automotive applications for representing lookup tables and characteristic curves [22]. This was addressed through a specialized table structure:

```
1 CREATE TABLE parameter_dimensions (
2    dimension_id BIGINT PRIMARY KEY,
3    parameter_id BIGINT REFERENCES parameters(
      ↳parameter_id) ON DELETE CASCADE,
4    dimension_index INTEGER NOT NULL,
5    default_value NUMERIC NOT NULL,
6    external_id INTEGER, -- \ac{PDD}reference
      ↳ID
7    UNIQUE (parameter_id, dimension_index)
```

```
8  );
```

Listing 5.3: Parameter Dimension Implementation

This implementation follows a modified entity-attribute-value (EAV) pattern while addressing the potential challenges typically associated with such models [13]. The introduction of a dimension index provides an ordered structure to parameter dimensions, enabling efficient representation of arrays and matrices while maintaining the relationship between parameters and their dimensional values.

5.1.2 Version Control Implementation

The version control implementation is a defining feature of the VMAP system, enabling parameter evolution management across different development phases. After evaluating multiple approaches described in Chapter 4, the phase-based versioning model was implemented, creating explicit relationships between parameters and development phases rather than using a generic temporal approach.

```
1   CREATE TABLE releases (
2       release_id INTEGER PRIMARY KEY,
3       name VARCHAR(50) NOT NULL UNIQUE, -- e.g.,
            "24.1", "24.3"
4       description TEXT,
5       is_active BOOLEAN DEFAULT true,
6       created_at TIMESTAMP WITH TIME ZONE DEFAULT
             CURRENT_TIMESTAMP,
7       created_by BIGINT REFERENCES users(user_id)
8   );
9
10  CREATE TABLE release_phases (
11      phase_id INTEGER PRIMARY KEY,
12      release_id INTEGER REFERENCES releases(
            release_id) ON DELETE CASCADE,
13      name VARCHAR(50) NOT NULL, -- e.g., "
            Initial", "PreTest1", "PreTest2", "Final
            "
14      sequence_number INTEGER NOT NULL, --
            Determines the order of phases
15      is_active BOOLEAN DEFAULT true,
```

```
16    created_at TIMESTAMP WITH TIME ZONE DEFAULT
        ↳ CURRENT_TIMESTAMP ,
17    created_by BIGINT REFERENCES users(user_id)
        ↳ ,
18    UNIQUE (release_id, name),
19    UNIQUE (release_id, sequence_number)
20  );
21
22  CREATE TABLE ecu_phases (
23    ecu_id INTEGER REFERENCES ecus(ecu_id) ON
        ↳DELETE CASCADE ,
24    phase_id INTEGER REFERENCES release_phases(
        ↳phase_id) ON DELETE CASCADE ,
25    is_active BOOLEAN DEFAULT true ,
26    is_frozen BOOLEAN DEFAULT false ,
27    frozen_at TIMESTAMP WITH TIME ZONE ,
28    frozen_by BIGINT REFERENCES users(user_id),
29    created_at TIMESTAMP WITH TIME ZONE DEFAULT
        ↳ CURRENT_TIMESTAMP ,
30    created_by BIGINT REFERENCES users(user_id)
        ↳ ,
31    PRIMARY KEY (ecu_id, phase_id)
32  );
```

Listing 5.4: Release and Phase Management Implementation

This implementation supports the bi-annual release cycle with four sequential phases per release. The `sequence_number` field provides explicit ordering of phases within a release, while unique constraints ensure consistency of phase naming and sequencing. The ECU-phase mapping implements the association between ECUs and specific release phases, supporting independent progression of different ECUs through the development cycle as described by Broy [8].

The phase-based approach was selected over temporal versioning for several reasons supported by database design principles:

- Direct phase associations align with the process-based nature of automotive development, where parameters evolve through explicitly defined development stages.

- The explicit phase relationships simplify the implementation of phase transitions and comparison operations, which are fundamental to the automotive development process as described by Pretschner et al. [29].

- The approach aligns with the mental model of automotive development engineers, who conceptualize parameter evolution in terms of distinct development phases rather than continuous time.

Phase management functionality was implemented through a stored procedure for phase transitions and parameter propagation:

```sql
CREATE OR REPLACE FUNCTION copy_phase_data(
    source_ecu_id INTEGER,
    source_phase_id INTEGER,
    target_ecu_id INTEGER,
    target_phase_id INTEGER,
    user_id BIGINT
)
RETURNS TABLE (
    variants_copied INTEGER,
    segments_copied INTEGER
) AS $$
DECLARE
    variant_count INTEGER := 0;
    segment_count INTEGER := 0;
    transaction_id BIGINT;
BEGIN
    -- Get a transaction ID for change tracking
    SELECT nextval('
      change_history_transaction_id_seq') INTO
       transaction_id;

    -- Set the transaction ID for tracking in
      logging trigger
    PERFORM set_config('app.transaction_id',
      transaction_id::text, true);
    PERFORM set_config('app.user_id', user_id::
      text, true);

```

```
24      -- Create a temporary table to map source
          ↳variant IDs to target variant IDs
25      CREATE TEMPORARY TABLE variant_mapping (
26          source_variant_id BIGINT,
27          target_variant_id BIGINT
28      ) ON COMMIT DROP;
29
30      -- Copy variants with modified attributes
          ↳for target phase
31      WITH inserted_variants AS (
32          INSERT INTO variants (
33              pid_id, ecu_id, phase_id, name,
                  ↳code_rule,
34              created_at, created_by, updated_by
35          )
36          SELECT
37              v.pid_id, target_ecu_id,
                  ↳target_phase_id, v.name, v.
                  ↳code_rule,
38              CURRENT_TIMESTAMP, user_id, user_id
39          FROM variants v
40          WHERE v.ecu_id = source_ecu_id AND v.
              ↳phase_id = source_phase_id
41          RETURNING variant_id, pid_id, name,
              ↳code_rule
42      )
43      INSERT INTO variant_mapping (
          ↳source_variant_id, target_variant_id)
44      SELECT v.variant_id, iv.variant_id
45      FROM variants v
46      JOIN inserted_variants iv ON v.pid_id = iv.
          ↳pid_id
47                              AND v.name = iv.
                                  ↳name
48                              AND v.code_rule =
                                  ↳iv.code_rule
49      WHERE v.ecu_id = source_ecu_id AND v.
          ↳phase_id = source_phase_id;
50
51      -- Get count of copied variants
```

```
52    SELECT COUNT(*) INTO variant_count FROM
        ↳variant_mapping;
53
54    -- Copy segments with parameter mapping
55    -- [Additional implementation details
        ↳omitted for brevity]
56
57    RETURN QUERY SELECT variant_count,
        ↳segment_count;
58 END;
59 $$ LANGUAGE plpgsql;
```

Listing 5.5: Phase Transition Function

This implementation provides an atomic transaction for phase transition, copying parameter variants and segments from one phase to another while maintaining proper references and tracking the operation in the change history. The function returns counts of affected entities, providing immediate feedback on the operation's scope. The use of temporary tables for mapping entities between phases implements the identity map pattern described by Fowler [16], ensuring proper relationship preservation during complex operations.

5.1.3 Variant and Segment Management

The variant and segment management implementation realizes the core parameter customization capabilities of the VMAP system:

```
1 CREATE TABLE variants (
2     variant_id BIGINT PRIMARY KEY,
3     pid_id BIGINT REFERENCES pids(pid_id) ON
        ↳DELETE CASCADE,
4     ecu_id INTEGER,
5     phase_id INTEGER,
6     name VARCHAR(100) NOT NULL,
7     code_rule TEXT,
8     created_at TIMESTAMP WITH TIME ZONE DEFAULT
        ↳ CURRENT_TIMESTAMP,
9     created_by BIGINT REFERENCES users(user_id)
        ↳,
```

```
10    updated_by BIGINT REFERENCES users(user_id)
          ↳,
11    FOREIGN KEY (ecu_id, phase_id) REFERENCES
          ↳ecu_phases(ecu_id, phase_id),
12    UNIQUE (phase_id, pid_id, name),
13    UNIQUE (phase_id, pid_id, code_rule)
14  );
15
16  CREATE TABLE segments (
17    segment_id BIGINT PRIMARY KEY,
18    variant_id BIGINT REFERENCES variants(
          ↳variant_id) ON DELETE CASCADE,
19    parameter_id BIGINT REFERENCES parameters(
          ↳parameter_id) ON DELETE CASCADE,
20    dimension_index INTEGER NOT NULL,
21    decimal NUMERIC NOT NULL,
22    created_at TIMESTAMP WITH TIME ZONE DEFAULT
          ↳ CURRENT_TIMESTAMP,
23    created_by BIGINT REFERENCES users(user_id)
          ↳,
24    updated_by BIGINT REFERENCES users(user_id)
25  );
```

Listing 5.6: Variant and Segment Implementation

Each variant is associated with a specific PID, ECU, and phase, creating a three-way relationship that places the variant in both functional and temporal contexts. The uniqueness constraints prevent duplicate variant names or code rules within the same PID and phase, enforcing a critical business rule identified during requirements analysis. The code_rule field stores boolean expressions determining when a variant applies based on vehicle configuration codes, implementing a domain-specific rule language for variant applicability.

Segments form the foundation of parameter customization, linking parameters to variants and storing modified values. The canonical decimal representation for all parameter values, regardless of native data type, implements the canonical model pattern described by Hohpe and Woolf [18], simplifying data manipulation while ensuring consistent value handling.

For parameter value resolution, a specialized database function was implemented:

```sql
1  CREATE OR REPLACE FUNCTION
     ↳resolve_parameter_value(
2      p_parameter_id BIGINT,
3      p_dimension_index INTEGER,
4      p_variant_ids BIGINT[]
5  )
6  RETURNS NUMERIC AS $$
7  DECLARE
8      v_value NUMERIC;
9      v_default_value NUMERIC;
10 BEGIN
11     -- First try to find a segment with the
         ↳given parameter and variant
12     SELECT s.decimal INTO v_value
13     FROM segments s
14     WHERE s.parameter_id = p_parameter_id
15       AND s.dimension_index = p_dimension_index
16       AND s.variant_id = ANY(p_variant_ids)
17     ORDER BY array_position(p_variant_ids, s.
         ↳variant_id)
18     LIMIT 1;
19
20     -- If no segment found, get the default
         ↳value
21     IF v_value IS NULL THEN
22         SELECT pd.default_value INTO
             ↳v_default_value
23         FROM parameter_dimensions pd
24         WHERE pd.parameter_id = p_parameter_id
25           AND pd.dimension_index =
               ↳p_dimension_index;
26
27         -- If no dimension record, try
             ↳parameter default
28         IF v_default_value IS NULL THEN
29             SELECT p.default_value INTO
                 ↳v_default_value
30             FROM parameters p
31             WHERE p.parameter_id =
```

```
           ↳p_parameter_id;
32        END IF;
33
34        v_value := COALESCE(v_default_value, 0)
            ↳;
35     END IF;
36
37     RETURN v_value;
38 END;
39 $$ LANGUAGE plpgsql;
```

Listing 5.7: Parameter Resolution Function

This function implements the value resolution logic required for parameter file genera-
tion, handling the complex logic of finding applicable parameter values based on variant
precedence. The array position-based ordering ensures that variants are applied in the
correct priority sequence, a critical requirement for automotive parameter management
as described by Staron [37].

To support documentation and compliance requirements, a snapshot mechanism was
implemented to capture complete parameter configurations at specific points in time:

```
1 CREATE TABLE documentation_snapshots (
2     snapshot_id INTEGER PRIMARY KEY,
3     name VARCHAR(255) NOT NULL,
4     description TEXT,
5     ecu_id INTEGER,
6     phase_id INTEGER,
7     variant_count INTEGER DEFAULT 0,
8     segment_count INTEGER DEFAULT 0,
9     created_at TIMESTAMP WITH TIME ZONE DEFAULT
        ↳ CURRENT_TIMESTAMP,
10     created_by BIGINT REFERENCES users(user_id)
        ↳,
11     FOREIGN KEY (ecu_id, phase_id) REFERENCES
        ↳ecu_phases(ecu_id, phase_id)
12 );
13
14 CREATE TABLE snapshot_variants (
15     snapshot_variant_id INTEGER PRIMARY KEY,
```

```
16      snapshot_id INTEGER REFERENCES
            ↳documentation_snapshots(snapshot_id) ON
            ↳DELETE CASCADE,
17      original_variant_id BIGINT, -- Reference to
            ↳ the original variant
18      pid_id BIGINT REFERENCES pids(pid_id) ON
            ↳DELETE CASCADE,
19      name VARCHAR(100) NOT NULL,
20      code_rule TEXT,
21      created_at TIMESTAMP WITH TIME ZONE,
22      created_by BIGINT REFERENCES users(user_id)
23  );

24
25  CREATE TABLE snapshot_segments (
26      snapshot_segment_id INTEGER PRIMARY KEY,
27      snapshot_id INTEGER REFERENCES
            ↳documentation_snapshots(snapshot_id) ON
            ↳DELETE CASCADE,
28      snapshot_variant_id INTEGER REFERENCES
            ↳snapshot_variants(snapshot_variant_id)
            ↳ON DELETE CASCADE,
29      original_segment_id BIGINT, -- Reference to
            ↳ the original segment
30      parameter_id BIGINT REFERENCES parameters(
            ↳parameter_id) ON DELETE CASCADE,
31      dimension_index INTEGER NOT NULL,
32      decimal NUMERIC NOT NULL,
33      created_at TIMESTAMP WITH TIME ZONE,
34      created_by BIGINT REFERENCES users(user_id)
35  );
```

Listing 5.8: Documentation Snapshot Implementation

This implementation follows the snapshot pattern described by Fowler [16], creating a complete copy of variant and segment data at specific points in time. Rather than using a temporal database approach with validity periods, the system explicitly materializes historical states, ensuring they remain accessible regardless of subsequent modifications to live data. References to original entities enable traceability between snapshot and live data, implementing the origin tracking pattern described by Tichy [40].

The snapshot creation function automates the process of capturing parameter configurations:

```sql
CREATE OR REPLACE FUNCTION
    ↳create_documentation_snapshot(
    p_name VARCHAR(255),
    p_description TEXT,
    p_ecu_id INTEGER,
    p_phase_id INTEGER,
    p_user_id BIGINT
)
RETURNS INTEGER AS $$
DECLARE
    v_snapshot_id INTEGER;
BEGIN
    -- Create the snapshot record
    INSERT INTO documentation_snapshots (
        name, description, ecu_id, phase_id,
        variant_count, segment_count,
            ↳created_at, created_by
    ) VALUES (
        p_name, p_description, p_ecu_id,
            ↳p_phase_id,
        0, 0, CURRENT_TIMESTAMP, p_user_id
    ) RETURNING snapshot_id INTO v_snapshot_id;

    -- Copy all variants for the ECU and phase
    INSERT INTO snapshot_variants (
        snapshot_id, original_variant_id,
            ↳pid_id, name,
        code_rule, created_at, created_by
    )
    SELECT
        v_snapshot_id, v.variant_id, v.pid_id,
            ↳v.name,
        v.code_rule, v.created_at, v.created_by
    FROM variants v
    WHERE v.ecu_id = p_ecu_id AND v.phase_id =
        ↳p_phase_id;

```

```
32    -- Copy all segments for the variants
33    -- [Implementation details omitted for
        ↳brevity]
34
35    -- Update the counts in the snapshot record
36    -- [Implementation details omitted for
        ↳brevity]
37
38    RETURN v_snapshot_id;
39  END;
40  $$ LANGUAGE plpgsql;
```

Listing 5.9: Documentation Snapshot Function

In the automotive industry, these snapshots serve multiple purposes: providing immutable records of parameter configurations at significant development milestones, supporting quality assurance processes and regulatory compliance requirements, and facilitating comparative analysis between development phases.

5.2 Access Control Implementation

The access control implementation realizes the hybrid role-permission model described in Chapter 4, providing a flexible foundation for managing user permissions across the system.

5.2.1 Role-Based Permission Model

The core RBAC implementation follows the structure defined by Sandhu et al. [32], with tables for users, roles, permissions, and their relationships:

```
1  CREATE TABLE roles (
2      role_id INTEGER PRIMARY KEY,
3      name VARCHAR(255) UNIQUE NOT NULL,
4      description TEXT
5  );
6
7  CREATE TABLE permissions (
```

```
8      permission_id INTEGER PRIMARY KEY,
9      name VARCHAR(255) UNIQUE NOT NULL,
10     description TEXT
11 );
12
13 CREATE TABLE role_permissions (
14     role_id INTEGER REFERENCES roles(role_id)
          ↳ON DELETE CASCADE,
15     permission_id INTEGER REFERENCES
          ↳permissions(permission_id) ON DELETE
          ↳CASCADE,
16     granted_at TIMESTAMP WITH TIME ZONE DEFAULT
          ↳ CURRENT_TIMESTAMP,
17     granted_by BIGINT REFERENCES users(user_id)
          ↳,
18     PRIMARY KEY (role_id, permission_id)
19 );
20
21 CREATE TABLE user_roles (
22     user_id BIGINT REFERENCES users(user_id) ON
          ↳ DELETE CASCADE,
23     role_id INTEGER REFERENCES roles(role_id)
          ↳ON DELETE CASCADE,
24     granted_at TIMESTAMP WITH TIME ZONE DEFAULT
          ↳ CURRENT_TIMESTAMP,
25     granted_by BIGINT REFERENCES users(user_id)
          ↳,
26     PRIMARY KEY (user_id, role_id)
27 );
```

Listing 5.10: Core RBAC Implementation

The model is extended with direct user permissions to implement the hybrid approach:

```
1 CREATE TABLE user_permissions (
2     user_permission_id INTEGER PRIMARY KEY,
3     user_id BIGINT REFERENCES users(user_id) ON
          ↳ DELETE CASCADE,
4     permission_id INTEGER REFERENCES
          ↳permissions(permission_id) ON DELETE
```

```
         ↳CASCADE,
5    granted_by BIGINT REFERENCES users(user_id)
        ↳,
6    granted_at TIMESTAMP WITH TIME ZONE DEFAULT
        ↳ CURRENT_TIMESTAMP,
7    UNIQUE (user_id, permission_id)
8 );
```

Listing 5.11: User-Permission Implementation

This hybrid model allows permissions to be granted either through role assignments or directly to users, providing the flexibility to handle exceptional cases without creating specialized roles, as recommended by Ferraiolo et al. [15].

The permission checking function implements the core security verification logic:

```
1  CREATE OR REPLACE FUNCTION has_permission(
2      p_user_id BIGINT,
3      p_permission_name VARCHAR
4  )
5  RETURNS BOOLEAN AS $$
6  BEGIN
7      RETURN EXISTS (
8          -- Check role-based permissions
9          SELECT 1
10         FROM user_roles ur
11         JOIN role_permissions rp ON ur.role_id
             ↳= rp.role_id
12         JOIN permissions p ON rp.permission_id
             ↳= p.permission_id
13         WHERE ur.user_id = p_user_id AND p.name
             ↳ = p_permission_name
14
15         UNION
16
17         -- Check direct user permissions
18         SELECT 1
19         FROM user_permissions up
20         JOIN permissions p ON up.permission_id
             ↳= p.permission_id
```

```
21          WHERE up.user_id = p_user_id AND p.name
              ↳ = p_permission_name
22      );
23  END;
24  $$ LANGUAGE plpgsql;
```

Listing 5.12: Permission Check Implementation

The UNION-based approach combines role-based and direct permission checks in a single database query, implementing an efficient verification mechanism. This implementation supports the separation of permission checking from application logic, allowing security policies to be enforced consistently across different access paths.

5.2.2 Module-Based Access Control

To complement the role-based permission model, a module-based access control system was implemented to restrict write access based on module assignments:

```
1  CREATE TABLE user_access (
2      user_id BIGINT REFERENCES users(user_id) ON
          ↳ DELETE CASCADE,
3      ecu_id INTEGER REFERENCES ecus(ecu_id) ON
          ↳DELETE CASCADE,
4      module_id INTEGER REFERENCES modules(
          ↳module_id) ON DELETE CASCADE,
5      write_access BOOLEAN DEFAULT true,
6      created_at TIMESTAMP WITH TIME ZONE DEFAULT
          ↳ CURRENT_TIMESTAMP,
7      created_by BIGINT REFERENCES users(user_id)
          ↳,
8      PRIMARY KEY (user_id, ecu_id, module_id),
9      CONSTRAINT user_access_ecu_module_fk
          ↳FOREIGN KEY (ecu_id, module_id)
10     REFERENCES ecu_modules(ecu_id, module_id)
11 );
```

Listing 5.13: Module Access Control Implementation

This implementation establishes a three-way relationship between users, ECUs, and modules, with a boolean flag distinguishing between read and write access. The constraint ensures access is granted only for valid ECU-module combinations, enforcing structural integrity. Database functions verify both role-based permissions and module-based access rights during critical operations:

```
CREATE OR REPLACE FUNCTION
    has_module_write_access(
    p_user_id BIGINT,
    p_ecu_id INTEGER,
    p_module_id INTEGER
)
RETURNS BOOLEAN AS $$
BEGIN
    -- First check if user has administrator
        role (bypasses module restrictions)
    IF EXISTS (
        SELECT 1
        FROM user_roles ur
        JOIN roles r ON ur.role_id = r.role_id
        WHERE ur.user_id = p_user_id AND r.name
            = 'administrator'
    ) THEN
        RETURN TRUE;
    END IF;

    -- Check module-specific write access
    RETURN EXISTS (
        SELECT 1
        FROM user_access ua
        WHERE ua.user_id = p_user_id
          AND ua.ecu_id = p_ecu_id
          AND ua.module_id = p_module_id
          AND ua.write_access = TRUE
    );
END;
$$ LANGUAGE plpgsql;
```

Listing 5.14: Module Access Check Implementation

70

This implementation combines role-based and attribute-based access control approaches, creating what Ferraiolo et al. describe as policy-enhanced RBAC [15]. The function first checks for administrator privileges, which bypass module-specific restrictions, then verifies specific module access rights if necessary.

5.3 Query Optimization Implementation

The complex relationships and large data volumes in automotive parameter management necessitated a comprehensive indexing strategy to support common query patterns.

5.3.1 Indexing Implementation

The indexing strategy focuses on accelerating hierarchical navigation and parameter retrieval operations:

```
1  -- Hierarchical path indexes
2  CREATE INDEX idx_pids_ecu_module ON pids(ecu_id
     ↳, module_id);
3  CREATE INDEX idx_parameters_pid_phase ON
     ↳parameters(pid_id, phase_id);
4  CREATE INDEX idx_variants_pid_phase ON variants
     ↳(pid_id, phase_id);
5  CREATE INDEX idx_segments_variant ON segments(
     ↳variant_id);
6  CREATE INDEX idx_segments_parameter ON segments
     ↳(parameter_id);
7
8  -- Phase-specific indexes
9  CREATE INDEX idx_parameters_phase ON parameters
     ↳(phase_id);
10 CREATE INDEX idx_variants_phase ON variants(
     ↳phase_id);
11
12 -- Text search optimization
13 CREATE EXTENSION IF NOT EXISTS pg_trgm;
14 CREATE INDEX idx_parameters_name_trgm
```

```
15      ON parameters USING gin (name gin_trgm_ops)
          ↳;
16 CREATE INDEX idx_variants_name_trgm
17      ON variants USING gin (name gin_trgm_ops);
```

Listing 5.15: Core Index Implementation

These indexes match the natural navigation paths in the data model, implementing the access path optimization pattern described by Molinaro [25]. The hierarchical path indexes support efficient traversal from ECUs through modules and PIDs to parameters, while the phase-specific indexes accelerate queries filtered by release phase, a common operation in the system.

The trigram-based text search indexes utilize PostgreSQL's specialized text search capabilities as described by Obe and Hsu [28]. These indexes enable efficient pattern matching and similarity-based searches, accommodating the fuzzy search requirements identified during user interviews. According to Karwin [21], these specialized indexes can significantly improve the performance of text-based queries, which are common in parameter management systems where engineers need to locate parameters by name or partial name.

5.3.2 Function-Based Optimization

To optimize complex parameter retrieval operations, several specialized database functions were implemented:

```
1 CREATE OR REPLACE FUNCTION search_parameters(
2     search_term TEXT,
3     max_results INT DEFAULT 20,
4     ecu_id_filter INT DEFAULT NULL,
5     phase_id_filter INT DEFAULT NULL
6 )
7 RETURNS TABLE (
8     parameter_id BIGINT,
9     name VARCHAR,
10    parameter_name VARCHAR,
11    pid_id BIGINT,
12    pid_name VARCHAR,
13    module_id INTEGER,
```

```
14      module_name VARCHAR,
15      ecu_id INTEGER,
16      ecu_name VARCHAR,
17      phase_id INTEGER,
18      relevance REAL
19  ) AS $$
20  BEGIN
21    RETURN QUERY
22    SELECT
23      p.parameter_id,
24      p.name,
25      p.parameter_name,
26      pid.pid_id,
27      pid.name AS pid_name,
28      m.module_id,
29      m.name AS module_name,
30      e.ecu_id,
31      e.name AS ecu_name,
32      p.phase_id,
33      -- Calculate relevance score based on
          ↳different matching criteria
34      (CASE
35          -- Exact matches get highest score
36          WHEN p.name ILIKE search_term OR p.
              ↳parameter_name ILIKE search_term
              ↳THEN 1.0
37          -- Starts with gets high score
38          WHEN p.name ILIKE (search_term || '%')
              ↳OR p.parameter_name ILIKE (
              ↳search_term || '%') THEN 0.8
39          -- Contains gets medium score
40          ELSE GREATEST(
41              similarity(p.name, search_term),
42              similarity(p.parameter_name,
                  ↳search_term)
43          )
44      END) AS relevance
45    FROM
46      parameters p
47    JOIN
```

```
48        pids\ac{PID}ON p.pid_id = pid.pid_id
49     JOIN
50        modules m ON pid.module_id = m.module_id
51     JOIN
52        ecus e ON pid.ecu_id = e.ecu_id
53     WHERE
54        -- Apply filters if provided
55        (ecu_id_filter IS NULL OR e.ecu_id =
            ↳ecu_id_filter) AND
56        (phase_id_filter IS NULL OR p.phase_id =
            ↳phase_id_filter) AND
57        -- Match parameter name or display name
58        (p.name ILIKE '%' || search_term || '%' OR
59         p.parameter_name ILIKE '%' || search_term
            ↳|| '%') AND
60        -- Only active parameters
61        p.is_active = true
62     ORDER BY
63        relevance DESC,
64        p.name ASC
65     LIMIT max_results;
66  END;
67  $$ LANGUAGE plpgsql;
```

Listing 5.16: Parameter Search Function

This function implements a sophisticated parameter search algorithm that combines exact matches, prefix matches, and similarity-based matching to provide relevant search results. The relevance scoring approach ensures that the most relevant parameters appear first in the results, following the information retrieval principles described by Obe and Hsu [28].

For phase comparison operations, a specialized function implements the parameter difference detection logic:

```
1  CREATE OR REPLACE FUNCTION
      ↳get_pids_with_changes_between_phases(
2      p_ecu_id INTEGER,
3      p_module_id INTEGER,
4      p_source_phase_id INTEGER,
```

```
5        p_target_phase_id INTEGER
6  )
7  RETURNS TABLE (
8        pid_id BIGINT,
9        name VARCHAR,
10       has_param_changes BOOLEAN,
11       has_variant_changes BOOLEAN,
12       has_segment_changes BOOLEAN
13 ) AS $
14 BEGIN
15       RETURN QUERY
16       WITH common_pids AS (
17            -- Get PIDs available in at least one
                  ↳of the phases
18            SELECT DISTINCT p.pid_id, p.name
19            FROM pids p
20            LEFT JOIN pid_phases pp1 ON p.pid_id =
                  ↳pp1.pid_id AND pp1.phase_id =
                  ↳p_source_phase_id
21            LEFT JOIN pid_phases pp2 ON p.pid_id =
                  ↳pp2.pid_id AND pp2.phase_id =
                  ↳p_target_phase_id
22            WHERE p.ecu_id = p_ecu_id AND p.
                  ↳module_id = p_module_id
23            AND (pp1.phase_id IS NOT NULL OR pp2.
                  ↳phase_id IS NOT NULL)
24       ),
25       -- Parameter changes detection
26       param_changes AS (
27            SELECT cp.pid_id
28            FROM common_pids cp
29            WHERE EXISTS (
30                 -- Parameters that exist in target
                       ↳but not in source
31                 SELECT 1 FROM parameters pt
32                 WHERE pt.pid_id = cp.pid_id
33                   AND pt.phase_id =
                         ↳p_target_phase_id
34                   AND NOT EXISTS (
35                       SELECT 1 FROM parameters ps
```

```
36              WHERE ps.pid_id = cp.pid_id
37                AND ps.phase_id =
                    ↳p_source_phase_id
38                AND ps.external_id = pt.
                    ↳external_id
39            )
40        ) OR EXISTS (
41            -- Parameters that exist in source
              ↳but not in target
42            SELECT 1 FROM parameters ps
43            WHERE ps.pid_id = cp.pid_id
44              AND ps.phase_id =
                  ↳p_source_phase_id
45              AND NOT EXISTS (
46                  SELECT 1 FROM parameters pt
47                  WHERE pt.pid_id = cp.pid_id
48                    AND pt.phase_id =
                      ↳p_target_phase_id
49                    AND pt.external_id = ps.
                      ↳external_id
50              )
51        ) OR EXISTS (
52            -- Parameters that exist in both
              ↳but have differences
53            SELECT 1 FROM parameters ps
54            JOIN parameters pt ON ps.
              ↳external_id = pt.external_id
55            WHERE ps.pid_id = cp.pid_id
56              AND pt.pid_id = cp.pid_id
57              AND ps.phase_id =
                  ↳p_source_phase_id
58              AND pt.phase_id =
                  ↳p_target_phase_id
59              AND (
60                  ps.name != pt.name OR
61                  ps.parameter_name != pt.
                    ↳parameter_name OR
62                  ps.type_id != pt.type_id OR
63                  ps.factor != pt.factor OR
64                  ps.bias_offset != pt.
```

76

```
                            ↳bias_offset OR
65                      ps.unit != pt.unit
66                  )
67          )
68      ),
69      -- Variant changes detection
70      variant_changes AS (
71          SELECT cp.pid_id
72          FROM common_pids cp
73          WHERE EXISTS (
74              -- Different variant counts between
                    ↳ phases
75              SELECT 1
76              FROM (
77                  SELECT pid_id, count(*) as cnt
78                  FROM variants
79                  WHERE pid_id = cp.pid_id AND
                        ↳phase_id = p_source_phase_id
80                  GROUP BY pid_id
81
82                  UNION ALL
83
84                  SELECT pid_id, -count(*) as cnt
85                  FROM variants
86                  WHERE pid_id = cp.pid_id AND
                        ↳phase_id = p_target_phase_id
87                  GROUP BY pid_id
88              ) v
89              GROUP BY v.pid_id
90              HAVING sum(v.cnt) != 0
91          ) OR EXISTS (
92              -- Variants with same name but
                    ↳different code rules
93              SELECT 1
94              FROM variants vs
95              JOIN variants vt ON vs.name = vt.
                    ↳name
96              WHERE vs.pid_id = cp.pid_id
97                AND vt.pid_id = cp.pid_id
98                AND vs.phase_id =
```

```
                             ↳p_source_phase_id
 99              AND vt.phase_id =
                             ↳p_target_phase_id
100              AND vs.code_rule != vt.code_rule
101          )
102      ),
103      -- Segment changes detection
104      segment_changes AS (
105          SELECT cp.pid_id
106          FROM common_pids cp
107          JOIN parameters ps ON ps.pid_id = cp.
                 ↳pid_id AND ps.phase_id =
                 ↳p_source_phase_id
108          JOIN parameters pt ON pt.pid_id = cp.
                 ↳pid_id AND pt.phase_id =
                 ↳p_target_phase_id
109                          AND pt.external_id =
                                 ↳ps.external_id
110          JOIN variants vs ON vs.pid_id = cp.
                 ↳pid_id AND vs.phase_id =
                 ↳p_source_phase_id
111          JOIN variants vt ON vt.pid_id = cp.
                 ↳pid_id AND vt.phase_id =
                 ↳p_target_phase_id
112                      AND vt.name = vs.name
113          JOIN segments ss ON ss.parameter_id =
                 ↳ps.parameter_id AND ss.variant_id =
                 ↳vs.variant_id
114          JOIN segments st ON st.parameter_id =
                 ↳pt.parameter_id AND st.variant_id =
                 ↳vt.variant_id
115                          AND st.dimension_index =
                                 ↳ ss.dimension_index
116          WHERE ss.decimal != st.decimal
117
118          UNION
119
120          -- Segments in source not in target
121          SELECT cp.pid_id
122          FROM common_pids cp
```

```sql
123         JOIN parameters ps ON ps.pid_id = cp.
              ↳pid_id AND ps.phase_id =
              ↳p_source_phase_id
124         JOIN parameters pt ON pt.pid_id = cp.
              ↳pid_id AND pt.phase_id =
              ↳p_target_phase_id
125                         AND pt.external_id =
                              ↳ps.external_id
126         JOIN variants vs ON vs.pid_id = cp.
              ↳pid_id AND vs.phase_id =
              ↳p_source_phase_id
127         JOIN variants vt ON vt.pid_id = cp.
              ↳pid_id AND vt.phase_id =
              ↳p_target_phase_id
128                         AND vt.name = vs.name
129         JOIN segments ss ON ss.parameter_id =
              ↳ps.parameter_id AND ss.variant_id =
              ↳vs.variant_id
130         WHERE NOT EXISTS (
131             SELECT 1 FROM segments st
132             WHERE st.parameter_id = pt.
                  ↳parameter_id
133                 AND st.variant_id = vt.variant_id
134                 AND st.dimension_index = ss.
                      ↳dimension_index
135         )
136
137         UNION
138
139         -- Segments in target not in source
140         SELECT cp.pid_id
141         FROM common_pids cp
142         JOIN parameters ps ON ps.pid_id = cp.
              ↳pid_id AND ps.phase_id =
              ↳p_source_phase_id
143         JOIN parameters pt ON pt.pid_id = cp.
              ↳pid_id AND pt.phase_id =
              ↳p_target_phase_id
144                         AND pt.external_id =
                              ↳ps.external_id
```

```
145        JOIN variants vs ON vs.pid_id = cp.
             ↳pid_id AND vs.phase_id =
             ↳p_source_phase_id
146        JOIN variants vt ON vt.pid_id = cp.
             ↳pid_id AND vt.phase_id =
             ↳p_target_phase_id
147                      AND vt.name = vs.name
148        JOIN segments st ON st.parameter_id =
             ↳pt.parameter_id AND st.variant_id =
             ↳vt.variant_id
149        WHERE NOT EXISTS (
150            SELECT 1 FROM segments ss
151            WHERE ss.parameter_id = ps.
               ↳parameter_id
152            AND ss.variant_id = vs.variant_id
153            AND ss.dimension_index = st.
               ↳dimension_index
154        )
155    )
156    -- Return\acp{PID} with change indicators
157    SELECT
158        cp.pid_id,
159        cp.name,
160        (cp.pid_id IN (SELECT pid_id FROM
             ↳param_changes)) AS has_param_changes
             ↳,
161        (cp.pid_id IN (SELECT pid_id FROM
             ↳variant_changes)) AS
             ↳has_variant_changes,
162        (cp.pid_id IN (SELECT pid_id FROM
             ↳segment_changes)) AS
             ↳has_segment_changes
163    FROM
164        common_pids cp
165    WHERE
166        -- Only return\acp{PID} that have at
             ↳least one type of change
167        cp.pid_id IN (SELECT pid_id FROM
             ↳param_changes)
168        OR cp.pid_id IN (SELECT pid_id FROM
```

```
              ↳variant_changes)
169        OR cp.pid_id IN (SELECT pid_id FROM
              ↳segment_changes)
170    ORDER BY
171        cp.name;
172 END;
173 $ LANGUAGE plpgsql;
```

Listing 5.17: Phase Comparison Function

This function implements an efficient algorithm for identifying differences between parameter configurations in different release phases. Unlike a naive approach that would require retrieving and comparing all parameters, variants, and segments, this function performs the comparison directly in the database using specialized subqueries for each type of change. This implementation follows the set-based processing pattern recommended by Date [12], leveraging the database engine's capabilities for efficient data comparison.

5.4 Change Tracking Implementation

The change tracking implementation provides comprehensive audit capabilities for all parameter data modifications, addressing both regulatory compliance requirements and supporting diagnostic analysis of parameter evolution.

5.4.1 Automatic Change Logging

A database trigger mechanism automatically records all modifications to critical entities:

```
1 CREATE TABLE change_history (
2     change_id BIGINT PRIMARY KEY,
3     user_id BIGINT REFERENCES users(user_id),
4     ecu_id INTEGER,
5     phase_id INTEGER,
6     entity_type VARCHAR(50) NOT NULL,
7     entity_id BIGINT NOT NULL,
8     change_type VARCHAR(50),
9     old_values JSONB,
```

```
10      new_values JSONB,
11      changed_at TIMESTAMP WITH TIME ZONE DEFAULT
            ↳ CURRENT_TIMESTAMP,
12      transaction_id BIGINT NOT NULL
13  );
14
15  CREATE OR REPLACE FUNCTION log_change()
16  RETURNS TRIGGER AS $
17  DECLARE
18      transaction_id BIGINT;
19      change_type VARCHAR(50);
20      old_values JSONB;
21      new_values JSONB;
22      entity_id BIGINT;
23      phase_id INTEGER;
24      ecu_id INTEGER;
25  BEGIN
26      -- Get the user ID from the application
            ↳context
27      user_id := NULLIF(current_setting('app.
            ↳user_id', TRUE), '')::BIGINT;
28
29      -- Get the current transaction ID or create
            ↳ a new one
30      SELECT COALESCE(NULLIF(current_setting('app
            ↳.transaction_id', TRUE), '')::BIGINT,
31                  nextval('
                        ↳change_history_transaction_id_seq
                        ↳')) INTO transaction_id;
32
33      -- Set the transaction ID for other
            ↳triggers in the same transaction
34      PERFORM set_config('app.transaction_id',
            ↳transaction_id::text, TRUE);
35
36      -- Determine entity_id and context
            ↳information based on table and operation
37      CASE TG_TABLE_NAME
38          WHEN 'variants' THEN
39              entity_id := CASE WHEN TG_OP = '
```

```
     ↳DELETE' THEN OLD.variant_id ELSE
     ↳ NEW.variant_id END;
40          ecu_id := CASE WHEN TG_OP = 'DELETE
     ↳' THEN OLD.ecu_id ELSE NEW.
     ↳ecu_id END;
41          phase_id := CASE WHEN TG_OP = '
     ↳DELETE' THEN OLD.phase_id ELSE
     ↳NEW.phase_id END;
42
43      WHEN 'segments' THEN
44          entity_id := CASE WHEN TG_OP = '
     ↳DELETE' THEN OLD.segment_id ELSE
     ↳ NEW.segment_id END;
45
46          -- Get phase_id from associated
     ↳variant
47          IF TG_OP = 'DELETE' THEN
48              SELECT v.phase_id, v.ecu_id
     ↳INTO phase_id, ecu_id
49              FROM variants v WHERE v.
     ↳variant_id = OLD.variant_id;
50          ELSE
51              SELECT v.phase_id, v.ecu_id
     ↳INTO phase_id, ecu_id
52              FROM variants v WHERE v.
     ↳variant_id = NEW.variant_id;
53          END IF;
54
55      -- Additional entity types handled here
56      -- [Implementation details omitted for
     ↳brevity]
57    END CASE;
58
59    -- Determine change type and capture entity
     ↳ state
60    IF TG_OP = 'INSERT' THEN
61        change_type := 'CREATE';
62        old_values := NULL;
63        new_values := to_jsonb(NEW);
64    ELSIF TG_OP = 'UPDATE' THEN
```

```
65          change_type := 'UPDATE';
66          old_values := to_jsonb(OLD);
67          new_values := to_jsonb(NEW);
68      ELSIF TG_OP = 'DELETE' THEN
69          change_type := 'DELETE';
70          old_values := to_jsonb(OLD);
71          new_values := NULL;
72      END IF;
73
74      -- Remove large or sensitive fields from
           ↳the JSON
75      IF old_values IS NOT NULL THEN
76          old_values := old_values - 'created_at'
               ↳ - 'updated_at';
77      END IF;
78
79      IF new_values IS NOT NULL THEN
80          new_values := new_values - 'created_at'
               ↳ - 'updated_at';
81      END IF;
82
83      -- Insert into change_history
84      INSERT INTO change_history (
85          user_id, ecu_id, phase_id, entity_type,
               ↳ entity_id,
86          change_type, old_values, new_values,
               ↳transaction_id, changed_at
87      ) VALUES (
88          user_id, ecu_id, phase_id,
               ↳TG_TABLE_NAME, entity_id,
89          change_type, old_values, new_values,
               ↳transaction_id, CURRENT_TIMESTAMP
90      );
91
92      RETURN NULL;
93  END;
94  $ LANGUAGE plpgsql;
95
96  -- Apply the trigger to critical entities
97  CREATE TRIGGER variants_change_trigger
```

```
98  AFTER INSERT OR UPDATE OR DELETE ON variants
99  FOR EACH ROW EXECUTE FUNCTION log_change();
100
101 CREATE TRIGGER segments_change_trigger
102 AFTER INSERT OR UPDATE OR DELETE ON segments
103 FOR EACH ROW EXECUTE FUNCTION log_change();
104
105 -- Additional triggers for other entities
106 -- [Implementation details omitted for brevity]
```

Listing 5.18: Change Tracking Trigger

This implementation captures complete entity states rather than just modified fields, storing them as JSONB documents. JSON path operators enable efficient extraction and querying of specific changes without complex joins, leveraging PostgreSQL's advanced document storage capabilities as described by Obe and Hsu [28]. The automatic trigger system ensures consistent logging regardless of how changes are made, preventing circumvention of audit controls.

The change tracking system implements several key patterns described by Fowler [16]: the unit of work pattern through transaction grouping, the state snapshot pattern through complete entity state capture, and the audit log pattern through comprehensive change recording. These patterns collectively support the sophisticated audit requirements common in regulated industries like automotive development.

5.4.2 Change Analysis Functions

Specialized functions enable analysis of parameter evolution over time:

```
1  CREATE OR REPLACE FUNCTION
       ↳get_parameter_history(
2      p_parameter_id BIGINT,
3      p_start_date TIMESTAMP WITH TIME ZONE
           ↳DEFAULT NULL,
4      p_end_date TIMESTAMP WITH TIME ZONE DEFAULT
           ↳ NULL
5  )
6  RETURNS TABLE (
7      change_id BIGINT,
```

```
 8        changed_at TIMESTAMP WITH TIME ZONE,
 9        user_name VARCHAR,
10        variant_name VARCHAR,
11        old_value NUMERIC,
12        new_value NUMERIC,
13        change_type VARCHAR,
14        transaction_id BIGINT,
15        phase_name VARCHAR
16   ) AS $
17   BEGIN
18        RETURN QUERY
19        SELECT
20            ch.change_id,
21            ch.changed_at,
22            u.first_name || ' ' || u.last_name AS
                 ↳user_name,
23            v.name AS variant_name,
24            (ch.old_values->>'decimal')::NUMERIC AS
                 ↳ old_value,
25            (ch.new_values->>'decimal')::NUMERIC AS
                 ↳ new_value,
26            ch.change_type,
27            ch.transaction_id,
28            rp.name AS phase_name
29        FROM
30            change_history ch
31            JOIN users u ON ch.user_id = u.user_id
32            JOIN segments s ON ch.entity_id = s.
                 ↳segment_id
33                        AND ch.entity_type = '
                              ↳segments'
34            JOIN variants v ON s.variant_id = v.
                 ↳variant_id
35            JOIN release_phases rp ON ch.phase_id =
                 ↳ rp.phase_id
36        WHERE
37            s.parameter_id = p_parameter_id
38            AND (p_start_date IS NULL OR ch.
                 ↳changed_at >= p_start_date)
39            AND (p_end_date IS NULL OR ch.
```

```
            ↳changed_at <= p_end_date)
40      ORDER BY
41          ch.changed_at DESC;
42  END;
43  $ LANGUAGE plpgsql;
```

Listing 5.19: Parameter History Function

This function provides a comprehensive view of how a parameter has evolved over time, including who made changes, which variants were affected, and the specific value modifications. The function uses JSON path operators to extract specific values from the JSONB store, implementing what Obe and Hsu [28] describe as the document extraction pattern. By including user and phase information in the result set, the function provides comprehensive context for each change, supporting both diagnostic and compliance requirements.

5.4.3 Partitioning Implementation

To address the potential performance and management challenges associated with the growing change history table, a partitioning strategy was implemented based on the phase model that forms the foundation of the VMAP system. PostgreSQL's declarative partitioning capabilities provide an efficient mechanism for dividing the large change_history table into smaller, more manageable segments based on logical boundaries [28].

Partition Design

The partitioning strategy leverages the phase-based organization of parameter data, creating individual partitions for each development phase to align the physical storage structure with the logical organization of the data. This approach offers several advantages over time-based or release-based partitioning alternatives.

Phase-based partitioning creates a direct correspondence between the database's physical organization and the domain's logical structure. According to Schwartz et al. [33], alignment between logical data organization and physical partitioning is an essential factor for successful partition design.

Queries filtering by phase (a predominant access pattern in the system) benefit from partition pruning, where PostgreSQL automatically eliminates irrelevant partitions from

consideration. This pruning mechanism reduces I/O and improves query response time for phase-specific operations [28].

Additionally, phase-based partitioning simplifies maintenance operations such as archiving or removing historical data for completed phases. When phases transition from active development to frozen status, their corresponding partitions can be managed accordingly without affecting ongoing development work.

The implementation uses PostgreSQL's LIST partitioning strategy, with phase_id as the partition key:

```sql
-- Create the partitioned change_history table
CREATE TABLE change_history (
    change_id BIGINT NOT NULL,
    user_id BIGINT REFERENCES users(user_id),
    ecu_id INTEGER,
    phase_id INTEGER,
    entity_type VARCHAR(50) NOT NULL,
    entity_id BIGINT NOT NULL,
    change_type VARCHAR(50),
    old_values JSONB,
    new_values JSONB,
    changed_at TIMESTAMP WITH TIME ZONE DEFAULT
        CURRENT_TIMESTAMP,
    transaction_id BIGINT NOT NULL
) PARTITION BY LIST (phase_id);

-- Create default partition for NULL phase_id values
-- (for system-level changes not tied to a specific phase)
CREATE TABLE change_history_default PARTITION
    OF change_history DEFAULT;
ALTER TABLE change_history_default ADD PRIMARY
    KEY (change_id);

-- Create indexes on the default partition
CREATE INDEX idx_ch_default_entity ON
    change_history_default(entity_type,
    entity_id, changed_at);
```

```
23  CREATE INDEX idx_ch_default_user ON
      ↳change_history_default(user_id, changed_at);
24  CREATE INDEX idx_ch_default_ecu ON
      ↳change_history_default(ecu_id, changed_at);
25  CREATE INDEX idx_ch_default_transaction ON
      ↳change_history_default(transaction_id,
      ↳changed_at);
26  CREATE INDEX idx_ch_default_changed_at ON
      ↳change_history_default(changed_at);
```

Listing 5.20: Change History Table Partitioning Definition

The default partition handles system-level changes that are not associated with specific phases, ensuring all change records have an appropriate storage location even if they don't fit the defined partition scheme. This approach aligns with PostgreSQL best practices for handling exceptions to the primary partitioning scheme [28].

Automated Partition Management

To ensure that appropriate partitions exist for each phase, an automated partition creation mechanism was implemented using database triggers. When a new phase is created, the system automatically creates a corresponding partition for change history records associated with that phase:

```
1   -- Procedure to safely create partitions by
      ↳moving data first
2   CREATE OR REPLACE PROCEDURE
      ↳create_change_history_partition_for_phase(
      ↳p_phase_id INTEGER, p_phase_name VARCHAR)
3   LANGUAGE plpgsql
4   AS $$
5   DECLARE
6       partition_name TEXT;
7       partition_exists BOOLEAN;
8       row_count INTEGER;
9   BEGIN
10      partition_name := 'change_history_p' ||
          ↳p_phase_id;
11
```

```
12      -- Check if partition already exists
13   SELECT EXISTS (
14       SELECT FROM pg_class c
15       JOIN pg_namespace n ON n.oid = c.
          ↳relnamespace
16       WHERE c.relname = partition_name AND n.
          ↳nspname = current_schema()
17   ) INTO partition_exists;
18
19   IF NOT partition_exists THEN
20       -- Check if we have rows in the default
          ↳ partition with this phase_id
21       EXECUTE 'SELECT COUNT(*) FROM
          ↳change_history_default WHERE
          ↳phase_id = $1'
22       INTO row_count
23       USING p_phase_id;
24
25       IF row_count > 0 THEN
26           -- Move the rows to a temporary
              ↳table
27           EXECUTE 'CREATE TEMPORARY TABLE
              ↳temp_change_history AS
28                   SELECT * FROM
                      ↳change_history_default
29                   WHERE phase_id = $1'
30           USING p_phase_id;
31
32           -- Delete those rows from the
              ↳default partition
33           EXECUTE 'DELETE FROM
              ↳change_history_default
34                   WHERE phase_id = $1'
35           USING p_phase_id;
36       END IF;
37
38       -- Create the partition
39       EXECUTE format(
40           'CREATE TABLE %I PARTITION OF
              ↳change_history FOR VALUES IN (%L
```

```
                    ↳)',
41              partition_name, p_phase_id
42          );

43
44          -- Add primary key
45          EXECUTE format(
46              'ALTER TABLE %I ADD PRIMARY KEY (
                    ↳change_id)',
47              partition_name
48          );

49
50          -- Add comment
51          EXECUTE format(
52              'COMMENT ON TABLE %I IS %L',
53              partition_name, 'Change history for
                    ↳ phase ' || p_phase_name
54          );

55
56          -- Add indexes
57          EXECUTE format(
58              'CREATE INDEX idx_%s_entity ON %I(
                    ↳entity_type, entity_id,
                    ↳changed_at)',
59              substring(partition_name from '
                    ↳change_history_(.*)'),
                    ↳partition_name
60          );

61
62          EXECUTE format(
63              'CREATE INDEX idx_%s_user ON %I(
                    ↳user_id, changed_at)',
64              substring(partition_name from '
                    ↳change_history_(.*)'),
                    ↳partition_name
65          );

66
67          EXECUTE format(
68              'CREATE INDEX idx_%s_ecu ON %I(
                    ↳ecu_id, changed_at)',
69              substring(partition_name from '
```

```
                            ↳change_history_(.*)'),
                            ↳partition_name
70          );

71

72          EXECUTE format(
73              'CREATE INDEX idx_%s_transaction ON
                    ↳ %I(transaction_id, changed_at)'
                    ↳,
74              substring(partition_name from '
                    ↳change_history_(.*)'),
                    ↳partition_name
75          );

76

77          EXECUTE format(
78              'CREATE INDEX idx_%s_changed_at ON
                    ↳%I(changed_at)',
79              substring(partition_name from '
                    ↳change_history_(.*)'),
                    ↳partition_name
80          );

81

82          -- If we moved rows, insert them back
               ↳into the new partition
83          IF row_count > 0 THEN
84              EXECUTE 'INSERT INTO ' ||
                    ↳partition_name || ' SELECT *
                    ↳FROM temp_change_history';

85

86              -- Clean up the temporary table
87              EXECUTE 'DROP TABLE
                    ↳temp_change_history';
88          END IF;

89

90          RAISE NOTICE 'Created partition % for
               ↳phase % (ID %). Moved % rows.',
91                       partition_name,
                            ↳p_phase_name, p_phase_id
                            ↳, row_count;
92      END IF;
93  END;
```

```
94   $$;
95
96   -- Function to trigger partition creation when
       ↳a new phase is created
97   CREATE OR REPLACE FUNCTION
       ↳create_phase_partition_trigger()
98   RETURNS TRIGGER
99   LANGUAGE plpgsql
100  AS $$
101  BEGIN
102      CALL
           ↳create_change_history_partition_for_phase
           ↳(NEW.phase_id, NEW.name);
103      RETURN NEW;
104  END;
105  $$;
106
107  -- Trigger to automatically create partitions
       ↳for new phases
108  CREATE TRIGGER trig_create_phase_partition
109  AFTER INSERT ON release_phases
110  FOR EACH ROW
111  EXECUTE FUNCTION create_phase_partition_trigger
       ↳();
```

Listing 5.21: Automatic Partition Creation

This implementation follows the approach described by Schwartz et al. [33] for maintaining partitioned tables in production environments. The procedure handles the complex task of safely creating new partitions, including migrating any existing records from the default partition to the newly created phase-specific partition. This migration step maintains data integrity during the partition creation process.

The partitioning implementation is particularly noteworthy for its attention to index creation. Each partition receives a complete set of indexes matching those on the default partition, ensuring consistent query performance across all partitions. According to Karwin [21], maintaining consistent indexing across partitions is essential for predictable query execution plans and optimal performance.

Query Performance Implications

The phase-based partitioning approach has significant performance implications for change history queries. For queries that include a phase_id filter criterion—a dominant pattern in the VMAP system—PostgreSQL can use partition pruning to eliminate irrelevant partitions from consideration. Salzberg and Tsotras [30] note that with effective partitioning, query performance can improve by an order of magnitude for large temporal datasets.

This performance benefit is particularly valuable for the most common access patterns in the system. When comparing parameters between phases, queries can focus exclusively on the relevant phase partitions. When auditing a user's activity within a specific phase, queries can be efficiently directed to the appropriate partition. When retrieving the modification history for a specific entity within a phase, partition pruning significantly reduces the search space.

Importantly, the partitioning scheme is transparent to application code, requiring no modifications to existing queries. The PostgreSQL query planner automatically applies partition pruning based on the WHERE clause predicates, allowing the application to benefit from partitioning without explicit partition selection in SQL statements.

Archiving Strategy

The phase-based partitioning strategy provides a natural foundation for long-term data archiving. As phases transition from active development to frozen status, their change history records become less frequently accessed but must be preserved for regulatory compliance and occasional historical analysis.

The partitioning implementation supports a future archiving strategy where frozen phase partitions could be compressed using PostgreSQL's table compression options to reduce storage requirements. According to Bhattacherjee et al. [4], efficient archiving strategies are essential for managing the growing storage requirements of versioned datasets.

Frozen phase partitions could be moved to alternative tablespaces, relocating older partitions to lower-cost storage media while maintaining accessibility. For very old phases that have progressed to production, partitions could be detached from the main table and accessed only when specifically required. These archiving capabilities align with the recommendations of Al-Kateb et al. [2] for managing historical data in temporal database systems, where different storage tiers can be leveraged based on data age and access frequency. Snodgrass [35] suggests that this tiered approach to historical

data management provides an optimal balance between accessibility and resource utilization.

At defined archive thresholds, typically when a release including multiple phases transitions to production status, historical data can be compressed or moved while preserving selective access capabilities. This approach ensures that historical reference data remains available for compliance purposes while optimizing system performance for active development phases.

5.5 Integration Implementation

The integration implementation connects VMAP with external enterprise systems providing parameter definitions and vehicle configuration data. Two primary integration points were implemented: synchronization with the PDD and communication with the Vehicle Configuration Database.

5.5.1 Parameter Definition Database Synchronization

Synchronization with the PDD was implemented through a structured process that maintains parameter definitions across different release phases. The database structure for tracking synchronization includes:

```
1  CREATE TABLE pdd_sync_history (
2      sync_id INTEGER PRIMARY KEY,
3      ecu_id INTEGER,
4      phase_id INTEGER,
5      database_name VARCHAR(255),
6      database_label VARCHAR(255),
7      sync_date TIMESTAMP WITH TIME ZONE DEFAULT
          ↳CURRENT_TIMESTAMP,
8      status VARCHAR(50) NOT NULL,
9      modules_count INTEGER DEFAULT 0,
10     pids_count INTEGER DEFAULT 0,
11     parameters_count INTEGER DEFAULT 0,
12     executed_by BIGINT REFERENCES users(user_id
          ↳),
13     transaction_id BIGINT,
```

```
14      entity_changes TEXT,
15      FOREIGN KEY (ecu_id, phase_id) REFERENCES
          ↳ecu_phases(ecu_id, phase_id)
16  );
```

Listing 5.22: Synchronization Tracking Implementation

This infrastructure implements the external system registry pattern described by Hohpe and Woolf [18], tracking synchronization operations with external database systems. The detailed history supports both operational monitoring and troubleshooting of integration issues.

The parameter loading implementation uses a bulk loading approach for efficiency:

```
1   CREATE OR REPLACE FUNCTION load_parameters_bulk
      ↳(
2       p_parameters JSONB,
3       p_ecu_id INTEGER,
4       p_phase_id INTEGER,
5       p_user_id BIGINT
6   )
7   RETURNS INTEGER AS $
8   DECLARE
9       v_count INTEGER;
10  BEGIN
11      -- Create temporary table for bulk loading
12      CREATE TEMPORARY TABLE temp_parameters (
13          parameter_id BIGINT,
14          pid_id BIGINT,
15          name VARCHAR(255),
16          parameter_name VARCHAR(255),
17          type_id INTEGER,
18          array_definition VARCHAR(50),
19          position INTEGER,
20          factor DECIMAL,
21          unit VARCHAR(50),
22          bias_offset DECIMAL,
23          external_id INTEGER
24      ) ON COMMIT DROP;
25
```

```sql
26        -- Load data from JSON into temporary table
27    INSERT INTO temp_parameters
28    SELECT
29        (elem->>'parameter_id')::BIGINT,
30        (elem->>'pid_id')::BIGINT,
31        elem->>'name',
32        elem->>'parameter_name',
33        (elem->>'type_id')::INTEGER,
34        elem->>'array_definition',
35        (elem->>'position')::INTEGER,
36        (elem->>'factor')::DECIMAL,
37        elem->>'unit',
38        (elem->>'bias_offset')::DECIMAL,
39        (elem->>'external_id')::INTEGER
40    FROM jsonb_array_elements(p_parameters) AS
         ↳elem;
41
42        -- Bulk insert into parameters table
43    INSERT INTO parameters (
44        parameter_id, pid_id, ecu_id, phase_id,
45        name, parameter_name, type_id,
             ↳array_definition,
46        position, factor, unit, bias_offset,
             ↳external_id,
47        created_by
48    )
49    SELECT
50        tp.parameter_id, tp.pid_id, p_ecu_id,
             ↳p_phase_id,
51        tp.name, tp.parameter_name, tp.type_id,
             ↳ tp.array_definition,
52        tp.position, tp.factor, tp.unit, tp.
             ↳bias_offset, tp.external_id,
53        p_user_id
54    FROM temp_parameters tp
55    ON CONFLICT (parameter_id) DO UPDATE SET
56        name = EXCLUDED.name,
57        parameter_name = EXCLUDED.
             ↳parameter_name,
58        type_id = EXCLUDED.type_id,
```

```
59        array_definition = EXCLUDED.
            ↳array_definition,
60        position = EXCLUDED.position,
61        factor = EXCLUDED.factor,
62        unit = EXCLUDED.unit,
63        bias_offset = EXCLUDED.bias_offset;
64
65    GET DIAGNOSTICS v_count = ROW_COUNT;
66    RETURN v_count;
67 END;
68 $ LANGUAGE plpgsql;
```

Listing 5.23: Bulk Parameter Loading Implementation

This bulk loading approach uses temporary tables and JSON parsing to efficiently process large volumes of parameter data, implementing the bulk transfer pattern described by Hohpe and Woolf [18]. The use of PostgreSQL's ON CONFLICT clause implements an efficient upsert operation, enabling both insertion of new parameters and updating of existing ones in a single operation.

5.5.2 Vehicle Configuration Database Integration

The VCD integration enables validation of variant code rules and supports parameter file generation for specific vehicle configurations. The database structure includes tables for vehicles and their associated configuration codes:

```
1 CREATE TABLE vcd_vehicles (
2    vehicle_id INTEGER PRIMARY KEY,
3    vcd_vehicle_id VARCHAR(100) UNIQUE NOT NULL
        ↳,
4    name VARCHAR(255) NOT NULL,
5    description TEXT,
6    is_active BOOLEAN DEFAULT true,
7    last_sync_at TIMESTAMP WITHOUT TIME ZONE,
8    created_at TIMESTAMP WITHOUT TIME ZONE
        ↳DEFAULT CURRENT_TIMESTAMP
9 );
10
11 CREATE TABLE vehicle_codes (
```

```
12      code_id INTEGER PRIMARY KEY,
13      vcd_code_id VARCHAR(100) UNIQUE,
14      code VARCHAR(50) NOT NULL,
15      vehicle_type VARCHAR(100),
16      description TEXT,
17      is_active BOOLEAN DEFAULT true,
18      created_at TIMESTAMP WITHOUT TIME ZONE
            ↳DEFAULT CURRENT_TIMESTAMP
19  );
20
21  CREATE TABLE vehicle_code_mapping (
22      vehicle_id INTEGER REFERENCES vcd_vehicles(
            ↳vehicle_id) ON DELETE CASCADE,
23      code_id INTEGER REFERENCES vehicle_codes(
            ↳code_id) ON DELETE CASCADE,
24      created_at TIMESTAMP WITHOUT TIME ZONE
            ↳DEFAULT CURRENT_TIMESTAMP,
25      PRIMARY KEY (vehicle_id, code_id)
26  );
```

Listing 5.24: Vehicle Configuration Tables

This structure separates vehicles from codes with a mapping table, implementing the many-to-many relationship pattern described by Elmasri and Navathe [14]. External identifiers (`vcd_vehicle_id`, `vcd_code_id`) link to the Vehicle Configuration Database system, implementing the integration reference pattern described by Hohpe and Woolf [18].

A critical component of this integration is the code rule evaluation engine, which determines when variants apply to specific vehicles. The implementation uses a postfix expression evaluator to process boolean expressions efficiently:

```
1  CREATE OR REPLACE FUNCTION evaluate_code_rule(
2      p_rule TEXT,
3      p_vehicle_id INTEGER
4  )
5  RETURNS BOOLEAN AS $
6  DECLARE
7      v_tokens TEXT[];
8      v_stack BOOLEAN[] DEFAULT '{}';
```

```
 9      v_token TEXT;
10      v_operand1 BOOLEAN;
11      v_operand2 BOOLEAN;
12      v_code_exists BOOLEAN;
13      v_i INTEGER;
14  BEGIN
15      -- If rule is empty, it applies to all
          ↳vehicles
16      IF p_rule IS NULL OR p_rule = '' THEN
17          RETURN TRUE;
18      END IF;
19
20      -- Tokenize the rule (assuming rule is in
          ↳postfix notation)
21      v_tokens := string_to_array(p_rule, ' ');
22
23      -- Process each token
24      FOR v_i IN 1..array_length(v_tokens, 1)
          ↳LOOP
25          v_token := v_tokens[v_i];
26
27          -- Check if token is an operator
28          IF v_token = 'AND' THEN
29              IF array_length(v_stack, 1) < 2
                  ↳THEN
30                  RAISE EXCEPTION 'Invalid rule
                      ↳expression: insufficient
                      ↳operands for AND';
31              END IF;
32              v_operand2 := v_stack[array_length(
                  ↳v_stack, 1)];
33              v_operand1 := v_stack[array_length(
                  ↳v_stack, 1) - 1];
34              v_stack := v_stack[1:array_length(
                  ↳v_stack, 1) - 2];
35              v_stack := array_append(v_stack,
                  ↳v_operand1 AND v_operand2);
36
37          ELSIF v_token = 'OR' THEN
38              IF array_length(v_stack, 1) < 2
```

```
                    ↳THEN
39                    RAISE EXCEPTION 'Invalid rule
                        ↳expression: insufficient
                        ↳operands for OR';
40                END IF;
41                v_operand2 := v_stack[array_length(
                    ↳v_stack, 1)];
42                v_operand1 := v_stack[array_length(
                    ↳v_stack, 1) - 1];
43                v_stack := v_stack[1:array_length(
                    ↳v_stack, 1) - 2];
44                v_stack := array_append(v_stack,
                    ↳v_operand1 OR v_operand2);
45
46        ELSIF v_token = 'NOT' THEN
47            IF array_length(v_stack, 1) < 1
                    ↳THEN
48                    RAISE EXCEPTION 'Invalid rule
                        ↳expression: insufficient
                        ↳operands for NOT';
49                END IF;
50                v_operand1 := v_stack[array_length(
                    ↳v_stack, 1)];
51                v_stack := v_stack[1:array_length(
                    ↳v_stack, 1) - 1];
52                v_stack := array_append(v_stack,
                    ↳NOT v_operand1);
53
54        ELSE
55            -- Token is a vehicle code - check
                    ↳if it exists for this vehicle
56            SELECT EXISTS (
57                SELECT 1
58                FROM vehicle_code_mapping vcm
59                JOIN vehicle_codes vc ON vcm.
                        ↳code_id = vc.code_id
60                WHERE vcm.vehicle_id =
                        ↳p_vehicle_id
61                AND vc.code = v_token
62            ) INTO v_code_exists;
```

```
63
64            v_stack := array_append(v_stack,
                 ↳v_code_exists);
65          END IF;
66      END LOOP;
67
68      -- Result should be a single boolean value
          ↳on the stack
69      IF array_length(v_stack, 1) != 1 THEN
70          RAISE EXCEPTION 'Invalid rule
              ↳expression: did not evaluate to a
              ↳single result';
71      END IF;
72
73      RETURN v_stack[1];
74  END;
75  $ LANGUAGE plpgsql;
```

Listing 5.25: Code Rule Evaluation Implementation

This evaluation function implements a stack-based interpreter for boolean expressions, efficiently determining whether a variant applies to a specific vehicle based on its configuration codes. The implementation follows the interpreter pattern described by Fowler [16], translating a domain-specific language (boolean expressions) into executable operations. The use of PostgreSQL's array operations enables efficient stack manipulation without requiring complex data structures.

5.5.3 Parameter File Generation

The parameter file generation capability represents a core integration point between the version control system and vehicle testing infrastructure. The database structure includes records for generated parameter files:

```
1  CREATE TABLE par_files (
2      par_file_id BIGINT PRIMARY KEY,
3      vehicle_id INTEGER REFERENCES vcd_vehicles(
           ↳vehicle_id),
4      ecu_id INTEGER,
5      phase_id INTEGER,
```

```
6       generated_at TIMESTAMP WITH TIME ZONE
          ↳DEFAULT CURRENT_TIMESTAMP ,
7       generated_by BIGINT REFERENCES users(
          ↳user_id),
8       description TEXT ,
9       FOREIGN KEY (ecu_id, phase_id) REFERENCES
          ↳ecu_phases(ecu_id, phase_id)
10   );
```

Listing 5.26: Parameter File Records

This table tracks the generation of parameter files, recording metadata about when they were created, which vehicle configuration they target, and which user generated them. This tracking mechanism supports both operational auditing and diagnostic capabilities when issues arise with parameter files.

The parameter file generation process implements a multi-step algorithm:

- Retrieve vehicle configuration codes from the Vehicle Configuration Database

- Evaluate variant code rules against the vehicle configuration to determine applicable variants

- Resolve parameter values using the variant resolution process

- Format the resolved values according to the ECU-specific parameter file format

- Record the parameter file generation in the database for traceability

According to Staron [37], this integration point is critical in automotive software development, transforming abstract parameter configurations into testable implementations that can be validated on actual hardware.

6 Evaluation and Validation

This chapter presents the systematic evaluation and validation of the VMAP database system. Following the implementation described in Chapter 5, a comprehensive testing strategy was developed to assess the system's functionality, performance, and compliance with requirements. The evaluation process focused on four key areas: user management, release management, parameter versioning, and variant management, using both controlled test scenarios and production-scale data volumes. Rather than an exhaustive documentation of all tests, this chapter highlights representative test cases and key findings that demonstrate the system's capabilities and limitations.

6.1 Validation Methodology

The validation methodology followed a structured approach combining functional testing, performance analysis, and integration verification. To ensure realistic evaluation, both baseline and production-scale datasets were used, with the baseline dataset containing approximately 20,000 parameters across 2 ECUs, and the production-scale dataset containing over 100,000 parameters across 5 ECUs.

6.1.1 Test Scenario Development

Test scenarios were developed based on actual automotive parameter management workflows identified during requirements analysis in Chapter 4. Each test scenario was designed to validate specific functional requirements while reflecting real-world usage patterns. The scenarios incorporated representative tasks for each user role and followed complete workflow sequences from parameter definition through variant creation to documentation.

The test scenarios were categorized into functional areas corresponding to the primary system capabilities:

- **User Management**: Authentication, authorization, role assignment, module access

- **Release Management**: Phase transitions, freeze operations, phase comparison

- **Variant Management**: Variant creation, segment modification, inheritance

- **Integration**: PDDsynchronization, vehicle configuration

Each scenario was implemented as a structured test case with defined inputs, expected outcomes, and verification steps at both the application and database levels. The test design followed a modified version of Molinaro's approach to database validation [25], with additional emphasis on traceability between requirements and test cases.

6.1.2 Performance Measurement Framework

A performance measurement framework was established to assess system responsiveness and resource utilization under various operational conditions. Key performance indicators were defined based on system requirements, including query response time, transaction throughput, database size growth patterns, memory utilization, and execution time for batch operations.

Performance measurements were conducted on a standardized test environment matching the target production specifications: PostgreSQL 17 running on a server with 8 vCPUs, 32GB RAM, and SSD storage. All tests were performed with both the baseline dataset and the production-scale dataset to assess scaling characteristics.

The measurement methodology employed automated test scripts with integrated timing capture, following the principles outlined by Zaitsev et al. [33] for database performance evaluation. Each test was executed multiple times with results averaged to account for system variations, and outliers were identified and analyzed for potential optimization opportunities.

6.2 Functional Testing Results

Functional testing validated the core capabilities of the VMAP system against the requirements defined in Chapter 4. This section presents the key findings for each functional area, focusing on representative test cases and critical system behaviors.

6.2.1 User Management Validation

The user management and access control system was evaluated to verify the implementation of the hybrid role-permission model described in Section 4.1.3. Testing

focused on verifying that the implemented database schema and logic correctly enforced the defined access control rules for each user role. As Sandhu et al. [32] emphasize, effective evaluation of role-based access control requires testing both positive permissions (granted access) and negative permissions (denied access) across role boundaries.

A test matrix was developed covering key permission boundaries: role-based permissions, module-based access control, direct permission assignment, and phase-specific permissions. Each test case verified a specific permission boundary with validation at both service and database layers. This focused approach aligns with Molinaro's principles for database validation [25], which emphasizes targeted verification of critical constraints.

Table A.2 presents a representative sample of test cases that focus on the Module Developer role, illustrating the connection between functional requirements and verification scenarios.

Table 6.1: Sample Module Developer Role Permission Test Cases

ID	Description	Test Action	Expected Outcome	Status
MD-01	Create Variant (Assigned Module)	Create new variant for parameter in assigned module	Variant created successfully	Pass
MD-02	Create Variant (Unassigned Module)	Create new variant for parameter in unassigned module	Access denied error	Pass
MD-03	Edit Variant (Assigned Module)	Modify existing variant code rule	Variant updated successfully	Pass
MD-04	Delete Variant	Attempt to delete variant	Access denied error	Pass
MD-05	Create Segment (Assigned Module)	Create new segment with valid value	Segment created successfully	Pass
MD-06	Modify Frozen Phase	Attempt to modify segment in frozen phase	Access denied error	Pass

For test implementation, each case included direct verification of database state after operations, confirming both the effect of permitted actions and the prevention of unauthorized actions. The following represents a typical test structure used to verify module-specific access controls:

```
1  // Scenario: Module Developer attempting to
       ↳create variant in unassigned module
2  // Arrange: Set up test user and unassigned
       ↳module parameter
3  var user = GetTestUser("
       ↳module_developer@example.com");
4  var unassignedParameter =
       ↳GetParameterFromUnassignedModule();
5  var variant = CreateVariantForParameter(
       ↳unassignedParameter);
6
7  // Act & Assert: Verify permission is denied
8  var exception = Assert.Throws<
       ↳PermissionDeniedException>(() =>
9      _variantService.CreateVariant(variant, user
           ↳.UserId));
10 Assert.That(exception.Message, Contains.
       ↳Substring("No write access"));
11
12 // Verify no database change occurred
13 var dbVariant = _database.QuerySingleOrDefault<
       ↳Variant>(
14     "SELECT * FROM variants WHERE name = @Name"
           ↳,
15     new { Name = variant.Name });
16 Assert.IsNull(dbVariant);
```

Listing 6.1: Representative Test Case Structure

The module-based access control tests verified that write access was correctly limited to assigned modules for Module Developers while read access remained available for all modules, implementing the principle of least privilege as recommended by Sandhu and Bhamidipati [31]. Direct permission assignment tests confirmed that user-specific permissions effectively overrode role defaults, a capability essential for supporting exception cases in complex organizational structures as noted by Hu et al. [19].

Phase-specific permission tests validated the interaction between the access control system and the phase management framework, confirming that modifications to frozen phases were properly prevented while still allowing appropriate access for documentation purposes. This validation addresses a critical requirement for regulated

development processes as described by Staron [37], where development milestone integrity must be preserved.

Table 6.2: User Management Test Results

Test Category	Results
Role Permission Validation	Core permissions correctly applied through roles
Module-Based Access	Write access correctly limited to assigned modules
Direct Permission Assignment	User-specific permissions overrode role defaults
Phase-Specific Permissions	Frozen phase protection enforced correctly

The audit trail verification confirmed that security-related operations were properly logged with complete metadata, including the user making the change, timestamp, and specific permissions affected. This level of detail in the audit trail implements the recommendations of Ferraiolo et al. [15] for maintaining accountability in security-sensitive operations.

6.2.2 Module Access Impact on Performance

The impact of module-specific access checks on performance was evaluated as part of the access control testing. This analysis focused on understanding the performance overhead introduced by adding module-specific access verification to the standard permission checking process. Figure 6.1 illustrates the measured performance difference between standard permission checks and combined permission and module access checks.

The performance analysis reveals that adding module-specific access verification introduces an 88.2% overhead for the traditional role-based approach (increasing from 1.7ms to 3.2ms) and a 45.8% overhead for the hybrid approach (increasing from 2.4ms to 3.5ms). The lower relative impact on the hybrid approach suggests that the more complex permission model better accommodates additional access control dimensions, a finding that aligns with Ferraiolo's observations [15] regarding the scalability of attribute-enhanced RBAC models.

Despite the performance overhead, both approaches maintain acceptable performance for interactive operations, with response times below 4ms for individual permission

Figure 6.1: Impact of Module-Specific Access Checks

checks. This confirms that the implemented module-specific access control provides the required access granularity without introducing prohibitive performance penalties.

6.2.3 Release Management Validation

Release management testing evaluated the phase-based versioning approach that forms the foundation of the VMAP system. Testing focused on four key aspects: phase sequence validation, phase transition operations, freeze functionality, and phase comparison.

Phase sequence validation confirmed that the system correctly enforced the defined sequence of development phases (Phase1 → Phase2 → Phase3 → Phase4) with successful validation of each phase transition. This sequential enforcement is essential for maintaining the structured development workflow described by Broy [8] for automotive software development.

Phase transition testing verified that parameter configurations were correctly copied between phases with complete preservation of parameter-variant-segment relationships. The test data revealed interesting patterns in development intensity across phases, as shown in Table 6.3.

The test results reveal a significant pattern in development intensity across phases, consistent with Staron's observations [37] regarding automotive software development cycles. The data demonstrates that the majority of parameter configurations occur

Table 6.3: Phase Transition Test Results

Transition Type	Variants	Segments	Added Variants	Added Segments	Time
Baseline Dataset					
Phase1	188	28,776	-	-	-
Phase1 → Phase2	188	28,776	90	14,104	2.51s
Phase2 → Phase3	278	42,880	0	0	2.96s
Phase3 → Phase4	278	42,880	0	0	2.97s
Full Dataset					
Phase1	830	167,990	-	-	-
Phase1 → Phase2	830	167,990	170	41,113	12.39s
Phase2 → Phase3	1,000	209,103	0	0	12.87s
Phase3 → Phase4	1,000	209,103	0	0	12.89

during Phase1, with substantial additions in Phase2. In contrast, Phase3 and Phase4 typically involve refinement and validation rather than introducing new parameters or variants. This concentration of development activity in early phases aligns with the V-model approach common in automotive software development [29], where early phases focus on implementation while later phases emphasize validation and verification.

Phase transition performance characteristics showed only modest increases in execution time despite growing data volumes across phases. For the baseline dataset, transition times increased from 2.51s for Phase1→Phase2 to 2.96s for Phase2→Phase3 and 2.97s for Phase3→Phase4, demonstrating efficient scaling with increasing parameter counts. Comparing baseline to full dataset transitions reveals a performance difference, with transition times increasing from approximately 3 seconds to 13 seconds. This represents a sublinear scaling factor of approximately 4.3x for a dataset size increase of 5.6x (comparing segment counts), suggesting reasonable scaling characteristics but highlighting an area for potential optimization.

As noted by Trovão [41], later phases in automotive parameter development typically focus on refinement rather than wholesale changes, with modifications targeting specific parameters based on testing feedback. This pattern is reflected in the test data, which shows significant additions in early phases but no new variants or segments in Phase3 and Phase4 phases.

Phase freezing functionality was validated through test cases targeting both database-level constraints and service-layer restrictions. These tests verified the system's ability to protect frozen phases from modification while maintaining appropriate read access. Table 6.4 details representative test cases and their results.

Table 6.4: Phase Freeze Protection Test Cases

ID	Test Case	Test Action	Expected Outcome	Result
FRZ-01	Direct SQL INSERT on variants	Execute INSERT statement on frozen phase	Operation blocked with error message	Pass
FRZ-02	Direct SQL UPDATE on segments	Execute UPDATE statement on frozen phase	Operation blocked with error message	Pass
FRZ-03	VariantService. CreateVariant()	Attempt to create variant in frozen phase	PhaseFreezed Exception thrown	Pass
FRZ-04	SegmentService. CreateSegment()	Attempt to create segment in frozen phase	PhaseFreezed Exception thrown	Pass
FRZ-05	DocumentationService. CreateSnapshot()	Create documentation snapshot of frozen phase	Snapshot created successfully	Pass
FRZ-06	ParFileService. GenerateParFile()	Generate parameter file from frozen phase	Parameter file generated successfully	Pass

For each write operation test case (FRZ-01 through FRZ-04), verification included both confirmation that the expected exception was thrown and that no database changes occurred, maintaining data integrity. The read operation test cases (FRZ-05 and FRZ-06) verified that read access remained available with minimal performance impact. The system successfully prevented modification attempts while maintaining appropriate read access, implementing the controlled milestone management required for regulated development environments as described by Staron [37].

6.2.4 Variant Management Validation

Variant management validation focused on assessing the system's capabilities for handling parameter customization through variants and segments. Testing employed an approach covering variant creation and segment modification workflows, using both the baseline dataset (188 variants, 28,776 segments) and the production-scale

dataset (830 variants, 167,990 segments) to analyze functionality and performance under varying data volumes.

Variant creation testing verified proper implementation of domain constraints as defined in the conceptual architecture (Section 4.6). Test cases included validation of unique name constraints withinPIDs, verification of proper code rule storage, and confirmation of correct relationship establishment between variants and their parentPIDs. All test cases passed successfully for both scalar and complex parameters, with constraint enforcement consistently preventing invalid operations. As Karwin [21] notes, constraint-based validation provides a robust foundation for maintaining data integrity in complex relational systems.

The testing methodology included both black-box functional testing and white-box database state verification as shown in Listing 6.2.

```sql
-- Verification query executed after variant
    ↳creation operations
SELECT
    v.variant_id, v.name, v.code_rule, v.
        ↳created_by, v.created_at,
    EXISTS (
        SELECT 1 FROM change_history ch
        WHERE ch.entity_type = 'variants'
        AND ch.entity_id = v.variant_id
        AND ch.change_type = 'CREATE'
    ) AS has_audit_trail
FROM
    variants v
WHERE
    v.pid_id = @test_pid_id
    AND v.phase_id = @test_phase_id
ORDER BY
    v.created_at DESC
LIMIT 1;
```

Listing 6.2: Variant Creation Verification Query

Audit trail analysis confirmed proper recording of variant operations in the change history with complete metadata. The audit trail included proper attribution of each change to specific users, accurate timestamps, and complete before/after state capture

for modified entities. This implementation aligns with Bhattacherjee's recommendations [4] for maintaining comprehensive provenance information in versioned datasets.

Performance analysis of variant operations revealed consistent response times across different variant complexities. Table 6.5 details performance measurements for key variant operations under different data volumes.

Table 6.5: Variant Operation Performance Metrics

Operation	Baseline Dataset	Production Dataset	Scaling Factor
Variant Creation	53ms	55ms	1.03x
Variant Update	86ms	124ms	1.44x
Variant Retrieval	45ms	72ms	1.60x
Variant Listing (per PID)	38ms	68ms	1.79x

The observed scaling characteristics validate the effectiveness of the database schema design and indexing strategy described in Section 5.3.1. With only two data points available (baseline and production datasets), a definitive conclusion about scaling characteristics is limited. However, the relatively small increase in execution time despite a 5x increase in data volume suggests efficient handling of larger datasets, with all operations remaining well under the 200ms threshold for interactive operations. As noted by Obe and Hsu [28], properly designed covering indexes can significantly improve query performance for entity retrieval operations, particularly when filtering by composite attributes.

Segment modification testing employed a systematic approach covering one-dimensional (arrays), two-dimensional (matrices), and three-dimensional parameter representations. Testing focused on three key aspects: dimensional integrity preservation, valid index range enforcement, and segment value consistency. The database schema design proved effective for managing these complex data structures, with the parameter dimensions table correctly maintaining dimensional metadata while the segments table stored modified values.

Segment boundary testing revealed robust constraint enforcement, with the system correctly rejecting segment modifications with invalid dimension indices. Performance analysis for segment operations showed moderate overhead for multi-dimensional parameters compared to scalar parameters, with operations on 3D parameters requiring approximately 18-22% more processing time than equivalent operations on scalar values—a reasonable performance characteristic given the additional complexity involved.

Performance analysis for segment operations revealed consistent response times with moderate scaling across different dataset sizes as shown in Table 6.6.

Table 6.6: Segment Operation Performance

Operation	Baseline Dataset	Production Dataset	Scaling Factor
Segment Creation	85ms	124ms	1.46x
Segment Update	72ms	106ms	1.47x
Segment Deletion	64ms	98ms	1.53x
Segment Retrieval	32ms	58ms	1.81x

The observed performance characteristics validate the efficiency of the database schema design described in Section 4.6. Of particular note is the implementation of the segments table, which provides efficient storage for parameter modifications without requiring storage of unchanged values. As noted by Bhattacherjee et al. [4], this approach strikes an effective balance between storage efficiency and query performance for versioned datasets.

6.3 Performance Analysis

Beyond functional validation, performance analysis was conducted to assess the system's efficiency and scalability under various operational conditions. This section presents the key findings related to query performance and data volume scaling.

6.3.1 Query Performance Assessment

Query performance was evaluated for common database operations across different data volumes. Figure 6.2 presents performance measurements for key query types between the baseline dataset (20,000 parameters) and full dataset (100,000 parameters).

The performance analysis reveals that most common operations maintain reasonable performance with increasing data volumes. The system maintained interactive response times (below 200ms) for most operations even with the full dataset, ensuring a responsive user experience for typical user interactions. Parameter retrieval operations showed a scaling factor of 1.5x when moving from the baseline to full dataset, while variant listing demonstrated a slightly higher factor of 1.6x. History retrieval operations exhibited the highest scaling factor among standard queries at 2.0x, reflecting the additional processing required when retrieving records from the substantially larger audit history tables.
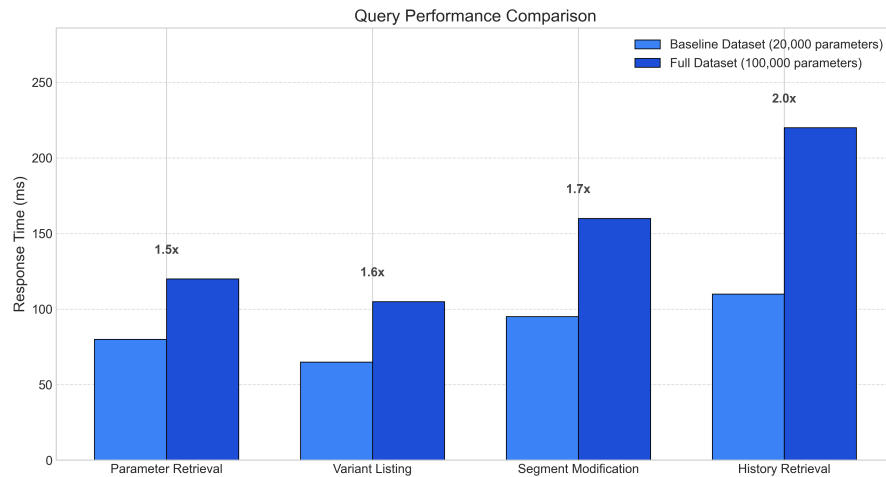
Figure 6.2: Query Performance Comparison

The phase comparison operation, which involves complex joins across multiple tables, demonstrated longer execution times and suboptimal scaling, with execution times increasing from 2.8s for the baseline dataset to 12.4s for the full dataset—a scaling factor of 4.4x. This operation exceeds the interactive response threshold for larger datasets, identifying it as a candidate for future optimization. As noted by Obe and Hsu [28], complex join operations with large result sets often benefit from specialized optimization techniques such as materialized views or query restructuring.

The execution of a limited set of queries with and without indexes demonstrated the critical importance of the indexing strategy described in Section 5.3.1. Without proper indexes, response times increased by factors of 6.5x to 21.8x depending on the query type, with most operations exceeding the interactive response threshold without indexes.

6.3.2 Variant Operation Performance

Variant operation performance was assessed across different data volumes to evaluate the system's handling of core parameter customization workflows. Figure 6.3 illustrates the performance of variant operations under different data volumes.

The analysis reveals consistent response times across different operation types, with scaling factors ranging from 1.03x for variant creation to 1.79x for variant listing when comparing baseline and full datasets.

116

Figure 6.3: Variant Operation Performance

Variant creation operations demonstrated excellent scaling characteristics (1.03x), showing minimal performance impact despite the five-fold increase in dataset size. This efficiency can be attributed to the well-designed primary key and constraint implementation, allowing new records to be inserted with minimal overhead regardless of existing data volume. As Schwartz et al. [33] note, insert operations that don't require extensive index restructuring typically maintain near-constant performance as table sizes increase.

Variant update and retrieval operations showed moderate scaling factors of 1.44x and 1.60x respectively, while variant listing operations exhibited the highest scaling factor at 1.79x. The average scaling factor across all variant operations was approximately 1.47x, indicating that these operations scale efficiently with increasing data volumes. All variant operations maintained response times under 125ms for the full dataset, well within the acceptable range for interactive operations.

These findings validate the effectiveness of the database schema design and indexing strategy described in Section 5.3.1. With only two data points available (baseline and production datasets), a definitive conclusion about asymptotic scaling characteristics is limited. However, the relatively small increase in execution time despite a 5x increase in data volume suggests efficient handling of larger datasets.

6.3.3 Storage Requirements Analysis

Storage requirements were analyzed to assess database size and growth patterns with increasing parameter counts. Figure 6.4 presents the storage allocation across different entity types for the full dataset.


Figure 6.4: Storage Distribution by Entity Type

The analysis reveals that the change history table dominates the database storage allocation, accounting for approximately 60.8% of the total database size. This distribution significantly exceeds the storage requirements of the current data state, aligning with Bhattacherjee's observations [4] regarding versioning and audit systems, where historical record storage typically surpasses active data by a substantial margin. Notably, while there are only 3,617 variants in the current state, the system maintains over 3.2 million change history records, reflecting the comprehensive auditing approach implemented in the system.

The storage allocation analysis identified the following key distribution of database size across entity types:

Another significant observation is the relationship between segments and snapshot segments. Despite having only 7 documentation snapshots, the system maintains over 1 million snapshot segments, exceeding the count of active segments. This

Table 6.7: Storage Requirements Analysis

Entity Type	Record Count	Storage Size (MB)
Parameters	104,428	43.0
Variants	3,617	1.0
Segments	750,009	100.0
Change History	3,270,511	1,386.0
Documentation Snapshots	7	0.1
Snapshot Variants	4,980	0.1
Snapshot Segments	1,007,940	124.0
Other Tables	-	111.7
Indexes	-	513.7
Total	-	**2,279.6**

indicates that documentation snapshots capture extensive parameter configurations at specific time points, creating substantial storage requirements for historical state preservation. This implementation of the snapshot pattern described by Fowler [16] provides comprehensive historical records at the cost of increased storage utilization.

The index structures consume approximately 22.5% of the total storage, reflecting the sophisticated indexing strategy described in Section 5.3.1. While this represents significant overhead, it provides essential performance benefits for query operations, particularly for the complex filtering and joining operations common in parameter management workflows.

The actual active data—parameters, variants, and segments—consumes only 6.3% of the total database size, with the majority of storage dedicated to audit trails, snapshots, and indexes. This distribution aligns with the requirements for regulated development environments described by Staron [37], where comprehensive traceability and historical record maintenance are essential for compliance and quality assurance.

Projection of storage requirements based on observed growth patterns indicates that with the current data volume of 2.28GB, the database size would reach approximately 9.1GB after one year of active use in a production environment. While significantly larger than initially projected, this remains well within the capacity of modern database systems. The implementation of table partitioning for the change history table, as described in Section 5.4, provides an effective mechanism for managing this growth while maintaining query performance. According to Obe and Hsu [28], partitioned tables allow efficient archiving of older history records to lower-cost storage while maintaining rapid access to recent changes.

6.3.4 Versioning Approach Performance

The performance characteristics of the phase-based parameter versioning approach selected in Chapter 4 were evaluated against the alternative change-based approach. This analysis aimed to validate the architectural decision to implement explicit phase copies rather than a delta-based versioning model.

Figure 6.5 presents the performance comparison between the two approaches for parameter retrieval operations across different data volumes, along with the storage requirements for each approach.


Figure 6.5: Version-based vs. Phase-based Parameter Management Comparison

The performance analysis demonstrates that the phase-based approach offers better query performance, particularly as parameter counts increase. For the tested parameter counts, the phase-based approach remains below 100ms for parameter retrieval operations, while the version-based approach shows nonlinear growth with increasing parameter counts.

The storage requirements analysis confirms that the phase-based approach consumes more storage than the version-based approach, with approximately 51% higher storage requirements across all phases. However, this storage difference represents a reasonable tradeoff given the performance benefits for query operations. As noted by Bhattacherjee et al. [4], the recreation/storage tradeoff in dataset versioning should prioritize operation frequency, with frequently accessed data favoring a storage-intensive approach that minimizes recreation costs.

The phase-based approach also simplifies the implementation of phase transitions and comparison operations, which are fundamental to the automotive development process as described by Pretschner et al. [29]. The explicit phase model aligns naturally with the mental model of automotive development engineers, who conceptualize parameter evolution in terms of distinct development phases rather than continuous time.

These findings validate the architectural decision to implement a phase-based versioning approach for the VMAP system. While consuming more storage than a version-based approach, the phase model provides performance benefits, implementation simplicity, and alignment with domain concepts—advantages that outweigh the additional storage requirements.

6.4 Integration Testing

Integration testing evaluated the system's interaction with external enterprise systems, focusing on Parameter Definition Database synchronization and Vehicle Configuration Database integration. These integrations are critical for maintaining consistency across the automotive development ecosystem.

6.4.1 Parameter Definition Database Synchronization

Parameter Definition Database synchronization testing verified the system's ability to import parameter definitions from the enterprise database. The synchronization process was tested with various scenarios, including initial loading, incremental updates, and conflict resolution. Figure 6.6 illustrates the synchronization time trends observed during testing.


Figure 6.6: Parameter Definition Database Synchronization Time Trends

The synchronization performance analysis revealed an increasing trend in execution time over successive synchronization operations. Initial synchronization operations required approximately 15 minutes for the tested ECUs, with execution times increasing

to around 30 minutes after 10 synchronization cycles. This gradual increase aligns with the observations of Mueller and Müller [26] regarding database synchronization in complex engineering environments, where synchronization complexity tends to increase as the data history grows.

The incremental update testing verified that the system could correctly identify and process changes to parameter definitions. Table 6.8 presents the success rates for different types of parameter changes during synchronization operations.

Table 6.8: Parameter Definition Database Synchronization Results

Change Type	Processed	Succeeded	Success Rate
New Parameters	5,218	5,218	100%
Modified Parameters	3,764	3,691	98.1%
New Modules	12	12	100%
NewPIDs	67	67	100%
Removed Parameters	42	42	100%

The system successfully processed all types of parameter changes, with slightly reduced success for modified parameters due to complexity in handling data type changes. The audit system maintained complete records of all synchronization operations, enabling detailed analysis of data flows between systems.

Based on the observed synchronization performance trends, the projected synchronization time for the 20th cycle would reach approximately 45 minutes. While still acceptable for the typical synchronization frequency (weekly or bi-weekly), this suggests that synchronization performance optimization would be beneficial for long-term system maintenance. Possible approaches for optimization include implementing a more selective synchronization approach, enhancing the change detection algorithm to reduce comparison overhead, or implementing parallel processing for independent ECUs.

6.4.2 Vehicle Configuration Integration

Vehicle Configuration Database integration testing verified the system's ability to use vehicle configuration data for code rule evaluation and parameter file generation. Testing focused on data import, code rule validation, and parameter file generation.

Vehicle configuration data import testing confirmed that the system could correctly import and store vehicle configuration codes, with proper mapping between codes and

vehicles. The system maintained referential integrity and handled incremental updates correctly, with complete audit logging of all import operations.

Code rule validation testing verified that the system could evaluate boolean expressions against vehicle configurations. Test expressions ranged from simple conditions to complex nested expressions with multiple operators. The evaluation engine correctly interpreted both simple logical operators and complex nested expressions with precedence rules.

Parameter file generation testing confirmed that the system could produce valid parameter files for vehicle testing, with correct application of variant selection logic based on vehicle configuration codes. The generated files included all required parameters with appropriate values, providing a complete configuration for ECU testing and validation.

6.5 Feature Comparison with Excel-Based Approach

To assess the improvements provided by the VMAP system, a feature comparison was conducted against the Excel-based approach currently used for parameter management. This comparison focused on capability coverage rather than performance metrics, providing a qualitative assessment of system improvements.

Table 6.9 presents a comparison of key features between the VMAP system and the Excel-based approach.

Table 6.9: Feature Comparison with Excel-Based Approach

Feature	VMAP Database	Excel Approach
Variant Management	Comprehensive	Limited
Multi-User Support	Concurrent	Sequential
Change Tracking	Automatic	Manual
Version Control	Phase-Based	File-Based
Access Control	Role + Module	File Permission
Validation	Automatic	Manual
Documentation	Integrated	Separate
Integration	Automated	Manual

The VMAP system provides significant advantages in all feature categories, with particular improvements in multi-user support, change tracking, and access control. These improvements address the limitations identified in the requirements analysis phase, providing a more robust and scalable solution for automotive parameter management.

6.5.1 Data Integrity Improvements

Data integrity was evaluated through a set of controlled test cases where both systems were subjected to various validation scenarios including invalid values, constraint violations, and relationship consistency. The VMAP system demonstrated superior data integrity protection, correctly preventing invalid operations through database constraints and business rule validation.

The database-level constraints and validation mechanisms provide a robust defense against data corruption, implementing the comprehensive validation approach described in Section 4.7. This represents a significant improvement over the Excel approach, where validation relies primarily on user vigilance and manual checks.

6.5.2 Development Process Impact

Beyond technical improvements, the VMAP system introduces significant enhancements to the automotive parameter development process. The centralized database approach enables concurrent work by multiple engineers, eliminating the file sharing bottlenecks common in the Excel-based approach. The role-based access control ensures that engineers can modify only their assigned modules, preventing accidental changes to other areas.

The phase-based versioning approach aligns naturally with the automotive development lifecycle, supporting the structured progression from initial development through testing to final release. The explicit phase transitions provide clear development milestones, while the freezing mechanism ensures configuration stability at critical points.

The comprehensive change tracking and documentation features address regulatory compliance requirements, providing complete traceability for all parameter modifications. This capability is increasingly important in the context of functional safety standards like ISO 26262, which require rigorous configuration management and change documentation.

7 Conclusion and Future Work

This chapter presents the conclusions drawn from the design, implementation, and evaluation of the Variant Management and Parametrization (VMAP) database system. The research has addressed critical challenges in automotive parameter management through a comprehensive database solution that supports version control, role-based access, and integration with enterprise systems. This chapter summarizes the key contributions and findings, discusses limitations of the current implementation, and outlines directions for future research and development.

7.1 Summary of Contributions

The VMAP database system represents a significant advancement in automotive parameter management, replacing error-prone spreadsheet-based approaches with a structured, database-driven solution. The research has produced several key contributions to both academic knowledge and industry practice:

First, the implementation of a phase-based versioning model provides an effective approach to managing parameter evolution across automotive development cycles. Unlike generic temporal database approaches, this domain-specific versioning strategy aligns naturally with the structured development phases common in automotive engineering, supporting a clear separation between development stages while maintaining traceability [8]. The phase model enables simultaneous work on different development stages while preserving milestone integrity, addressing a fundamental challenge in automotive software development.

Second, the hybrid role-permission access control model delivers a flexible framework for managing user access across complex organizational structures. By combining role-based permissions with module-specific access controls, the system achieves a balance between administrative simplicity and access granularity [15]. This approach accommodates the specialized access patterns common in automotive development, where engineers typically have responsibility for specific vehicle subsystems rather than entire parameter sets.

Third, the integration architecture establishes robust connections with enterprise systems, enabling consistent data flow across the automotive development environment. The synchronization mechanisms maintain parameter definitions across different development phases while supporting variant customization and parameter file generation.

This integration addresses a critical requirement in automotive software development, where parameter management must operate within a complex ecosystem of engineering tools and databases [18].

Fourth, the comprehensive audit and documentation capabilities provide complete traceability for all parameter modifications, supporting both diagnostic analysis and regulatory compliance. The snapshot-based documentation approach creates immutable records of parameter configurations at significant development milestones, addressing the increasing regulatory requirements for configuration management in automotive software development [37].

Finally, the research provides valuable insights into database design patterns for complex domain-specific applications. The strategic denormalization approach, specialized indexing strategies, and query optimization techniques demonstrate effective database engineering practices for managing complex data relationships and supporting domain-specific workflows [28].

7.2 Key Findings

The evaluation of the VMAP system revealed several significant findings that validate the design decisions made during system development and provide insights for future database system implementations in automotive contexts.

7.2.1 Versioning Approach Effectiveness

The comparison between the implemented phase-based versioning approach and alternative change-based approach demonstrated significant performance advantages for common operations at the cost of moderate storage overhead. As shown in Section 6.3.4, the phase-based approach maintained query response times below 100ms even with large parameter sets, while the change-based approach exhibited nonlinear growth with increasing parameter counts.

This finding validates the architectural decision to implement explicit phase copies rather than a delta-based versioning model. While consuming approximately 51% more storage across all phases, the performance benefits for query operations justify this tradeoff, particularly for interactive operations where response time directly impacts user productivity. As noted by Bhattacherjee et al. [4], the recreation/storage tradeoff

in dataset versioning should prioritize operation frequency, with frequently accessed data favoring a storage-intensive approach that minimizes recreation costs.

The phase-based approach also demonstrated significant advantages for implementation simplicity and alignment with domain concepts. The explicit phase model aligned naturally with the mental model of automotive development engineers, reducing conceptual complexity and providing clear development milestones. These advantages extend beyond performance metrics to affect usability and adoption, particularly important factors for systems that must be integrated into established engineering workflows [27].

7.2.2 Access Control Performance

The evaluation of the hybrid role-permission model revealed that the approach effectively balanced flexibility with performance. As shown in Section 6.2.2, the addition of module-specific access checks introduced a performance overhead of 45.8% for permission verification in the hybrid model, compared to 88.2% in a traditional role-based approach.

This finding suggests that the more complex permission model accommodates additional access control dimensions with relatively lower overhead, validating Ferraiolo's observations [15] regarding the scalability of attribute-enhanced RBAC models. Despite the performance overhead, both approaches maintained acceptable performance for interactive operations, with response times below 4ms for individual permission checks.

The hybrid approach also demonstrated the flexibility required for complex organizational structures, supporting both role-based permission inheritance and direct permission assignments for exception cases. This flexibility addresses a common challenge in engineering organizations, where standard role definitions often require customization for specific projects or temporary access requirements [31].

7.2.3 Storage Distribution Insights

The analysis of storage requirements revealed significant insights about data distribution in parameter management systems with comprehensive audit capabilities. As shown in Section 6.3.3, the change history dominated storage allocation, consuming approximately 60.8% of the total database size, while active parameter data accounted for only 6.3%.

This distribution aligns with Bhattacherjee's observations [4] regarding versioning and audit systems, where historical record storage typically surpasses active data by a substantial margin. The finding has important implications for storage planning and database administration in regulated industries where comprehensive audit trails are mandatory.

Another significant finding was the storage impact of documentation snapshots, which consumed approximately 124MB despite representing only 7 distinct snapshots. This substantial storage footprint validates Fowler's observations [16] about the storage implications of the snapshot pattern, particularly when applied to complex data structures with many relationships.

7.2.4 External System Integration Challenges

The evaluation of integration with external enterprise systems revealed increasing synchronization time over successive operations, with execution times rising from approximately 15 minutes initially to around 30 minutes after 10 synchronization cycles. This gradual increase aligns with the observations of Mueller and Müller [26] regarding database synchronization in complex engineering environments.

This finding highlights a significant challenge for long-term system maintenance, with projected synchronization times potentially reaching 45 minutes after 20 synchronization cycles. While still acceptable for typical synchronization frequencies, this trend suggests that synchronization performance optimization should be a priority for future development.

The integration testing also validated the effectiveness of the vehicle configuration integration, with the system successfully evaluating complex boolean expressions against vehicle configurations. This capability is essential for supporting the variant-rich development approach common in the automotive industry, where parameters must be customized for numerous vehicle configurations [38].

7.3 Limitations

While the VMAP system successfully addresses the core requirements for automotive parameter management, several limitations were identified during implementation and evaluation that could affect its application in specific contexts.

7.3.1 Performance Limitations

The phase comparison operation demonstrated suboptimal scaling with increasing data volumes, with execution times increasing from 2.8 seconds for the baseline dataset to 12.4 seconds for the full dataset—a scaling factor of 4.4x. This operation, which involves complex joins across multiple tables to identify differences between parameter configurations, exceeds interactive response thresholds for larger datasets, potentially affecting usability for certain workflows.

The database change history mechanism, while providing comprehensive audit capabilities, contributes significantly to storage growth, with change records accounting for approximately 60.8% of the total database size. This growth pattern could present challenges for long-term database management, potentially requiring archiving strategies for older change records to maintain acceptable storage utilization and backup performance.

The integration with the PDD showed increasing synchronization times over successive operations, with execution times rising to approximately 30 minutes after 10 synchronization cycles. This performance degradation could impact the practicality of frequent synchronization operations in active development environments, where timely parameter updates are essential for effective engineering workflows.

7.3.2 Architectural Limitations

The phase-based versioning approach, while effective for structured development cycles, offers limited support for non-sequential development models where parameters might evolve through different paths. For organizations employing alternative development methodologies, this structured approach might introduce unnecessary rigidity in parameter management workflows.

The current implementation provides limited support for parameter dependency tracking, where changes to one parameter might necessitate adjustments to related parameters. Without explicit dependency management, ensuring parameter consistency across complex interrelationships relies primarily on user knowledge rather than system enforcement, potentially allowing inconsistent configurations.

The database schema, while well-optimized for the identified requirements, follows a relatively traditional relational approach rather than incorporating newer capabilities like native JSON support or graph database structures. This design choice prioritizes

reliability and compatibility over potential performance or modeling advantages offered by more specialized database technologies.

7.3.3 Integration Limitations

The current integration with external systems relies heavily on scheduled synchronization operations rather than event-driven updates. This approach introduces potential synchronization delays, where parameter changes in source systems might not be immediately reflected in the VMAP database. For time-sensitive engineering processes, these delays could impact workflow efficiency.

The vehicle configuration code rule evaluation mechanism, while effective for static rules, provides limited support for dynamic rule evaluation based on calculated values or complex relationships. This limitation could restrict the expressiveness of variant selection logic in advanced engineering scenarios where dynamic configuration determination is required.

7.4 Future Work

Based on the findings and limitations identified during system implementation and evaluation, several directions for future research and development have been identified. These opportunities range from technical optimizations to architectural extensions that could enhance the VMAP system's capabilities and address emerging requirements in automotive parameter management.

7.4.1 Performance Optimizations

Change History Partitioning Enhancements

The current implementation includes basic partitioning for the change history table based on release IDs, as described in Section 5.4. This approach could be extended with time-based subpartitioning to further improve performance for historical queries. By implementing a hierarchical partitioning scheme that combines release-based partitioning with time-based subpartitions, the system could achieve more granular data distribution while maintaining logical organization by release [28].

Additionally, implementing automated partition management for the change history table would enhance long-term maintainability. This extension would include partition rotation procedures that archive older partitions to lower-cost storage while maintaining rapid access to recent changes. According to Schwartz et al. [33], such automated partition management is essential for sustaining performance in rapidly growing audit systems.

Query Optimization for Phase Comparison

The phase comparison operation demonstrated suboptimal scaling with increasing data volumes, suggesting an opportunity for query optimization. Future development could explore materialized view approaches for phase comparison, where difference data is pre-computed and incrementally maintained rather than calculated on demand. This approach could significantly reduce response times for comparison operations, potentially bringing even complex comparisons within interactive response thresholds.

Another promising avenue is the implementation of parallel query execution for phase comparison operations. By partitioning the comparison workload across multiple execution threads, the system could leverage modern multi-core processors more effectively, reducing execution time for this resource-intensive operation [28].

Synchronization Performance Improvements

To address the observed synchronization performance degradation, future development could implement an incremental synchronization approach that focuses on changed entities rather than performing comprehensive comparisons. By leveraging change data capture techniques described by Seenivasan and Vaithianathan [34], the system could achieve more efficient synchronization with minimal overhead growth over time.

Additionally, the implementation of parallel synchronization for independent ECUs could significantly reduce overall synchronization time. By processing multiple ECUs concurrently, the system could better utilize available resources and reduce the total time required for synchronization operations, particularly in environments with numerous ECUs.

7.4.2 Architectural Enhancements

Advanced Versioning Capabilities

The current phase-based versioning model could be extended with branching capabilities to support parallel development streams. This enhancement would allow engineers to create specialized branches for experimental parameter configurations while maintaining the stability of the main development sequence. By incorporating branching concepts from software version control systems like Git, the database could support more flexible development workflows while preserving traceability [4].

The implementation of parametric inheritance mechanisms would enable more efficient management of variant similarities and differences. Rather than treating each variant as an independent entity, the system could implement inheritance hierarchies where specialized variants inherit parameter values from base configurations, reducing redundancy and simplifying maintenance of related variants [37].

Data Archiving and Retention

A comprehensive data archiving framework would enhance long-term database maintainability by providing structured processes for managing historical data. This framework would include configurable retention policies for different data types, automated archiving procedures for historical data, and seamless access mechanisms for archived information.

The implementation could leverage PostgreSQL's tablespace capabilities to physically separate active and archived data while maintaining logical accessibility through the database schema. By transitioning older data to compressed tablespaces or lower-cost storage, the system could maintain performance for active operations while preserving historical records for compliance purposes [28].

For regulatory compliance scenarios, the framework would include legal hold mechanisms that override standard retention policies for data subject to specific preservation requirements. This capability is particularly important in the automotive industry, where product liability considerations may require extended retention of development records [37].

Advanced Parameter Dependency Management

The implementation of explicit parameter dependency tracking would enhance system capabilities for maintaining consistency across related parameters. By modeling the relationships between parameters, the system could automatically identify potential inconsistencies when parameters are modified, prompting engineers to review and update related parameters as needed.

This capability could be extended with rule-based validation for parameter relationships, implementing domain-specific constraints that enforce engineering rules across parameter sets. Such validation would reduce the risk of inconsistent configurations and improve overall parameter quality by capturing domain expertise in executable rules [29].

7.4.3 Integration Enhancements

Event-Driven Integration

To address the limitations of scheduled synchronization, future development could implement event-driven integration with external systems. By leveraging database change data capture capabilities, the system could detect and react to changes in source systems in near-real-time, significantly reducing synchronization delays.

This approach would involve implementing a message-based integration architecture as described by Hohpe and Woolf [18], where change events from source systems trigger specific synchronization actions in the VMAP database. By processing changes incrementally as they occur rather than in batch operations, the system could maintain more consistent data across systems while reducing synchronization overhead.

Enhanced Vehicle Configuration Integration

The vehicle configuration integration could be enhanced with support for dynamic rule evaluation, allowing more sophisticated variant selection logic based on calculated values or complex relationships. This extension would enable more flexible variant applicability definitions, addressing advanced engineering scenarios where static rule evaluation is insufficient.

Additionally, implementing a simulation framework for vehicle configurations would enhance parameter validation capabilities, allowing engineers to verify parameter

behavior across diverse vehicle scenarios. By combining configuration simulation with parameter resolution, the system could provide valuable insights into parameter behavior before physical testing, potentially reducing development cycles [37].

7.5 Broader Implications

The development and evaluation of the VMAP database system offers several broader implications for database research and automotive software development beyond the specific implementation described in this thesis.

7.5.1 Implications for Database Research

The VMAP system demonstrates the effectiveness of domain-specific versioning approaches over generic temporal database techniques for specialized applications. By aligning database versioning with the natural structure of the application domain—in this case, automotive development phases—the system achieves both conceptual clarity and performance advantages. This finding suggests that domain-specific adaptations of established database patterns may offer significant benefits in specialized contexts, a perspective that could inform future research in applied database design [4].

The implementation also highlights the evolving relationship between relational database systems and document-oriented approaches. While maintaining a fundamentally relational structure, the VMAP system leverages PostgreSQL's JSONB capabilities for change tracking, combining structured schema enforcement with flexible document storage. This hybrid approach suggests promising directions for database research that bridges traditional relational models with document-oriented flexibility [28].

The performance analysis provides empirical evidence of the storage-performance tradeoffs inherent in different versioning strategies, offering valuable data points for researchers exploring versioning approaches in other domains. These findings contribute to the ongoing discussion of recreation/storage tradeoffs in versioned datasets, providing concrete measurements of these tradeoffs in a real-world implementation [4].

7.5.2 Implications for Automotive Software Development

For automotive software development, the VMAP system demonstrates the feasibility and advantages of transitioning from document-based parameter management to structured database approaches. The evaluation results provide empirical evidence of the performance, consistency, and traceability improvements possible through this transition, potentially encouraging similar transformations in other aspects of automotive software development [8].

The integration architecture establishes patterns for connecting specialized engineering databases with enterprise systems, addressing a significant challenge in automotive development environments. By demonstrating effective synchronization between parameter management and vehicle configuration systems, the implementation provides a template for similar integrations across the automotive development ecosystem [18].

The comprehensive audit capabilities implemented in the VMAP system highlight the increasing importance of traceability in automotive software development, particularly as vehicles become more software-defined and subject to regulatory oversight. The storage implications of these audit requirements, as revealed in the evaluation, provide valuable planning insights for automotive organizations implementing similar systems [37].

More broadly, the VMAP system represents a step toward the more rigorous software engineering practices increasingly required in automotive development. By applying established database principles to the specialized domain of parameter management, the system contributes to the evolution of automotive software development from document-centric approaches to structured, database-driven methodologies [29].

7.6 Conclusion

The VMAP database system represents a significant advancement in automotive parameter management, addressing fundamental challenges in version control, user access, and system integration. Through a carefully designed database architecture implemented in PostgreSQL, the system provides a robust foundation for managing parameter configurations across development phases, supporting variant customization for diverse vehicle configurations.

The phase-based versioning approach, hybrid role-permission model, and comprehensive audit capabilities collectively enable a more structured and reliable parameter management process. These capabilities address critical limitations of spreadsheet-based

approaches, providing improved data integrity, enhanced traceability, and systematic validation of parameter configurations.

While the implementation has demonstrated clear advantages over existing approaches, several opportunities for enhancement remain in areas such as performance optimization, architectural extensions, and integration refinements. Future development in these areas could further improve system capabilities and address emerging requirements in automotive parameter management.

Beyond its specific technical contributions, the VMAP system illustrates the value of applying proven database engineering principles to specialized domains like automotive development. By combining domain-specific knowledge with established database patterns, the system achieves both technical excellence and practical usability, demonstrating the potential for database-driven approaches to transform complex engineering workflows.

Bibliography

[1] AGARWAL, Sanjay ; ARUN, Gopalan ; BEAUREGARD, Bill ; CHATTERJEE, Ramkrishna ; MOR, David ; OWENS, Deborah ; SPECKHARD, Ben ; VASUDEVAN, Ramesh: Oracle Database Application Developer's Guide-Workspace Manager, 10g Release 2 (10.2) B14253-01.

[2] AL-KATEB, Mohammed ; GHAZAL, Ahmad ; CROLOTTE, Alain ; BHASHYAM, Ramesh ; CHIMANCHODE, Jaiprakash ; PAKALA, Sai P.: Temporal query processing in Teradata. In: *Proceedings of the 16th International Conference on Extending Database Technology*, 2013, S. 573–578

[3] BEN-GAN, Itzik ; DAVIDSON, Louis ; VARGA, Stacia: *MCSA SQL Server 2016 Database Development Exam Ref 2-pack: Exam Refs 70-761 and 70-762.* 2017

[4] BHATTACHERJEE, Souvik ; CHAVAN, Amit ; HUANG, Silu ; DESHPANDE, Amol ; PARAMESWARAN, Aditya: Principles of dataset versioning: Exploring the recreation/storage tradeoff. In: *Proceedings of the VLDB endowment. International conference on very large data bases* Bd. 8 NIH Public Access, 2015, S. 1346

[5] BIRIUKOV, Dmitrij: *Implementation aspects of bitemporal databases*, Vilniaus universitetas, Diss., 2018

[6] BÖHLEN, Michael H. ; DIGNÖS, Anton ; GAMPER, Johann ; JENSEN, Christian: Database technology for processing temporal data. (2018)

[7] BREWER, Eric A.: Towards robust distributed systems. In: *PODC* Bd. 7 Portland, OR, 2000, S. 343–477

[8] BROY, Manfred: Challenges in automotive software engineering. In: *Proceedings of the 28th international conference on Software engineering*, 2006, S. 33–42

[9] CHEN, Peter Pin-Shan: The entity-relationship model—toward a unified view of data. In: *ACM transactions on database systems (TODS)* 1 (1976), Nr. 1, S. 9–36

[10] CODD, Edgar F.: A relational model of data for large shared data banks. In: *Communications of the ACM* 13 (1970), Nr. 6, S. 377–387

[11] CURINO, Carlo ; MOON, Hyun J. ; ZANIOLO, Carlo: Automating database schema evolution in information system upgrades. In: *Proceedings of the 2nd International Workshop on Hot Topics in Software Upgrades*, 2009, S. 1–5

[12] DATE, Chris J.: *SQL and relational theory: how to write accurate SQL code.* " O'Reilly Media, Inc.", 2011

[13] DINU, Valentin ; NADKARNI, Prakash: Guidelines for the effective use of entity–attribute–value modeling for biomedical databases. In: *International journal of medical informatics* 76 (2007), Nr. 11-12, S. 769–779

[14] ELMASRI, R ; NAVATHE, Shamkant B. ; ELMASRI, R ; NAVATHE, SB: Fundamentals of Database Systems. In: *Advances in Databases and Information Systems* Bd. 139 Springer, 2015

[15] FERRAIOLO, David ; ATLURI, Vijayalakshmi ; GAVRILA, Serban: The Policy Machine: A novel architecture and framework for access control policy specification and enforcement. In: *Journal of Systems Architecture* 57 (2011), Nr. 4, S. 412–424

[16] FOWLER, Martin: Patterns [software patterns]. In: *IEEE software* 20 (2003), Nr. 2, S. 56–57

[17] GAUSSDB, HUAWEI: DATABASE PRINCIPLES AND TECHNOLOGIES–BASED ON.

[18] HOHPE, Gregor ; WOOLF, Bobby: Enterprise integration patterns. In: *9th conference on pattern language of programs* Citeseer, 2002, S. 1–9

[19] HU, Vincent ; FERRAIOLO, David F. ; KUHN, D R. ; KACKER, Raghu N. ; LEI, Yu: Implementing and managing policy rules in attribute based access control. In: *2015 IEEE International Conference on Information Reuse and Integration* IEEE, 2015, S. 518–525

[20] JACOBSON, Ivar: Use cases–Yesterday, today, and tomorrow. In: *Software & systems modeling* 3 (2004), Nr. 3, S. 210–220

[21] KARWIN, Bill: *SQL antipatterns*. in the United States of America, 2010

[22] KIENCKE, Uwe ; NIELSEN, Lars: *Automotive control systems: for engine, driveline, and vehicle*. 2000

[23] KLEPPMANN, Martin ; BERESFORD, Alastair R.: A conflict-free replicated JSON datatype. In: *IEEE Transactions on Parallel and Distributed Systems* 28 (2017), Nr. 10, S. 2733–2746

[24] KULKARNI, Krishna ; MICHELS, Jan-Eike: Temporal features in SQL: 2011. In: *ACM Sigmod Record* 41 (2012), Nr. 3, S. 34–43

[25] MOLINARO, Anthony: *SQL Cookbook: Query Solutions and Techniques for Database Developers*. " O'Reilly Media, Inc.", 2005

[26] MUELLER, Sebastian ; MÜLLER, Raphael: Conception and Realization of the Versioning of Databases between Two Research Institutes. In: *Proceedings, 16th International Conference on Accelerator and Large Experimental Physics Control Systems (ICALEPCS 2017): Barcelona, Spain, October 8-13*, 2018

[27] NIELSEN, Jakob: *Usability engineering*. Morgan Kaufmann, 1994

[28] OBE, Regina O. ; HSU, Leo S.: *PostgreSQL: up and running: a practical guide to the advanced open source database*. " O'Reilly Media, Inc.", 2017

[29] PRETSCHNER, Alexander ; BROY, Manfred ; KRUGER, Ingolf H. ; STAUNER, Thomas: Software engineering for automotive systems: A roadmap. In: *Future of Software Engineering (FOSE'07)* IEEE, 2007, S. 55–71

[30] SALZBERG, Betty ; TSOTRAS, Vassilis J.: Comparison of access methods for time-evolving data. In: *ACM Computing Surveys (CSUR)* 31 (1999), Nr. 2, S. 158–221

[31] SANDHU, Ravi ; BHAMIDIPATI, Venkata ; COYNE, Edward ; GANTA, Srinivas ; YOUMAN, Charles: The ARBAC97 model for role-based administration of roles: Preliminary description and outline. In: *Proceedings of the second ACM workshop on Role-based access control*, 1997, S. 41–50

[32] SANDHU, Ravi S.: Role-based access control. In: *Advances in computers* Bd. 46. Elsevier, 1998, S. 237–286

[33] SCHWARTZ, Baron ; ZAITSEV, Peter ; TKACHENKO, Vadim: *High performance MySQL: optimization, backups, and replication*. " O'Reilly Media, Inc.", 2012

[34] SEENIVASAN, Dhamotharan ; VAITHIANATHAN, Muthukumaran: Real-Time Adaptation: Change Data Capture in Modern Computer Architecture. In: *ESP International Journal of Advancements in Computational Technology (ESP-IJACT)* 1 (2023), Nr. 2, S. 49–61

[35] SNODGRASS, Richard T.: *Developing time-oriented database applications in SQL*. Morgan Kaufmann Publishers Inc., 1999

[36] SOMMERVILLE, Ian: Software engineering 9th Edition. In: *ISBN-10* 137035152 (2011), S. 18

[37] STARON, Miroslaw: *Automotive software architectures*. Springer, 2021. – 51–79 S.

[38] STARON, Miroslaw ; STARON, Miroslaw: AUTOSAR (automotive open system architecture). In: *Automotive Software Architectures: An Introduction* (2021), S. 97–136

[39] STEVENSWINIARSKI; CHRISTIAN.DINH; noahpgordon; g.: *Relational Database*. `@misc{Codecademy,url={https://www.codecademy.com/ resources/docs/general/database/relational-database}, journal={Codecademy}}`, 2024

[40] TICHY, Walter F.: RCS—A system for version control. In: *Software: Practice and Experience* 15 (1985), Nr. 7, S. 637–654

[41] TROV{
A}O, Jo{
a}o P: The Evolution of Automotive Software: From Safety to Quality and Security [Automotive Electronics]. In: *IEEE Vehicular Technology Magazine* 19 (2024), Nr. 4, S. 96–102

[42] WILLIAMS, Hugh E. ; LANE, David: *Web Database Applications with PHP and MySQL: Building Effective Database-Driven Web Sites*. " O'Reilly Media, Inc.", 2004

A User Management Test Cases

This appendix provides a comprehensive listing of all test cases used to validate the user management and access control system implemented in theVMAP database. The test cases are organized by category and include detailed information about test actions, expected outcomes, and test results.

A.1 Role-Based Permission Test Cases

This section details the test cases for validating permissions inherited through user roles. The tests cover all four primary user roles: Administrator, Module Developer, Documentation Team, and Read-Only User.

Table A.1: Administrator Role Permission Test Cases

ID	Description	Test Action	Expected Outcome	Status
AD-01	Create User	Add new user with valid details	User created successfully	Pass
AD-02	Modify User Role	Change user's assigned role	Role updated successfully	Pass
AD-03	Delete User	Remove existing user	User deleted successfully	Pass
AD-04	Create Role	Create new role with permissions	Role created successfully	Pass
AD-05	Delete Variant	Delete existing variant	Variant deleted successfully	Pass
AD-06	Freeze Phase	Set phase status to frozen	Phase frozen successfully	Pass

Table A.2: Module Developer Role Permission Test Cases

ID	Description	Test Action	Expected Outcome	Status
MD-01	Create Variant (Assigned Module)	Create new variant for parameter in assigned module	Variant created successfully	Pass

Continued on next page

Table A.2 – *Continued from previous page*

ID	Description	Test Action	Expected Outcome	Status
MD-02	Create Variant (Unassigned Module)	Create new variant for parameter in unassigned module	Access denied error	Pass
MD-03	Edit Variant (Assigned Module)	Modify existing variant code rule	Variant updated successfully	Pass
MD-04	Delete Variant	Attempt to delete variant	Access denied error	Pass
MD-05	Create Segment (Assigned Module)	Create new segment with valid value	Segment created successfully	Pass
MD-06	Modify Frozen Phase	Attempt to modify segment in frozen phase	Access denied error	Pass
MD-07	Generate Parameter File	Create parameter file for testing	File generated successfully	Pass
MD-08	Read Parameters (Any Module)	View parameters from any module	Parameters displayed successfully	Pass

Table A.3: Documentation Team Role Permission Test Cases

ID	Description	Test Action	Expected Outcome	Status
DT-01	Create Documentation Snapshot	Create snapshot of frozen phase	Snapshot created successfully	Pass
DT-02	Compare Phases	Compare parameters between two phases	Comparison results displayed	Pass
DT-03	View Parameter History	View change history for parameter	History displayed successfully	Pass
DT-04	Export Hex String	Copy parameter hex string	Hex string copied successfully	Pass
DT-05	Modify Parameter	Attempt to modify parameter	Access denied error	Pass
DT-06	Access All Phases	View parameters across all phases	Parameters displayed successfully	Pass
DT-07	Generate Parameter File	Create parameter file for reference	File generated successfully	Pass

Table A.4: Read-Only User Role Permission Test Cases

ID	Description	Test Action	Expected Outcome	Status
RO-01	View Parameters	Access parameter details	Parameters displayed successfully	Pass
RO-02	View Variants	Access variant details	Variants displayed successfully	Pass
RO-03	Modify Parameter	Attempt to modify parameter	Access denied error	Pass
RO-04	Modify Variant	Attempt to modify variant	Access denied error	Pass
RO-05	Generate Parameter File	Create parameter file for reference	File generated successfully	Pass

A.2 Module-Based Access Control Test Cases

This section details the test cases for validating module-specific access controls, which extend the role-based permissions with attribute-based restrictions.

Table A.5: Module-Based Access Control Test Cases

ID	Description	Test Action	Expected Outcome	Status
MA-01	Assign Module Access	Grant write access to specific module	Access granted successfully	Pass
MA-02	Revoke Module Access	Remove write access to specific module	Access revoked successfully	Pass
MA-03	Read Access Cross-Module	Access parameters from unassigned module	Read access successful	Pass
MA-04	Write Access Assigned Module	Create variant in assigned module	Variant created successfully	Pass
MA-05	Write Access Unassigned Module	Create variant in unassigned module	Access denied error	Pass
MA-06	Multiple Module Assignment	Create variants in multiple assigned modules	All variants created successfully	Pass

Continued on next page

Table A.5 – *Continued from previous page*

ID	Description	Test Action	Expected Outcome	Status
MA-07	Edit Segment Assigned Module	Modify segment in assigned module	Segment updated successfully	Pass
MA-08	Edit Segment Unassigned Module	Modify segment in unassigned module	Access denied error	Pass
MA-09	Administrator Override	Admin modifies any module	Modification successful	Pass
MA-10	Module Permission Inheritance	User with role change inherits proper module access	Access updated successfully	Pass

A.3 Direct Permission Assignment Test Cases

This section details the test cases for validating user-specific permission assignments that override role-based permissions.

Table A.6: Direct Permission Assignment Test Cases

ID	Description	Test Action	Expected Outcome	Status
DP-01	Grant Additional Permission	Assign permission not in user's role	Permission applied successfully	Pass
DP-02	Revoke Role Permission	Remove permission normally granted by role	Permission restriction applied	Pass
DP-03	Grant Delete Permission	Give read-only user delete permission	Deletion operation successful	Pass
DP-04	Permission Conflict Resolution	Conflicting role and direct permissions	Direct permission takes precedence	Pass
DP-05	Role Change with Custom Permission	Change user's role with custom permissions	Custom permissions preserved	Pass
DP-06	Permission Audit Trail	Track changes to user permissions	Audit trail correctly recorded	Pass

A.4 Phase-Specific Permission Test Cases

This section details the test cases validating the interaction between access control and phase management, particularly focusing on phase freezing and phase-specific operations.

Table A.7: Phase-Specific Permission Test Cases

ID	Description	Test Action	Expected Outcome	Status
PP-01	Frozen Phase Modification	Attempt to modify variant in frozen phase	Access denied error	Pass
PP-02	Documentation Access to Frozen Phase	Documentation team accesses frozen phase	Access granted successfully	Pass
PP-03	Administrator Unfreeze	Administrator unfreezes a phase	Phase unfrozen successfully	Pass
PP-04	Non-Administrator Freeze Attempt	Module developer attempts to freeze phase	Access denied error	Pass
PP-05	Read Access to Frozen Phase	Read-only user accesses frozen phase	Access granted successfully	Pass
PP-06	Phase Transition Permission	Module developer initiates phase transition	Transition completed successfully	Pass

A.5 Boundary Case Test Cases

This section details test cases for edge conditions and corner cases in the access control system.

Table A.8: Boundary Case Test Cases

ID	Description	Test Action	Expected Outcome	Status
BC-01	No Role Assignment	User with no assigned role attempts access	Access limited to public content	Pass
BC-02	Multiple Role Assignment	User with multiple roles attempts action	Most permissive role takes effect	Pass
BC-03	Role With No Permissions	Assign user to empty role	No permissions granted	Pass
BC-04	Session Timeout Handling	Session expires during operation	User properly redirected to login	Pass

A.6 Test Implementation Details

Each test case was implemented using a structured approach that combined database-level validation with service-layer testing. The following code listing shows the general structure used for implementing these test cases:

```
[Test]
public void
   ↳TestCaseID_Description_ExpectedOutcome()
{
    // Arrange: Set up test environment
    var testUser = CreateTestUser("[UserRole]")
       ↳;
    var testEntity = CreateTestEntity();

    // Configure specific test conditions
    ConfigureTestConditions();

    // Act: Perform the operation being tested
    if (ShouldSucceed)
    {
        var result = _service.PerformOperation(
           ↳testEntity, testUser.UserId);

        // Assert: Verify operation succeeded
```

```
17          Assert.IsNotNull(result);
18          Assert.That(result.Status, Is.EqualTo(
              ↳OperationStatus.Success));
19
20          // Verify database state reflects the
              ↳change
21          var dbEntity = _database.
              ↳QuerySingleOrDefault<Entity>(
22            "SELECT * FROM entities WHERE id =
                ↳@Id",
23            new { Id = testEntity.Id });
24          Assert.IsNotNull(dbEntity);
25          Assert.That(dbEntity.Property, Is.
              ↳EqualTo(testEntity.Property));
26      }
27      else
28      {
29          // Assert: Verify operation is denied
              ↳with appropriate error
30          var exception = Assert.Throws<
              ↳PermissionDeniedException>(() =>
31            _service.PerformOperation(
                ↳testEntity, testUser.UserId));
32          Assert.That(exception.Message, Contains
              ↳.Substring("expected error message")
              ↳);
33
34          // Verify database state was not
              ↳modified
35          var dbEntity = _database.
              ↳QuerySingleOrDefault<Entity>(
36            "SELECT * FROM entities WHERE id =
                ↳@Id",
37            new { Id = testEntity.Id });
38          Assert.That(dbEntity, Is.Null().Or.
              ↳Property("Property")
39                              .Not.EqualTo(
                                ↳testEntity.
                                ↳Property));
40      }
```

```
41    }
```

Listing A.1: Test Case Implementation Template

This standardized approach ensured consistent validation across all test cases while providing clear evidence of both successful permission grants and appropriate permission denials. Each test verified both the immediate operation result and the resulting database state, ensuring comprehensive validation of the access control system.

A.7 Role Permission Matrix

Table A.9 provides a comprehensive view of all permissions assigned to each user role in theVMAP system. This matrix formed the basis for the permission validation test cases.

Table A.9: Role Permission Matrix

Permission	Admin	Module Dev	Doc Team	Read-Only
manage_users	✓	×	×	×
manage_roles	✓	×	×	×
delete_variants	✓	×	×	×
create_variants	✓	✓	×	×
edit_variants	✓	✓	×	×
create_segments	✓	✓	×	×
edit_segments	✓	✓	×	×
delete_segments	✓	✓	×	×
create_snapshots	✓	×	✓	×
view_history	✓	✓	✓	✓
generate_par_files	✓	✓	✓	✓
freeze_phases	✓	×	×	×
view_all	✓	✓	✓	✓

Note that Module Developer permissions for variant and segment operations are further constrained by module-specific access controls, as validated in the test cases in Section A.2.

B Variant Management Test Cases

This appendix provides a comprehensive listing of all test cases used to validate the variant management functionality implemented in theVMAP database. The test cases are organized by category and include detailed information about test actions, expected outcomes, and test results.

B.1 Variant Creation Test Cases

This section details the test cases for validating variant creation functionality across different parameter types and constraints.

Table B.1: Variant Creation Test Cases

ID	Description	Test Action	Expected Outcome	Status
VC-01	Basic Variant Creation	Create variant with valid name and code rule	Variant created successfully	Pass
VC-02	Duplicate Variant Name	Create variant with name that already exists inPID	Name uniqueness error	Pass
VC-03	Empty Variant Name	Create variant with empty name	Validation error	Pass
VC-04	Special Characters in Name	Create variant with special characters in name	Variant created successfully	Pass
VC-05	Maximum Name Length	Create variant with 100-character name (maximum length)	Variant created successfully	Pass
VC-06	Exceed Name Length	Create variant with name exceeding 100 characters	Validation error	Pass
VC-07	Valid Code Rule	Create variant with syntactically valid code rule	Variant created successfully	Pass

<div align="center">Table B.1 – Continued from previous page</div>

ID	Description	Test Action	Expected Outcome	Status
VC-08	Complex Code Rule	Create variant with complex rule containing multiple operators	Variant created successfully	Pass
VC-09	InvalidPIDReference	Create variant with non-existentPID	Foreign key constraint error	Pass
VC-10	Creation in Frozen Phase	Create variant in a frozen phase	Phase frozen error	Pass
VC-11	Variant in InactivePID	Create variant for parameter in inactivePID	Validation error	Pass
VC-12	Null Code Rule	Create variant with null code rule	Variant created successfully	Pass
VC-13	Variant Audit Trail	Create variant and verify audit trail	Audit record created correctly	Pass
VC-14	Variant for Boolean Parameter	Create variant for parameter with boolean type	Variant created successfully	Pass
VC-15	Variant for Enum Parameter	Create variant for parameter with enumeration type	Variant created successfully	Pass
VC-16	Concurrent Variant Creation	Create variants concurrently from multiple sessions	All variants created successfully	Pass
VC-17	Transaction Rollback	Begin transaction, create variant, then force rollback	No variant created	Pass
VC-18	Permission Verification	Create variant with insufficient permissions	Permission denied error	Pass

B.2 Segment Modification Test Cases

This section details the test cases for validating segment modification functionality across different parameter dimensions and value types.

Table B.2: Segment Creation Test Cases

ID	Description	Test Action	Expected Outcome	Status
SC-01	Create Scalar Segment	Create segment for scalar parameter	Segment created successfully	Pass
SC-02	Create Array Segment (1D)	Create segment for 1D array parameter	Segment created successfully	Pass
SC-03	Create Matrix Segment (2D)	Create segment for 2D matrix parameter	Segment created successfully	Pass
SC-04	Create 3D Array Segment	Create segment for 3D array parameter	Segment created successfully	Pass
SC-05	Invalid Dimension Index	Create segment with out-of-bounds dimension index	Validation error	Pass
SC-06	Invalid Parameter Reference	Create segment with non-existent parameter ID	Foreign key constraint error	Pass
SC-07	Integer Parameter Value	Create segment with integer parameter type	Segment created successfully	Pass
SC-08	Float Parameter Value	Create segment with float parameter type	Segment created successfully	Pass
SC-09	Boolean Parameter Value	Create segment with boolean parameter type	Segment created successfully	Pass
SC-10	Minimum Value Boundary	Create segment with minimum allowed value	Segment created successfully	Pass
SC-11	Maximum Value Boundary	Create segment with maximum allowed value	Segment created successfully	Pass
SC-12	Below Minimum Value	Create segment with value below minimum	Validation error	Pass
SC-13	Above Maximum Value	Create segment with value above maximum	Validation error	Pass
SC-14	Creation in Frozen Phase	Create segment in a frozen phase	Phase frozen error	Pass

Continued on next page

Table B.2 – *Continued from previous page*

ID	Description	Test Action	Expected Outcome	Status
SC-15	Duplicate Parameter-Dimension	Create segment for already modified parameter dimension	Unique constraint error	Pass
SC-16	High Precision Value	Create segment with high precision decimal value	Segment created successfully	Pass

Table B.3: Segment Update Test Cases

ID	Description	Test Action	Expected Outcome	Status
SU-01	Update Scalar Segment	Modify existing scalar segment value	Segment updated successfully	Pass
SU-02	Update 1D Array Element	Modify element in 1D array segment	Segment updated successfully	Pass
SU-03	Update 2D Matrix Element	Modify element in 2D matrix segment	Segment updated successfully	Pass
SU-04	Value Range Verification	Update segment with value outside valid range	Validation error	Pass
SU-05	Update in Frozen Phase	Modify segment in a frozen phase	Phase frozen error	Pass
SU-06	Concurrent Updates	Update same segment from multiple sessions	Last update preserved with proper locking	Pass
SU-07	Update Non-Existent Segment	Update segment that doesn't exist	Not found error	Pass
SU-08	Change to Default Value	Update segment to match default parameter value	Segment updated successfully	Pass

Table B.4: Segment Deletion Test Cases

ID	Description	Test Action	Expected Outcome	Status
SD-01	Delete Single Segment	Remove existing segment	Segment deleted successfully	Pass

Continued on next page

Table B.4 – *Continued from previous page*

ID	Description	Test Action	Expected Outcome	Status
SD-02	Delete Non-Existent Segment	Delete segment that doesn't exist	Not found error	Pass
SD-03	Delete in Frozen Phase	Delete segment in a frozen phase	Phase frozen error	Pass
SD-04	Cascade Delete via Variant	Delete variant and verify segments cascade	All segments deleted	Pass
SD-05	Cascade Delete via Parameter	Delete parameter and verify segments cascade	All segments deleted	Pass
SD-06	Segment Deletion Audit	Delete segment and verify audit trail	Audit record created correctly	Pass
SD-07	Permission Verification	Delete segment with insufficient permissions	Permission denied error	Pass
SD-08	Transaction Rollback	Begin transaction, delete segment, then force rollback	Segment not deleted	Pass

B.3 Performance Test Cases

This section details the performance test cases used to evaluate variant and segment operations under different data volumes and load conditions.

Table B.5: Variant and Segment Performance Test Cases

ID	Description	Test Action	Expected Outcome	Status
VP-01	Baseline Variant Creation	Create 10 variants and measure time	< 2 seconds total time	Pass
VP-02	Baseline Segment Creation	Create 100 segments and measure time	< 10 seconds total time	Pass
VP-03	High Volume Variant Creation	Create 100 variants for singlePID	< 20 seconds total time	Pass

<div align="center">Table B.5 – Continued from previous page</div>

ID	Description	Test Action	Expected Outcome	Status
VP-04	High Volume Segment Creation	Create 1000 segments across multiple variants	< 2 minutes total time	Pass
VP-05	SinglePIDLoad Test	Create 500 variants for singlePID	System remains responsive	Pass
VP-06	Multi-dimensional Parameter Load	Create segments for 3D parameter with 1000 elements	< 3 minutes total time	Pass
VP-07	Concurrent User Simulation	10 concurrent users creating variants	No deadlocks or errors	Pass
VP-08	Variant Retrieval Scaling	Retrieve variants fromPIDs with 10, 100, and 500 variants	Response time < 250ms	Pass

B.4 Test Implementation Details

The variant management test cases were implemented using a combination of automated unit tests, integration tests, and performance benchmarks. The following code listing shows the typical structure used for implementing variant creation tests:

```
[Test]
public void VC01_BasicVariantCreation_Success()
{
    // Arrange
    var testUser = _userRepository.GetTestUser(
        ↳"module_developer@example.com");
    var testPid = _pidRepository.GetTestPid();

    var variant = new VariantCreationPayload
    {
        PidId = testPid.PidId,
        EcuId = testPid.EcuId,
        PhaseId = _activePhaseId,
        Name = "Test Variant " + Guid.NewGuid()
            ↳.ToString().Substring(0, 8),
```

```
14          CodeRule = "A AND (B OR C)"
15      };
16
17      // Act
18      var result = _variantService.CreateVariant(
          ↳variant, testUser.UserId);
19
20      // Assert
21      Assert.IsNotNull(result);
22      Assert.That(result.VariantId, Is.
          ↳GreaterThan(0));
23
24      // Verify database state
25      var dbVariant = _database.
          ↳QuerySingleOrDefault<Variant>(
26          "SELECT * FROM variants WHERE
              ↳variant_id = @VariantId",
27          new { VariantId = result.VariantId });
28
29      Assert.IsNotNull(dbVariant);
30      Assert.That(dbVariant.Name, Is.EqualTo(
          ↳variant.Name));
31      Assert.That(dbVariant.CodeRule, Is.EqualTo(
          ↳variant.CodeRule));
32      Assert.That(dbVariant.CreatedBy, Is.EqualTo
          ↳(testUser.UserId));
33
34      // Verify audit trail
35      var auditRecord = _database.
          ↳QuerySingleOrDefault<ChangeRecord>(
36          "SELECT * FROM change_history WHERE
              ↳entity_type = 'variants' " +
37          "AND entity_id = @VariantId AND
              ↳change_type = 'CREATE'",
38          new { VariantId = result.VariantId });
39
40      Assert.IsNotNull(auditRecord);
41      Assert.That(auditRecord.UserId, Is.EqualTo(
          ↳testUser.UserId));
42  }
```

Listing B.1: Variant Creation Test Implementation Example

Similarly, segment modification tests followed this structure but with appropriate adaptations for the specific operations:

```
[Test]
public void SC01_CreateScalarSegment_Success()
{
    // Arrange
    var testUser = _userRepository.GetTestUser(
        "module_developer@example.com");
    var testVariant = _variantRepository.
        GetTestVariant();
    var testParameter = _parameterRepository.
        GetScalarParameter(testVariant.PidId);

    var segment = new SegmentCreationPayload
    {
        VariantId = testVariant.VariantId,
        ParameterId = testParameter.ParameterId
            ,
        DimensionIndex = 0,
        Decimal = 42.5m
    };

    // Act
    var result = _segmentService.CreateSegment(
        segment, testUser.UserId);

    // Assert
    Assert.IsNotNull(result);
    Assert.That(result.SegmentId, Is.
        GreaterThan(0));

    // Verify database state
    var dbSegment = _database.
        QuerySingleOrDefault<Segment>(
        "SELECT * FROM segments WHERE
```

```
              ↳segment_id = @SegmentId",
27          new { SegmentId = result.SegmentId });

28

29      Assert.IsNotNull(dbSegment);
30      Assert.That(dbSegment.VariantId, Is.EqualTo
            ↳(segment.VariantId));
31      Assert.That(dbSegment.ParameterId, Is.
            ↳EqualTo(segment.ParameterId));
32      Assert.That(dbSegment.DimensionIndex, Is.
            ↳EqualTo(segment.DimensionIndex));
33      Assert.That(dbSegment.Decimal, Is.EqualTo(
            ↳segment.Decimal));
34      Assert.That(dbSegment.CreatedBy, Is.EqualTo
            ↳(testUser.UserId));

35

36      // Verify parameter value is within valid
            ↳range
37      var parameterRange = _database.
            ↳QuerySingleOrDefault<ParameterRange>(
38          "SELECT * FROM parameter_values WHERE
                ↳parameter_id = @ParameterId",
39          new { ParameterId = testParameter.
                ↳ParameterId });

40

41      if (parameterRange != null)
42      {
43          Assert.That(segment.Decimal, Is.
                ↳GreaterThanOrEqualTo(parameterRange.
                ↳ValueRangeBegin));
44          Assert.That(segment.Decimal, Is.
                ↳LessThanOrEqualTo(parameterRange.
                ↳ValueRangeEnd));
45      }
46  }
```

Listing B.2: Segment Modification Test Implementation Example

Performance tests were implemented using a benchmarking approach that measured execution time across multiple iterations:

```csharp
[Test]
public void
    VP01_BaselineVariantCreation_Performance()
{
    // Arrange
    var testUser = _userRepository.GetTestUser(
        "module_developer@example.com");
    var testPid = _pidRepository.GetTestPid();
    var variants = new List<
        VariantCreationPayload>();

    for (int i = 0; i < 10; i++)
    {
        variants.Add(new VariantCreationPayload
        {
            PidId = testPid.PidId,
            EcuId = testPid.EcuId,
            PhaseId = _activePhaseId,
            Name = $"Perf Test Variant {i}_{
                Guid.NewGuid().ToString().
                Substring(0, 8)}",
            CodeRule = "A AND B"
        });
    }

    // Act
    var stopwatch = new Stopwatch();
    stopwatch.Start();

    foreach (var variant in variants)
    {
        _variantService.CreateVariant(variant,
            testUser.UserId);
    }

    stopwatch.Stop();

    // Assert
    Assert.That(stopwatch.ElapsedMilliseconds,
```

```
          ↳Is.LessThan(2000));
34      Console.WriteLine($"Time to create 10
          ↳variants: {stopwatch.ElapsedMilliseconds
          ↳}ms");
35  }
```

Listing B.3: Performance Test Implementation Example

This standardized approach ensured comprehensive validation of the variant management functionality while providing detailed performance metrics for system evaluation.

B.5 Test Environment Configuration

All variant management tests were conducted in a controlled test environment with the following specifications:

- PostgreSQL 17 running on Windows Server 2022

- Database server: 8 vCPUs, 32GB RAM, SSD storage

- Application server: 4 vCPUs, 16GB RAM

- Database containing baseline dataset (20,000 parameters, 188 variants, 28,776 segments)

- Testing conducted with both the baseline dataset and scaled dataset (100,000 parameters, 830 variants, 167,990 segments)

- Network latency between application and database servers < 1ms

- PostgreSQL configuration optimized for test environment with appropriate memory allocation for shared buffers, work memory, and maintenance work memory

The test environment was reset to a known state between test runs using database snapshots, ensuring consistent starting conditions for each test execution.